

# Announcements

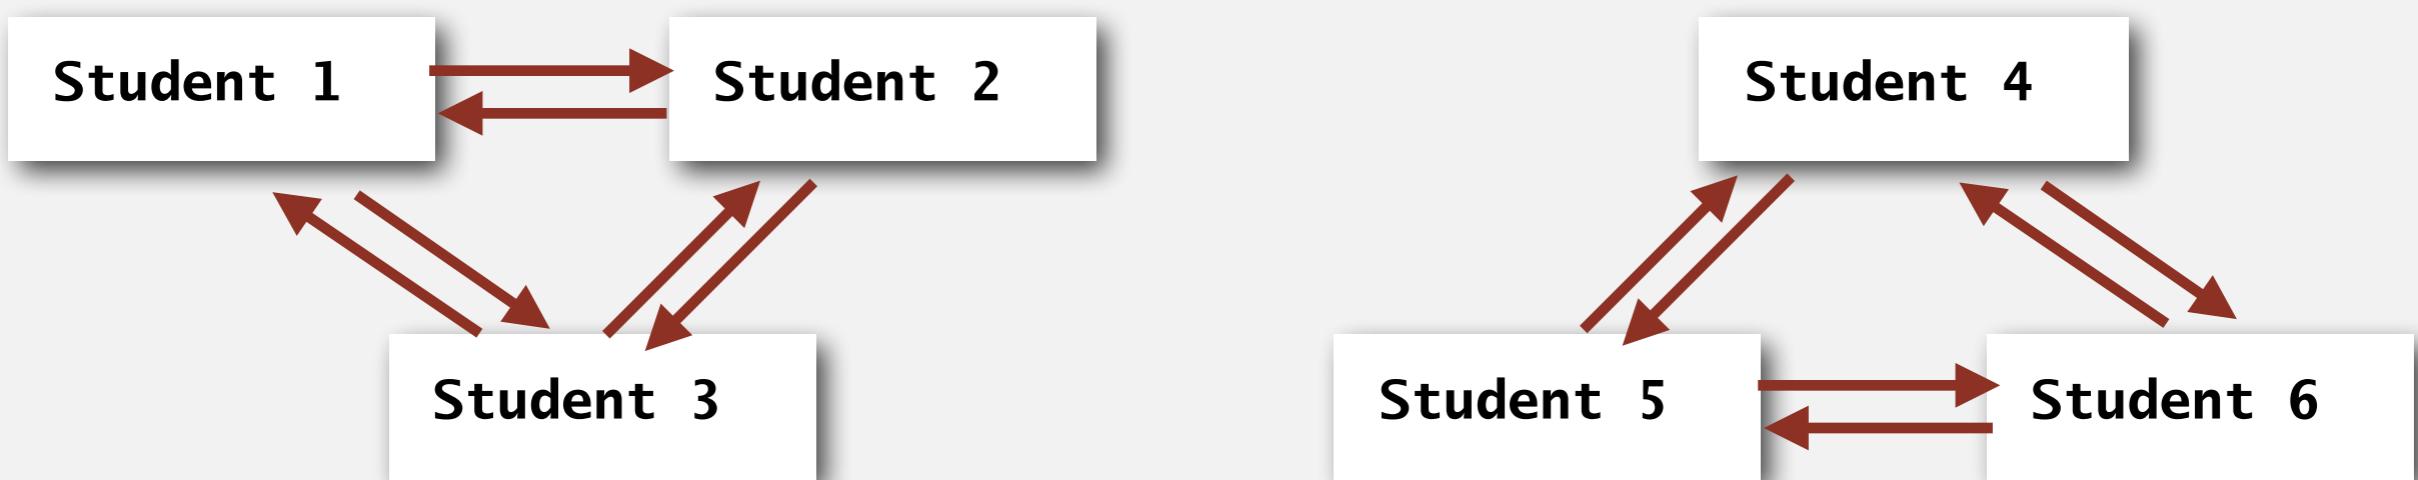
---

The term project will be released soon. The first deliverable will be due 1 week from today.

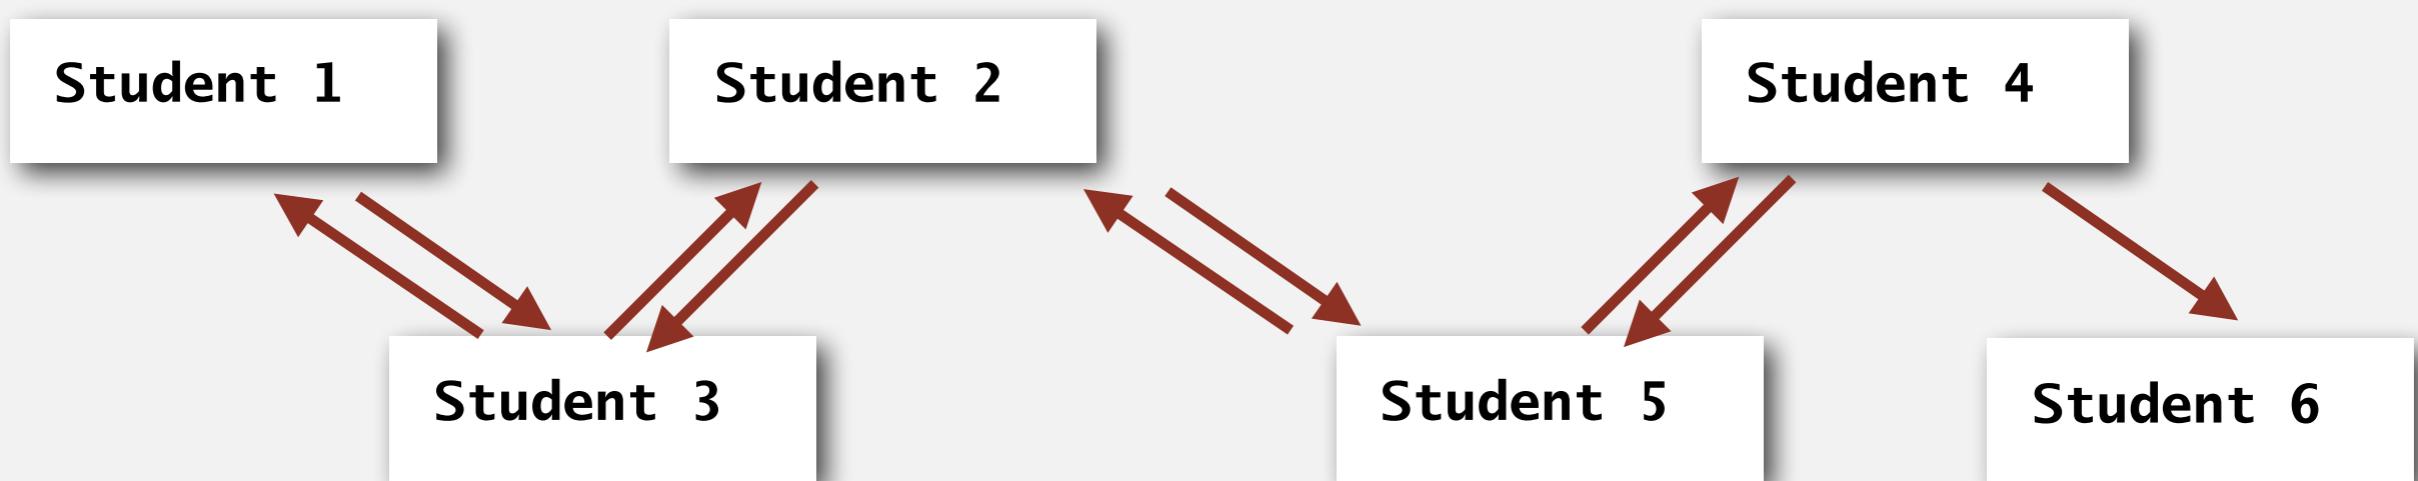
## Clarification on collaboration policy for written homework:

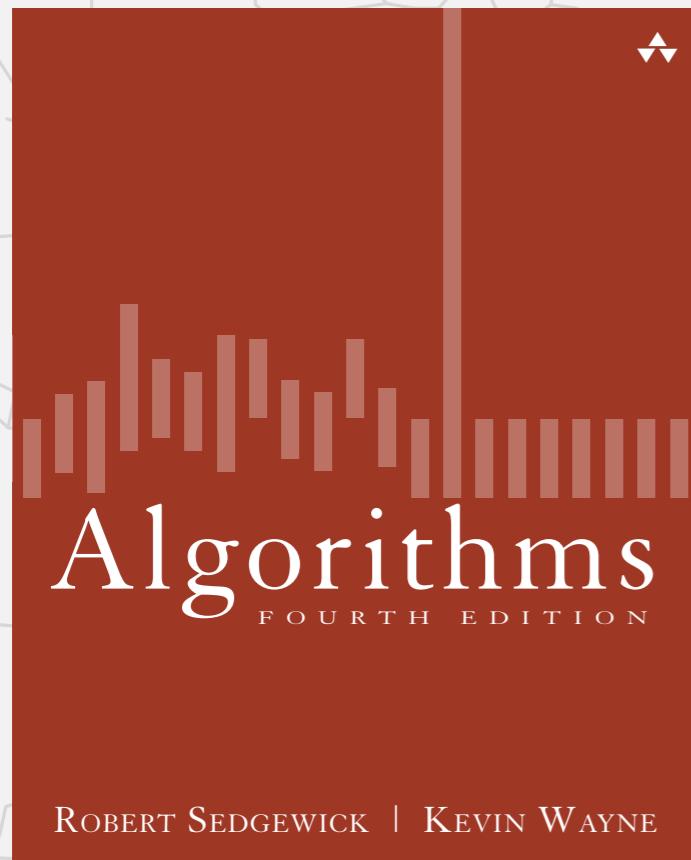
You are allowed to discuss solutions to problems in groups of three, *documenting* who you discussed with at the top of your assignment. You are allowed to ask anyone for LaTeX help (for instance, “How do I center a stack of equations?”). You are **not** allowed to write up the solutions together.

### Allowed



### Not Allowed





<http://algs4.cs.princeton.edu>

## 3.3 BALANCED SEARCH TREES

---

- ▶ *2–3 search trees*
- ▶ *red–black BSTs*

# BST: ordered symbol table operations summary

	sequential search	binary search	BST
search	$N$	$\log N$	$h$
insert	$N$	$N$	$h$
min / max	$N$	1	$h$
floor / ceiling	$N$	$\log N$	$h$
rank	$N$	$\log N$	$h$
select	$N$	1	$h$
ordered iteration	$N \log N$	$N$	$N$

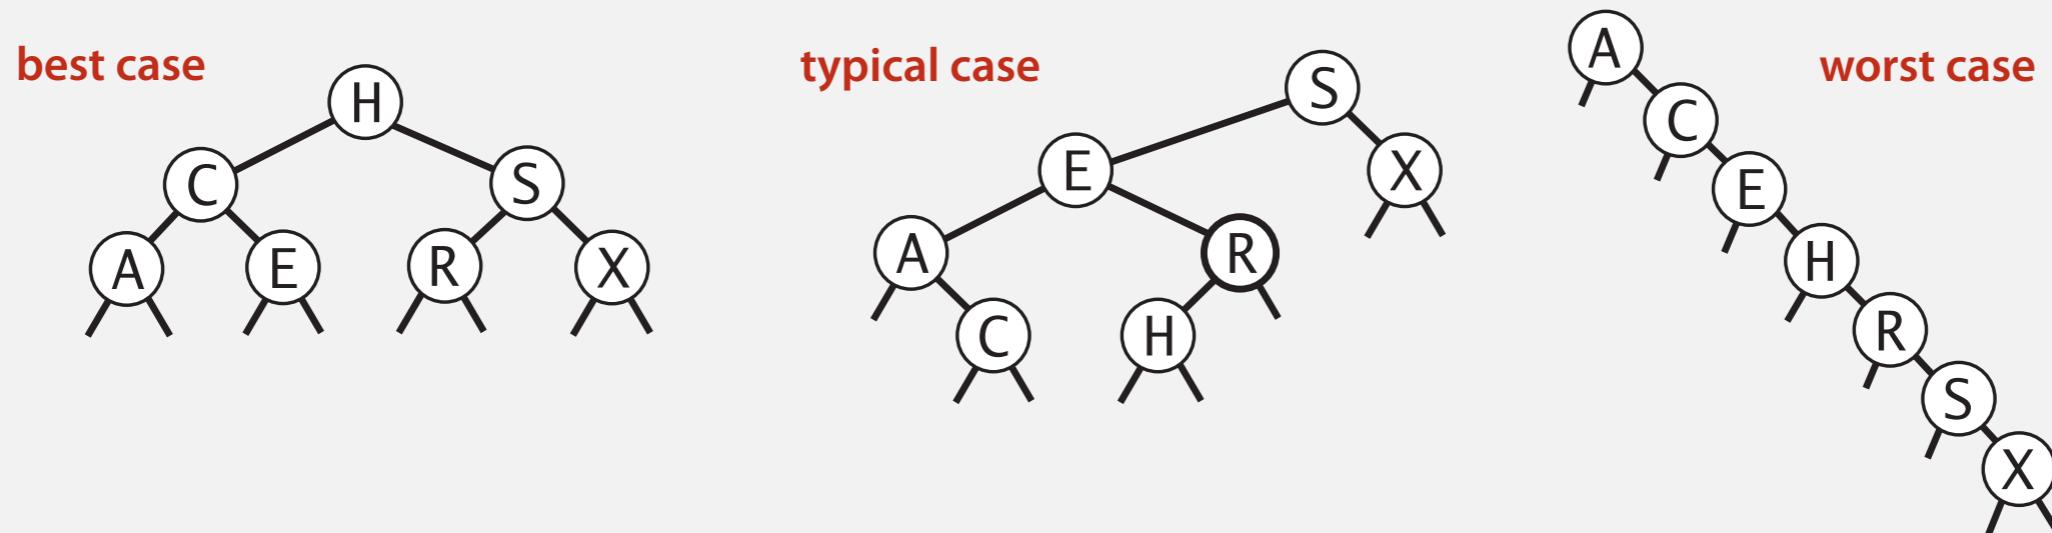
$h =$  height of BST  
(proportional to  $\log N$   
if keys inserted in random order)

order of growth of running time of ordered symbol table operations

# Tree shape

---

- Many BSTs correspond to same set of keys.
- Number of compares for search/insert = 1 + depth of node.



Bottom line. Tree shape depends on order of insertion.

# Symbol table review

implementation	guarantee			average case			ordered ops?	key interface
	search	insert	delete	search hit	insert	delete		
sequential search (unordered list)	$N$	$N$	$N$	$N$	$N$	$N$		<code>equals()</code>
binary search (ordered array)	$\log N$	$N$	$N$	$\log N$	$N$	$N$	✓	<code>compareTo()</code>
BST	$N$	$N$	$N$	$\log N$	$\log N$	$\sqrt{N}$	✓	<code>compareTo()</code>
goal	$\log N$	$\log N$	$\log N$	$\log N$	$\log N$	$\log N$	✓	<code>compareTo()</code>

Challenge. Guarantee performance.

optimized for teaching and coding;  
introduced to the world in this course!

This lecture. 2–3 trees and left-leaning red–black BSTs. Textbook: B-trees.

# Algorithms

ROBERT SEDGEWICK | KEVIN WAYNE

<http://algs4.cs.princeton.edu>

## 3.3 BALANCED SEARCH TREES

---

- ▶ 2–3 search trees
- ▶ red-black BSTs

## 2-3 tree

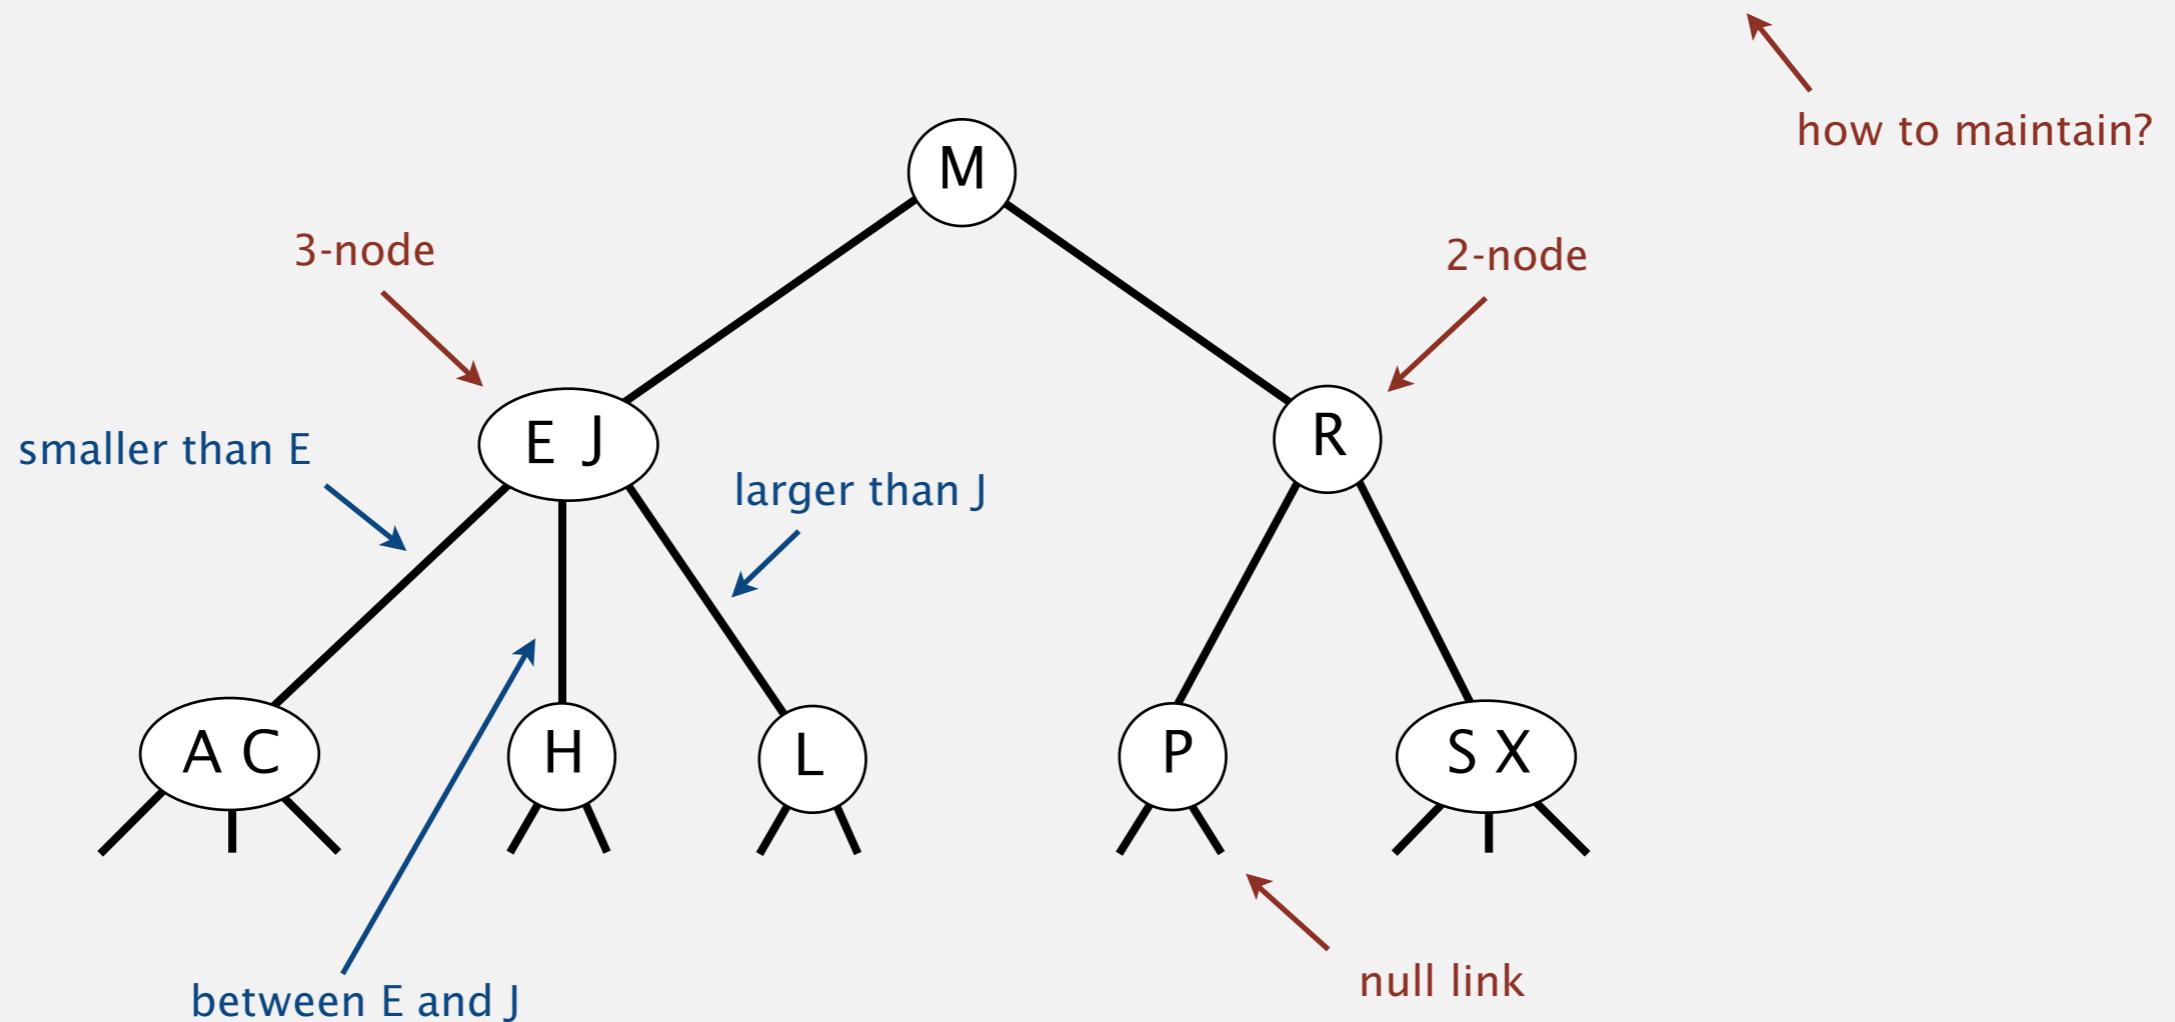
---

Allow 1 or 2 keys per node.

- 2-node: one key, two children.
- 3-node: two keys, three children.

Symmetric order. Inorder traversal yields keys in ascending order.

Perfect balance. Every path from root to null link has same length.



## 2-3 tree demo

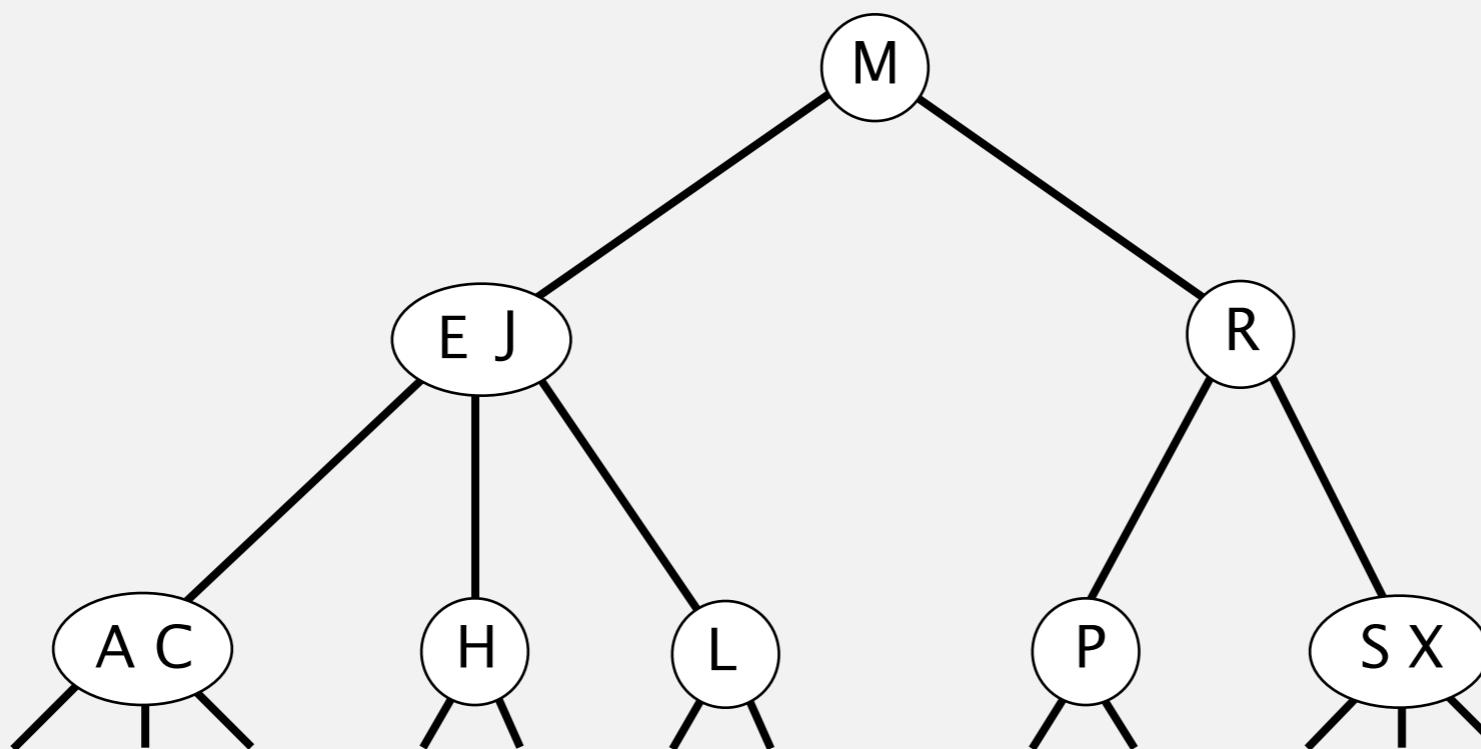
---

### Search.

- Compare search key against keys in node.
- Find interval containing search key.
- Follow associated link (recursively).



search for H



## 2-3 tree demo: search

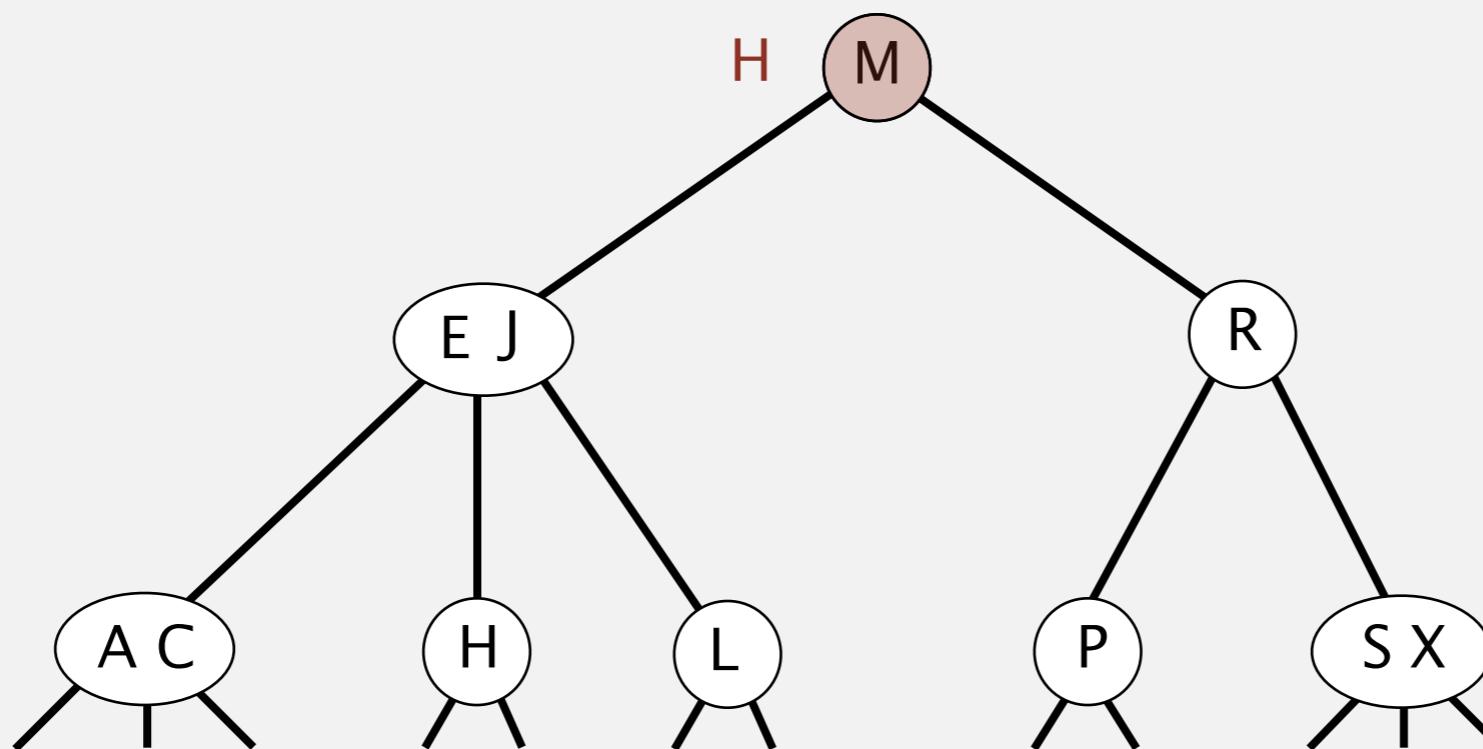
---

### Search.

- Compare search key against keys in node.
- Find interval containing search key.
- Follow associated link (recursively).

search for H

H is less than M  
(go left)



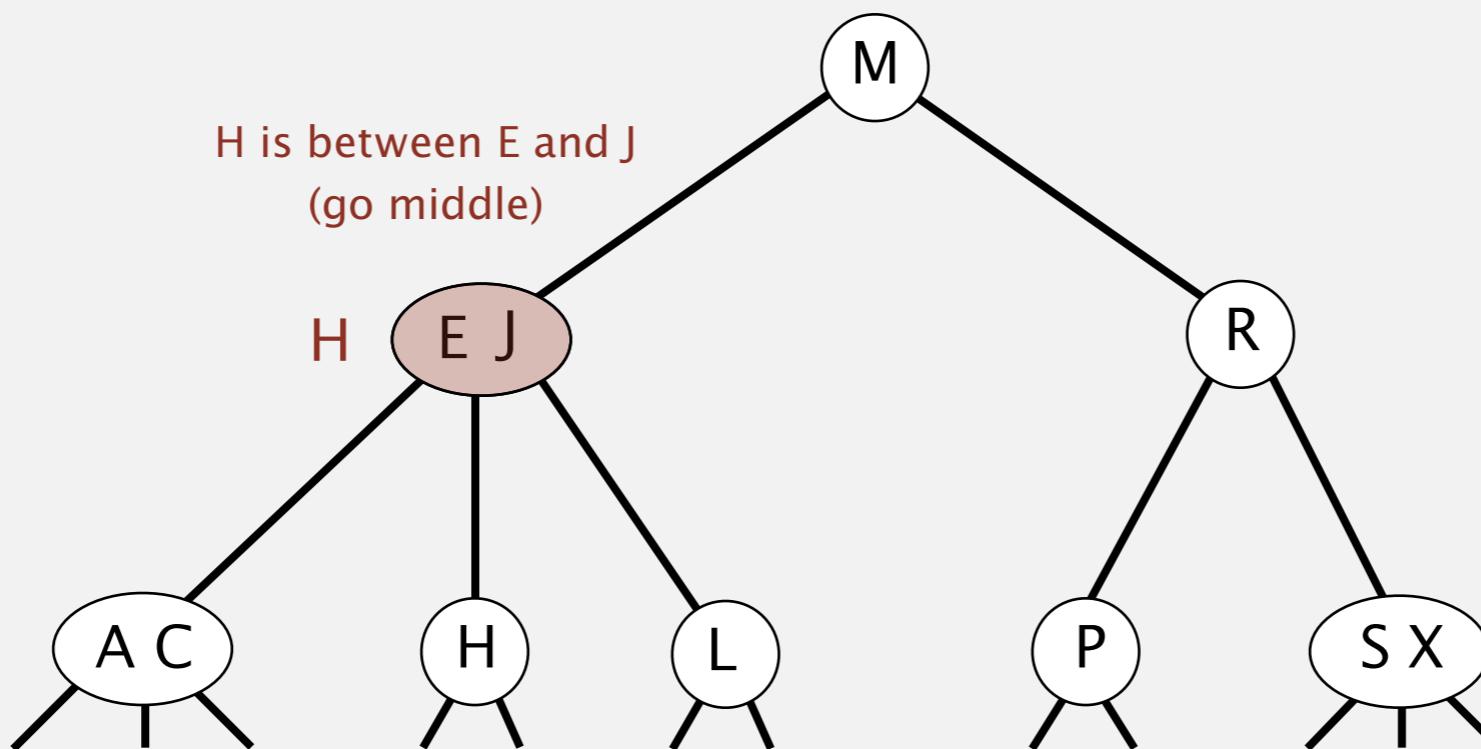
## 2-3 tree demo: search

---

### Search.

- Compare search key against keys in node.
- Find interval containing search key.
- Follow associated link (recursively).

search for H



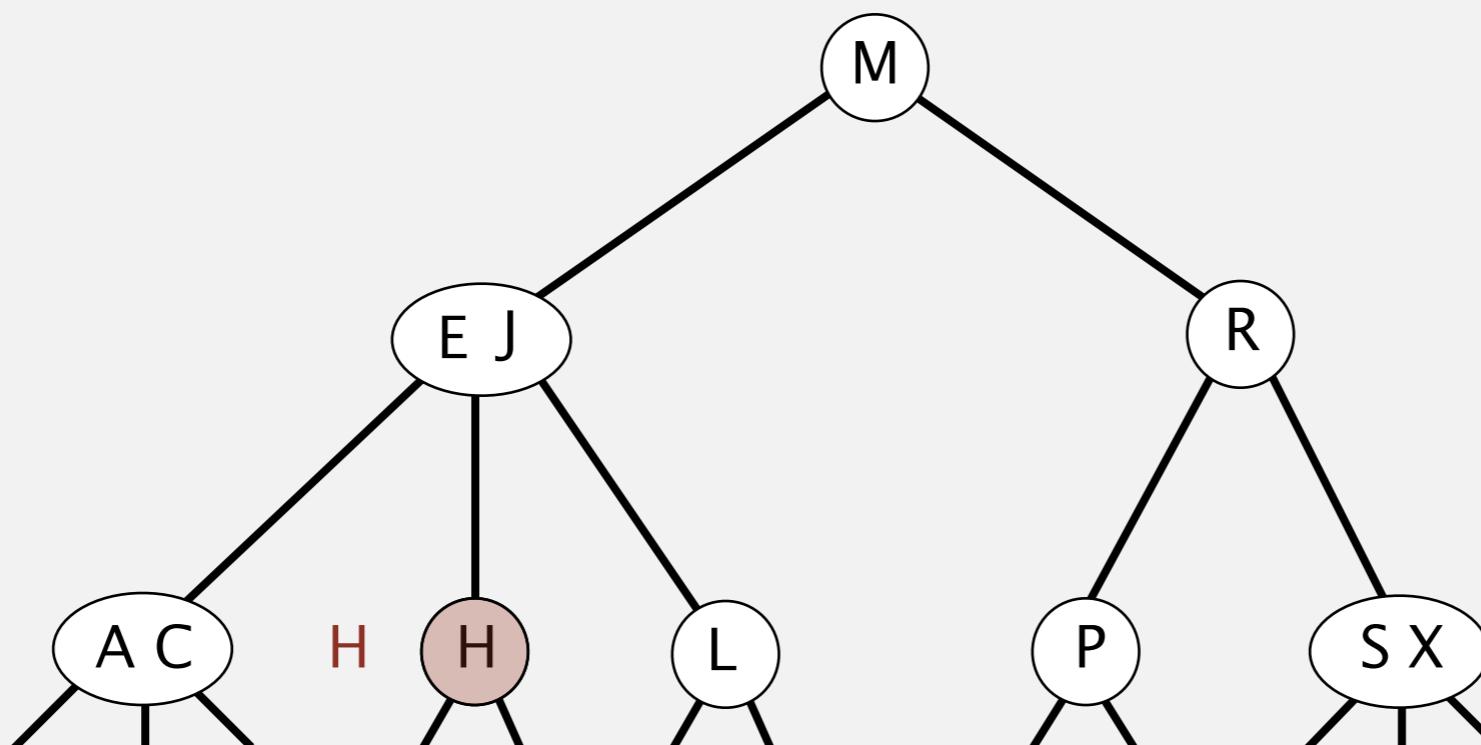
## 2-3 tree demo: search

---

### Search.

- Compare search key against keys in node.
- Find interval containing search key.
- Follow associated link (recursively).

search for H



found H  
(search hit)

## 2-3 tree demo: search

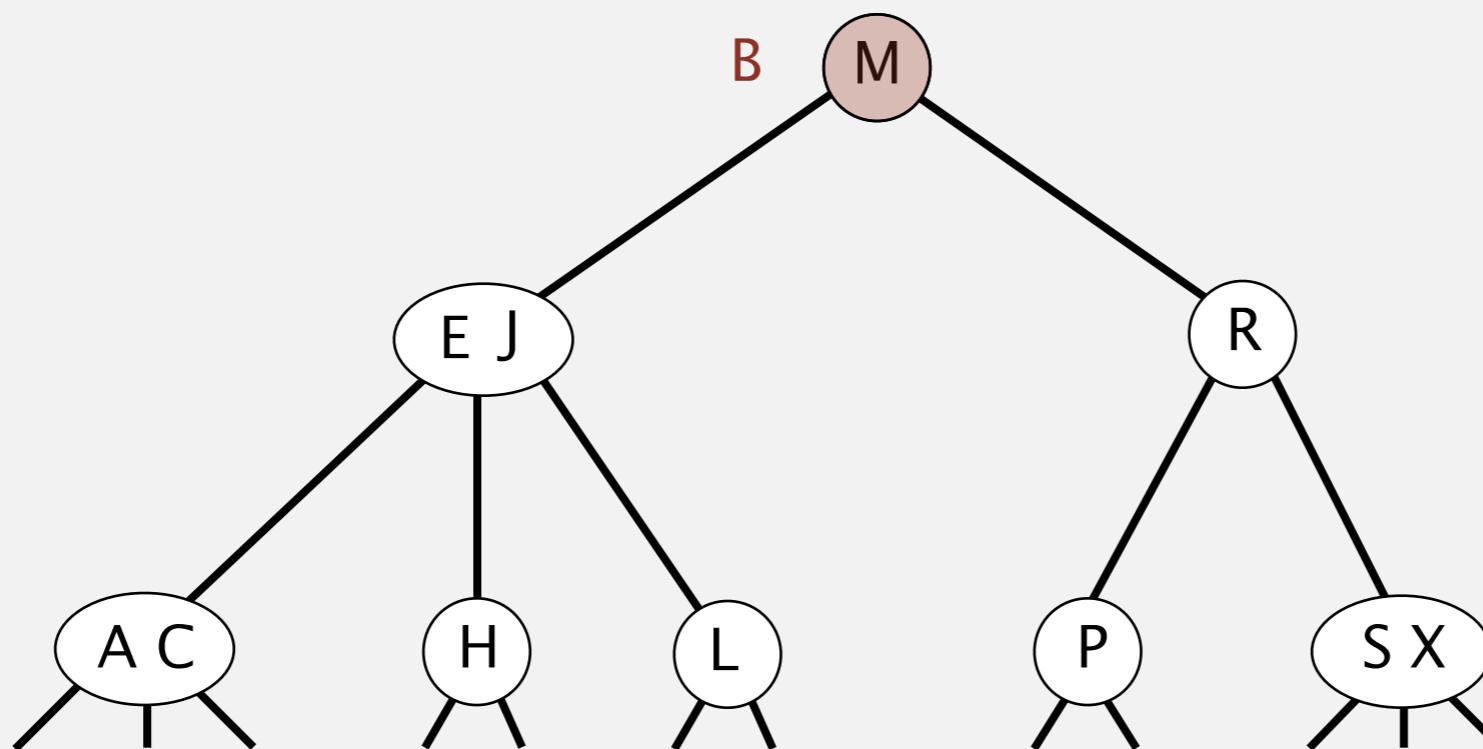
---

### Search.

- Compare search key against keys in node.
- Find interval containing search key.
- Follow associated link (recursively).

search for B

B is less than M  
(go left)



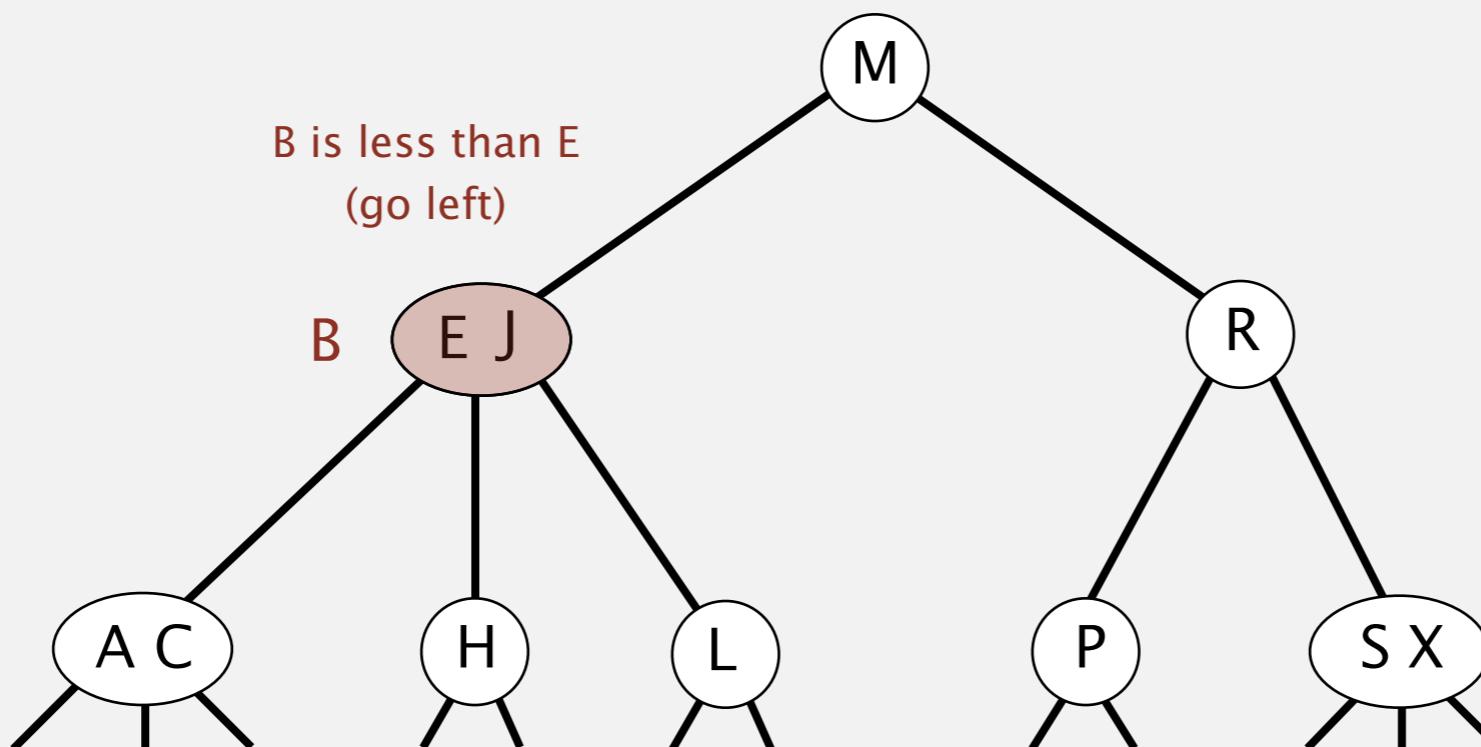
## 2-3 tree demo: search

---

### Search.

- Compare search key against keys in node.
- Find interval containing search key.
- Follow associated link (recursively).

search for B



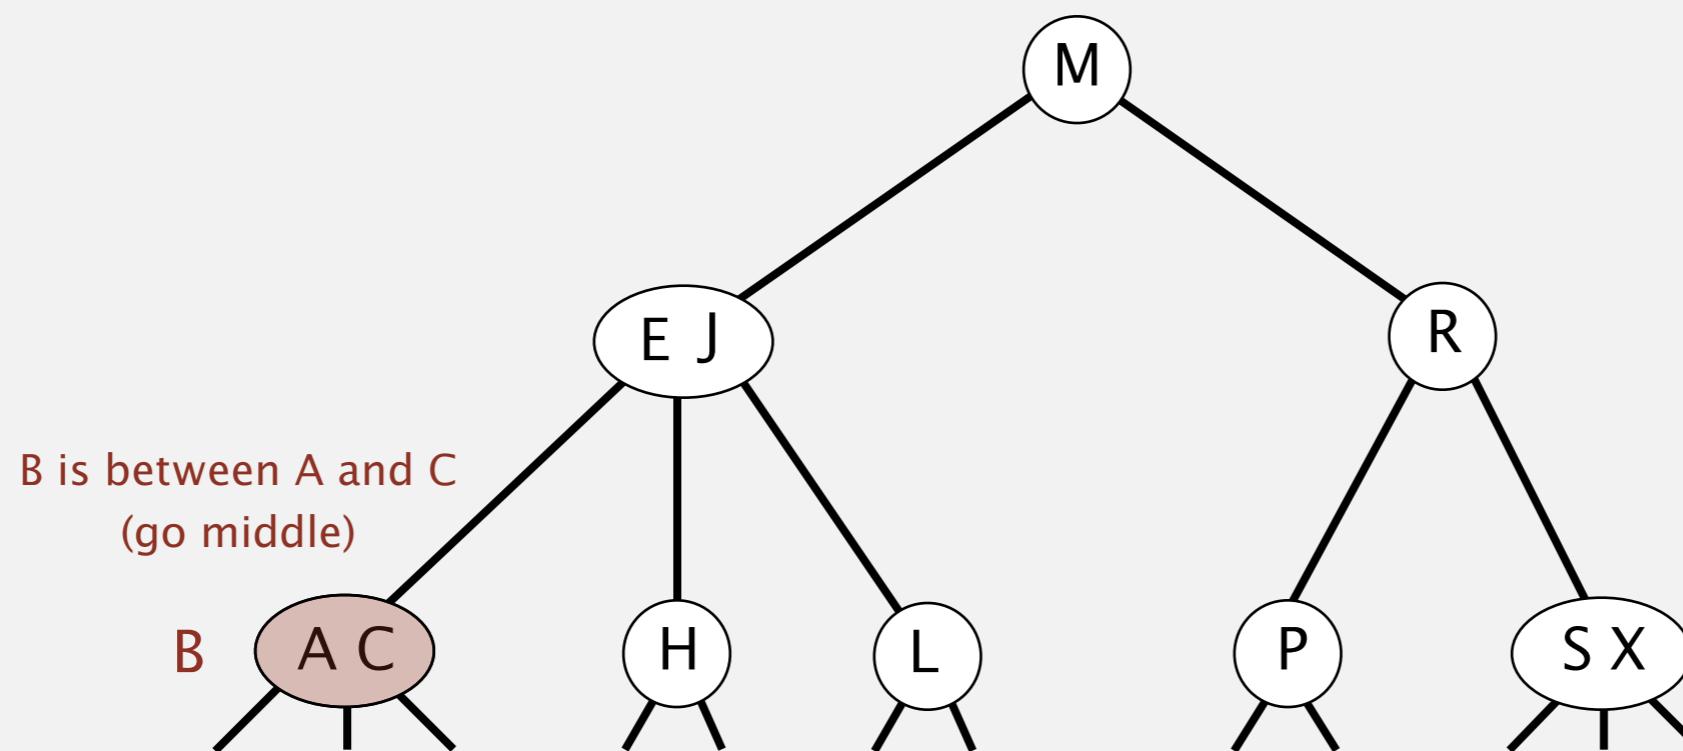
## 2-3 tree demo: search

---

### Search.

- Compare search key against keys in node.
- Find interval containing search key.
- Follow associated link (recursively).

search for B



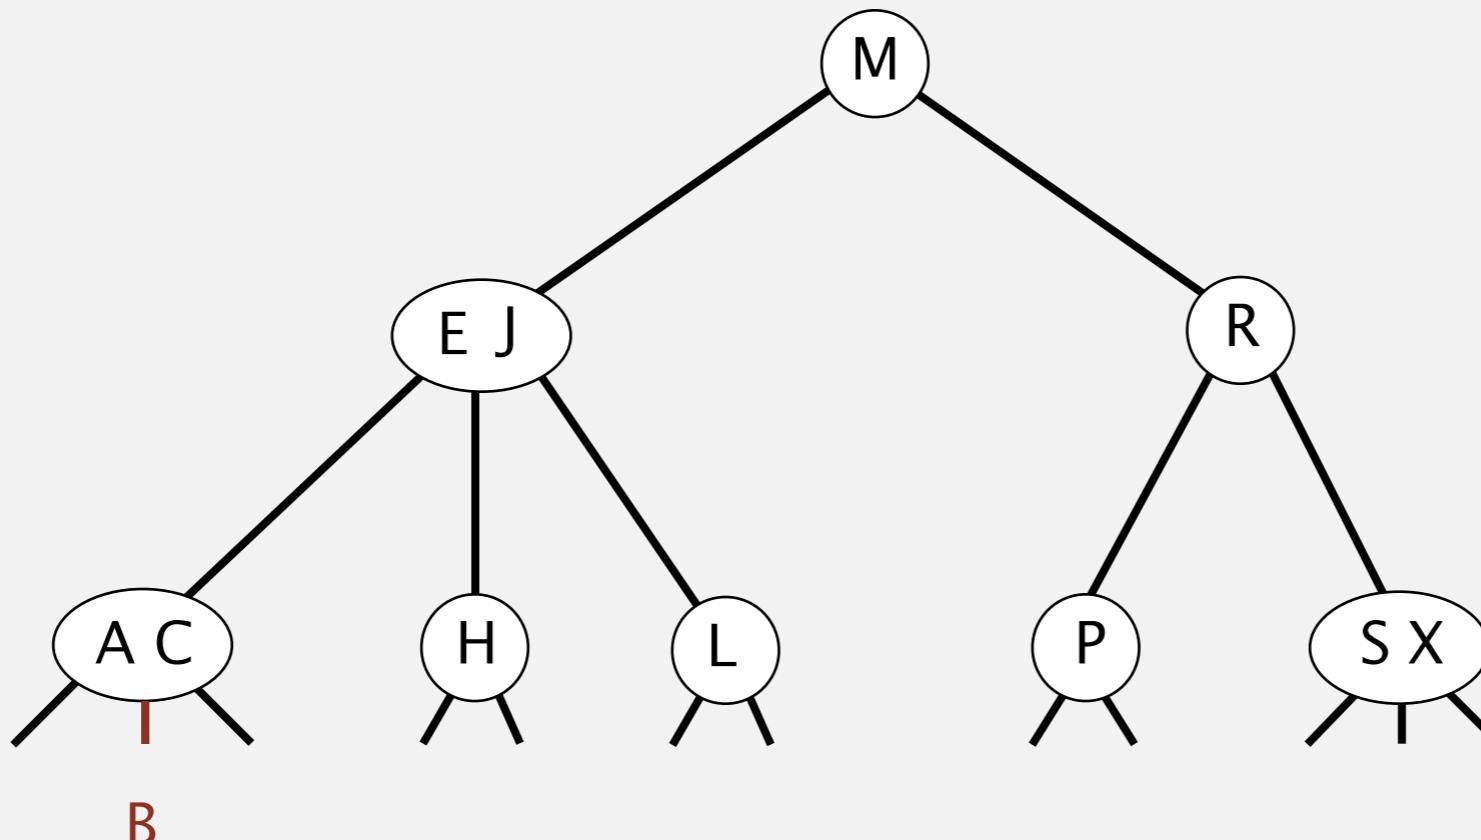
## 2-3 tree demo: search

---

### Search.

- Compare search key against keys in node.
- Find interval containing search key.
- Follow associated link (recursively).

search for B



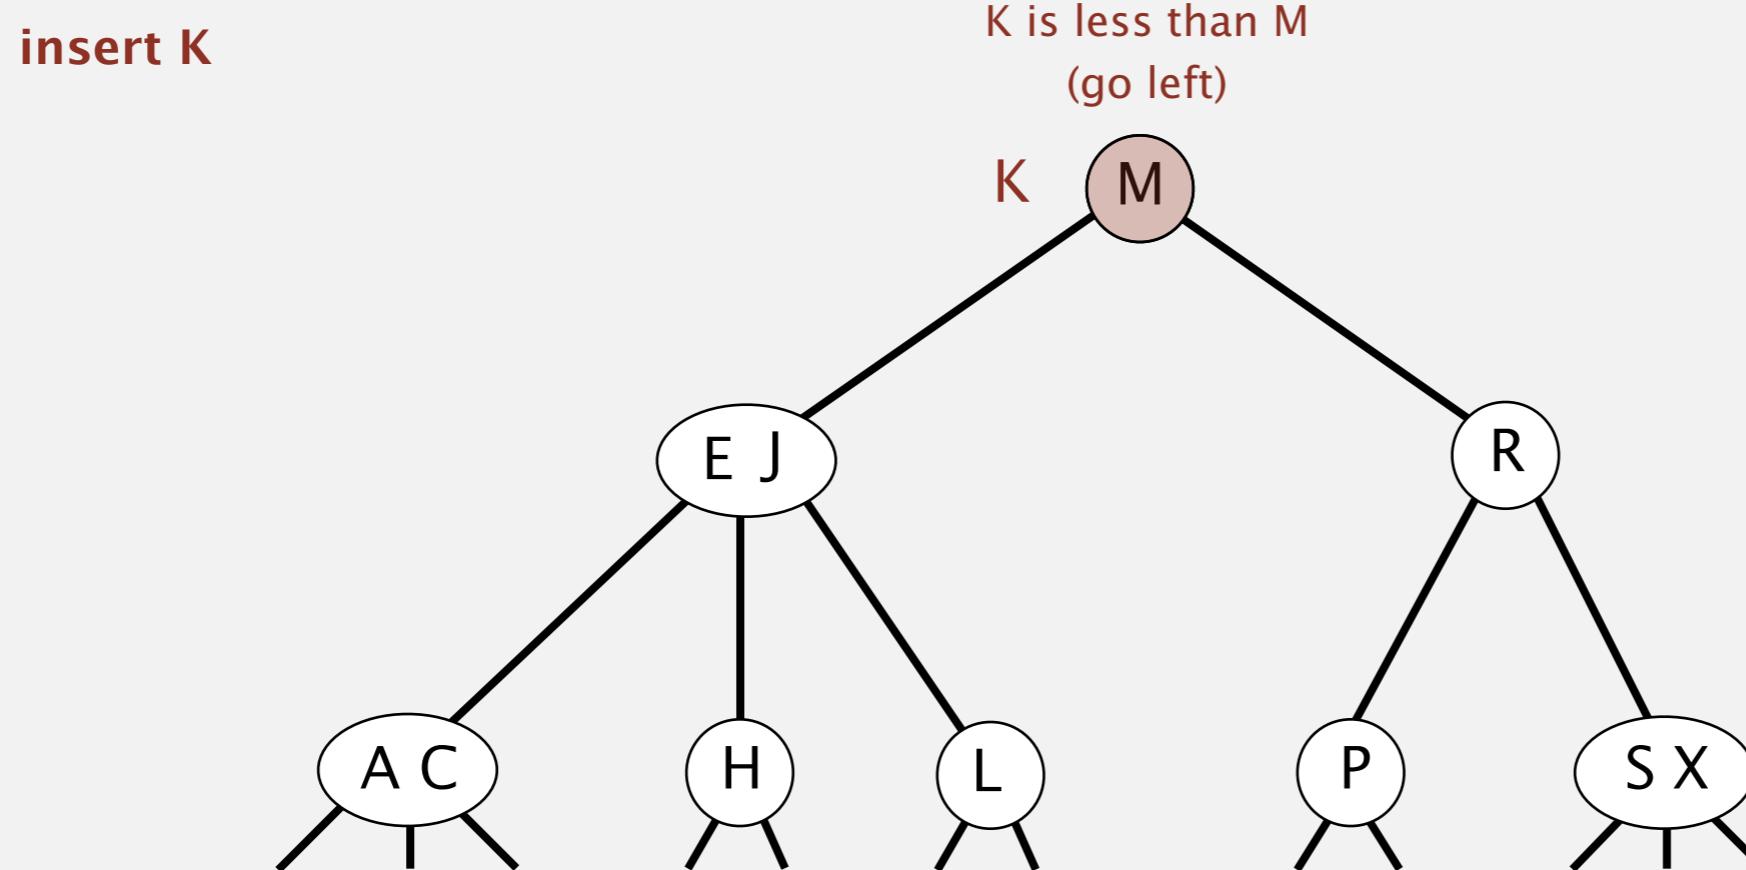
link is null  
(search miss)

## 2-3 tree demo: insertion

---

Insert into a 2-node at bottom.

- Search for key, as usual.
- Replace 2-node with 3-node.



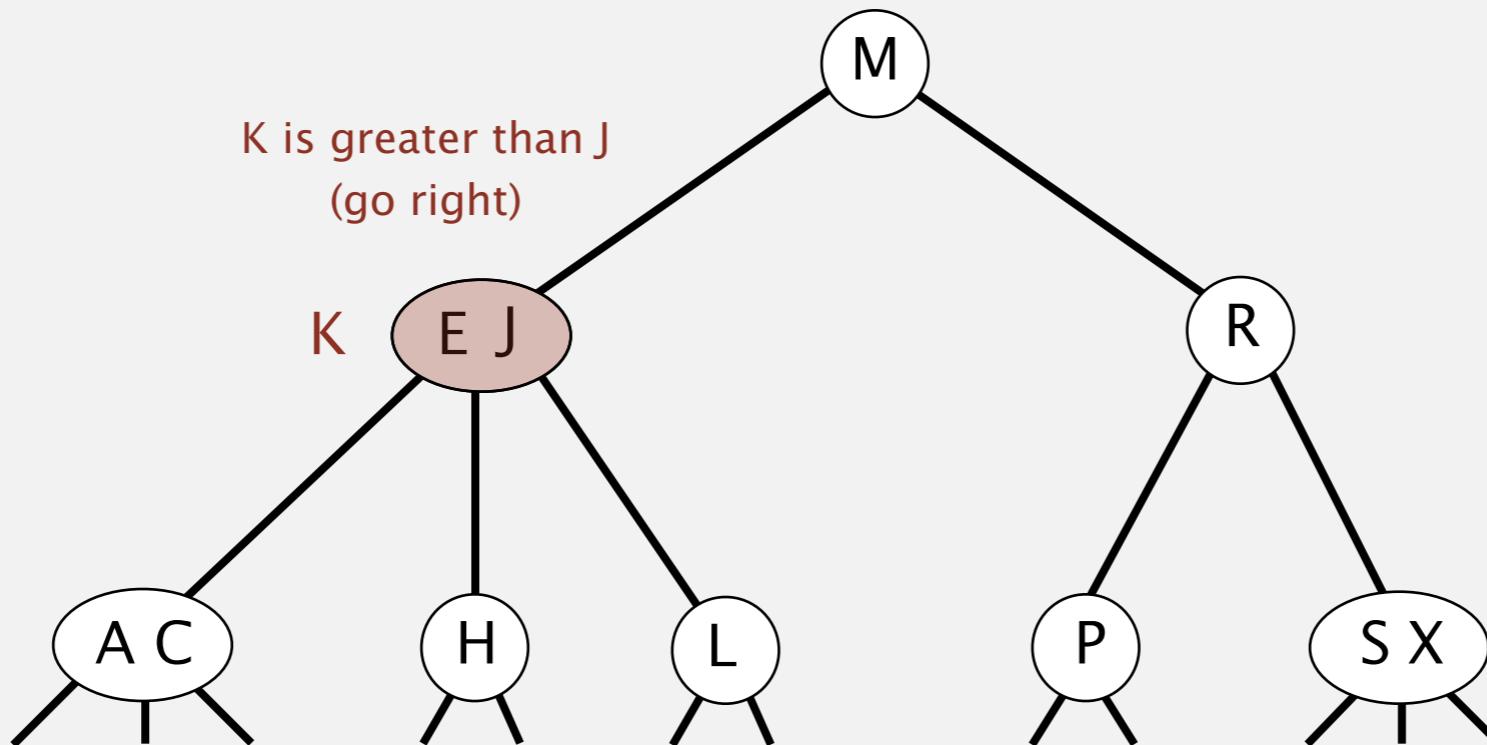
## 2-3 tree demo: insertion

---

Insert into a 2-node at bottom.

- Search for key, as usual.
- Replace 2-node with 3-node.

insert K



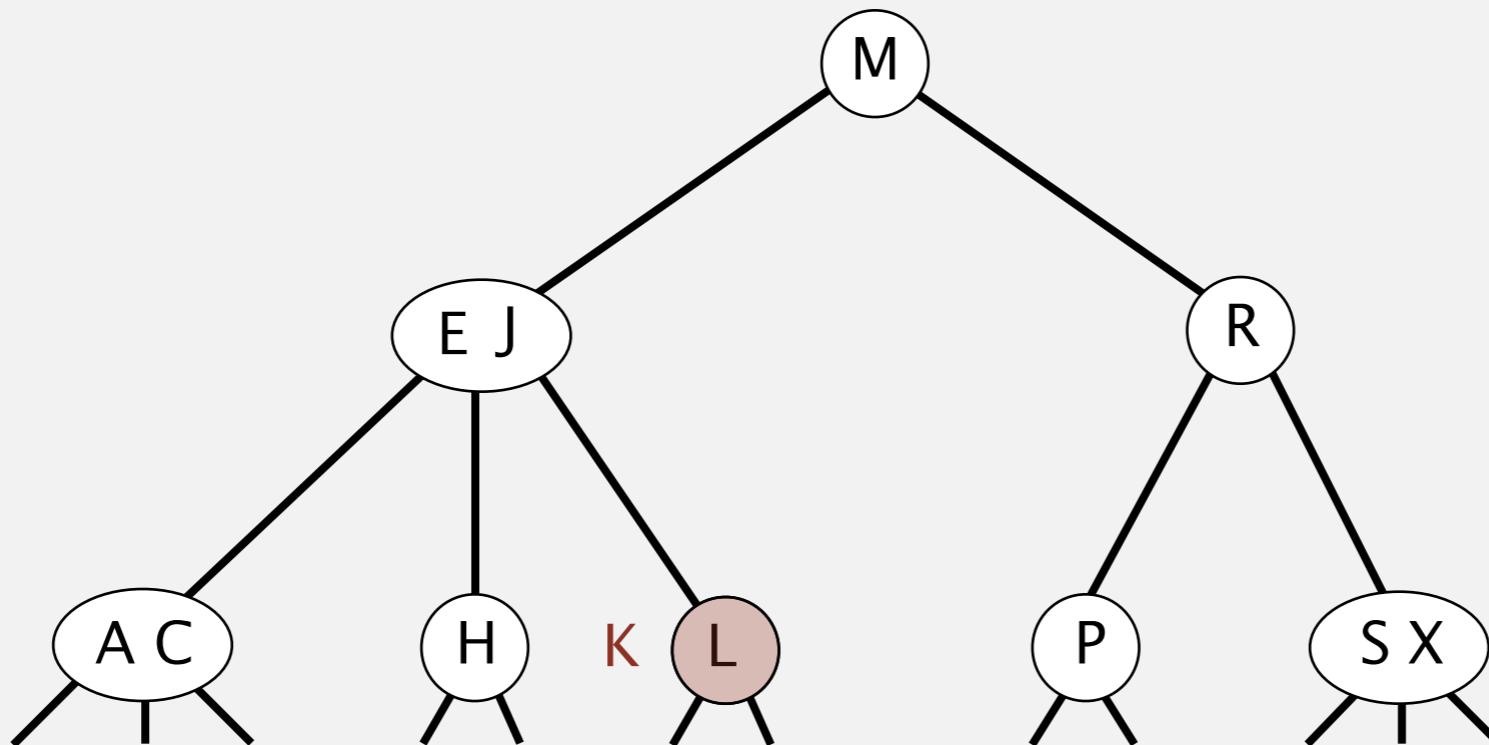
## 2-3 tree demo: insertion

---

Insert into a 2-node at bottom.

- Search for key, as usual.
- Replace 2-node with 3-node.

insert K



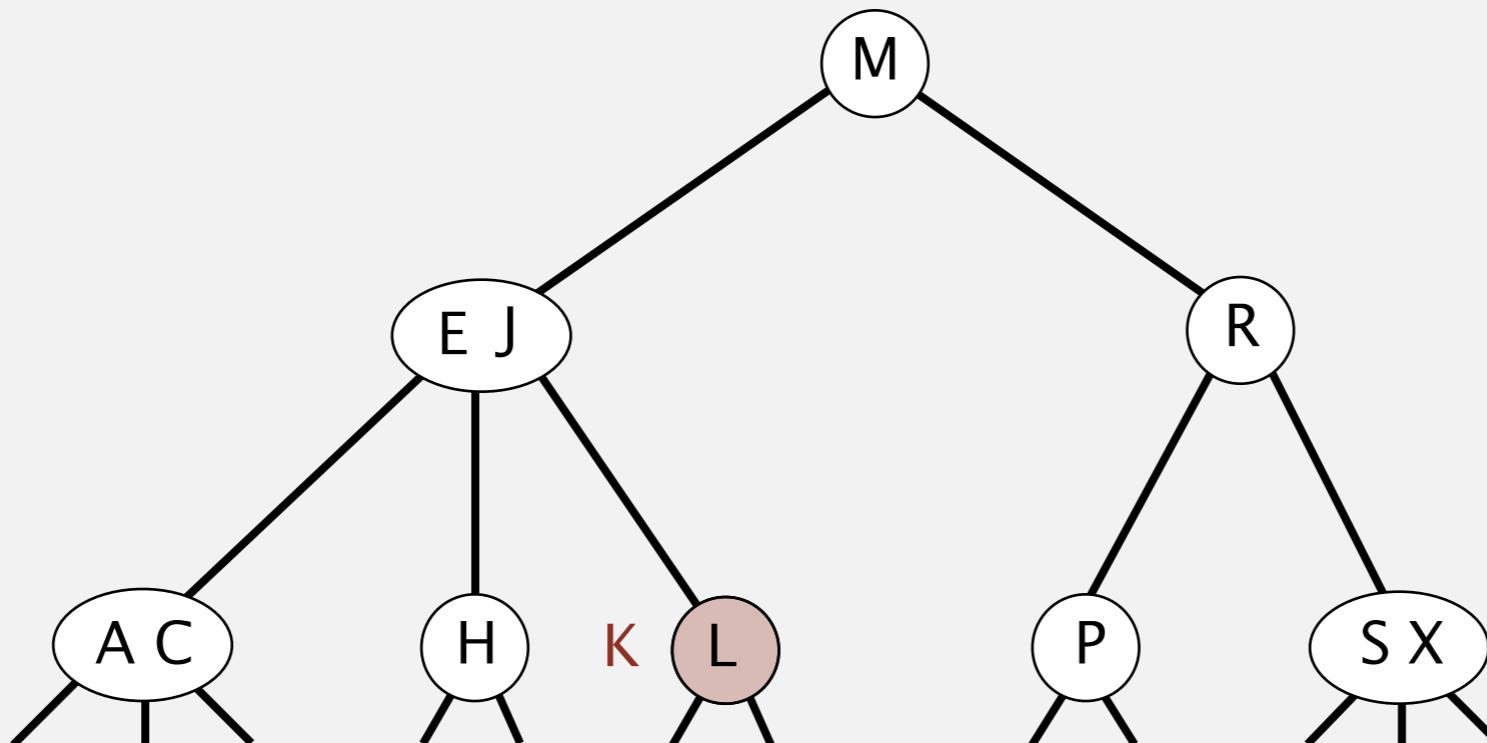
## 2-3 tree demo: insertion

---

Insert into a 2-node at bottom.

- Search for key, as usual.
- Replace 2-node with 3-node.

insert K



replace 2-node with  
3-node containing K

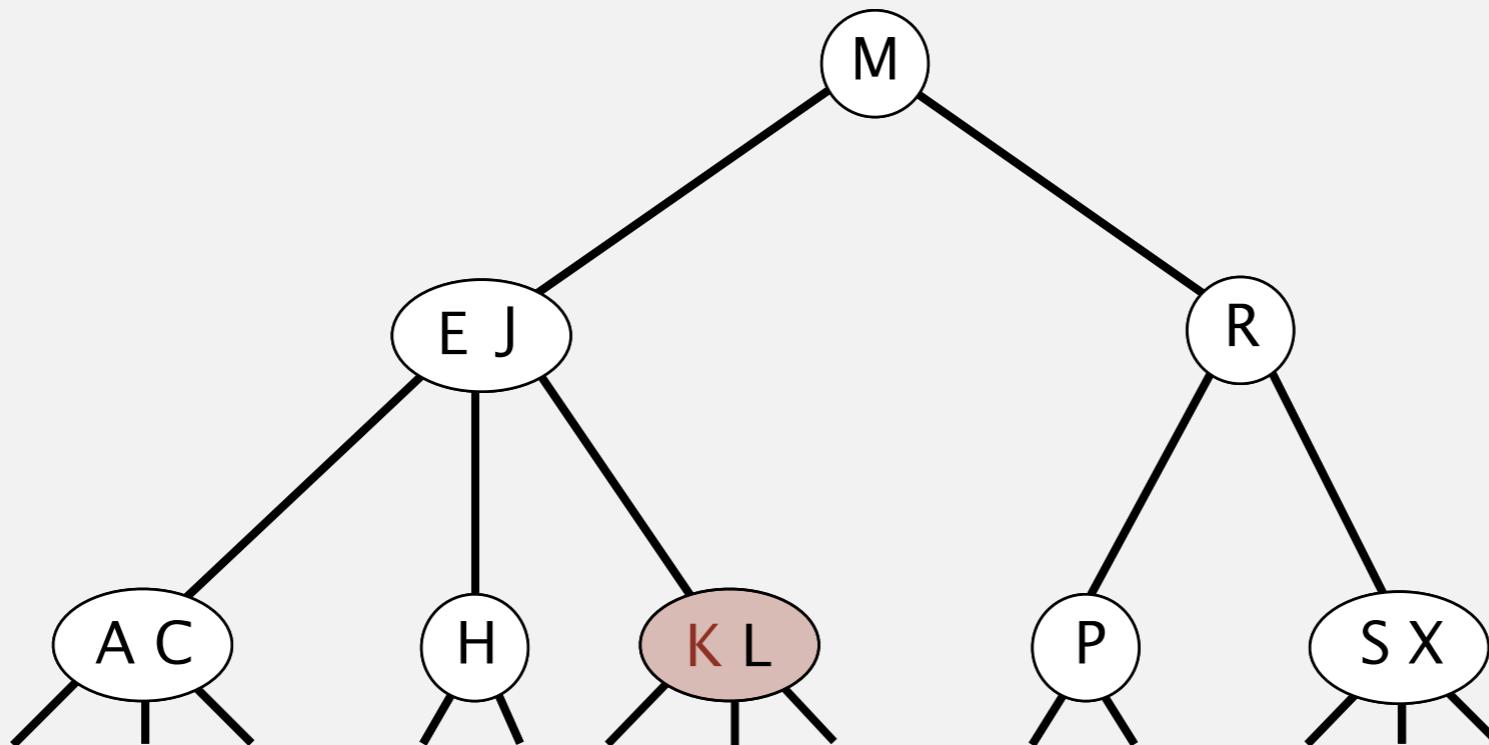
## 2-3 tree demo: insertion

---

Insert into a 2-node at bottom.

- Search for key, as usual.
- Replace 2-node with 3-node.

insert K



## 2-3 tree demo: insertion

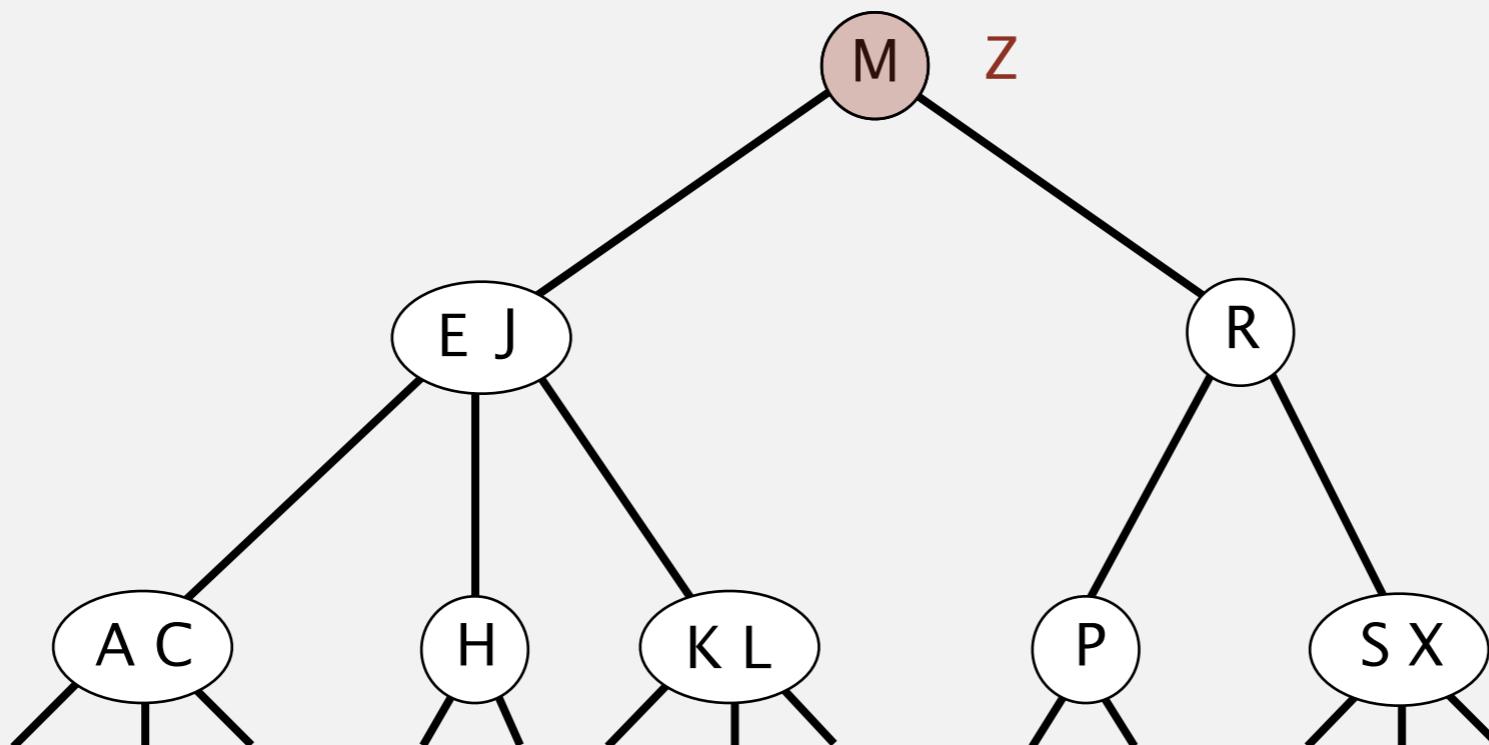
---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.

insert Z

Z is greater than M  
(go right)



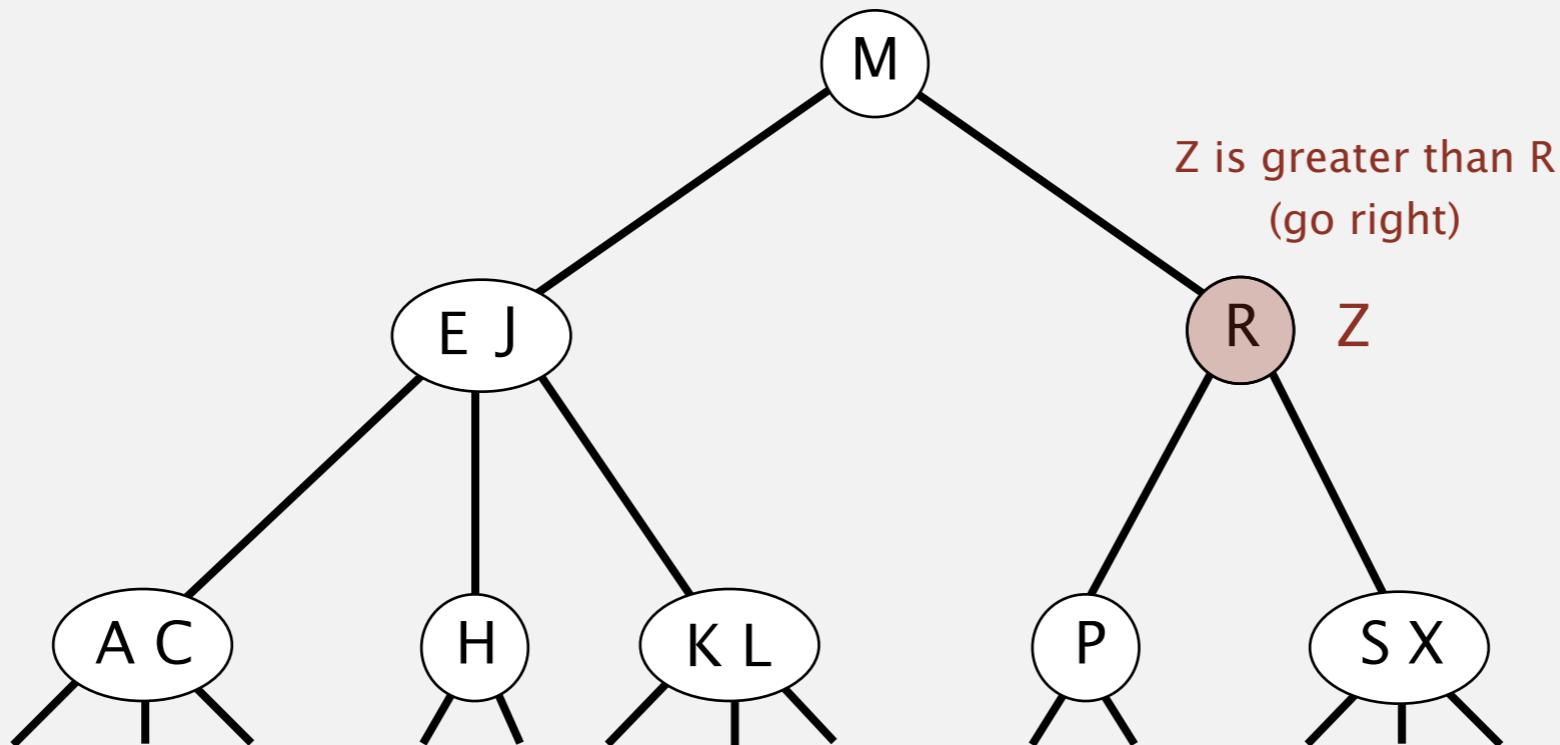
## 2-3 tree demo: insertion

---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.

insert Z



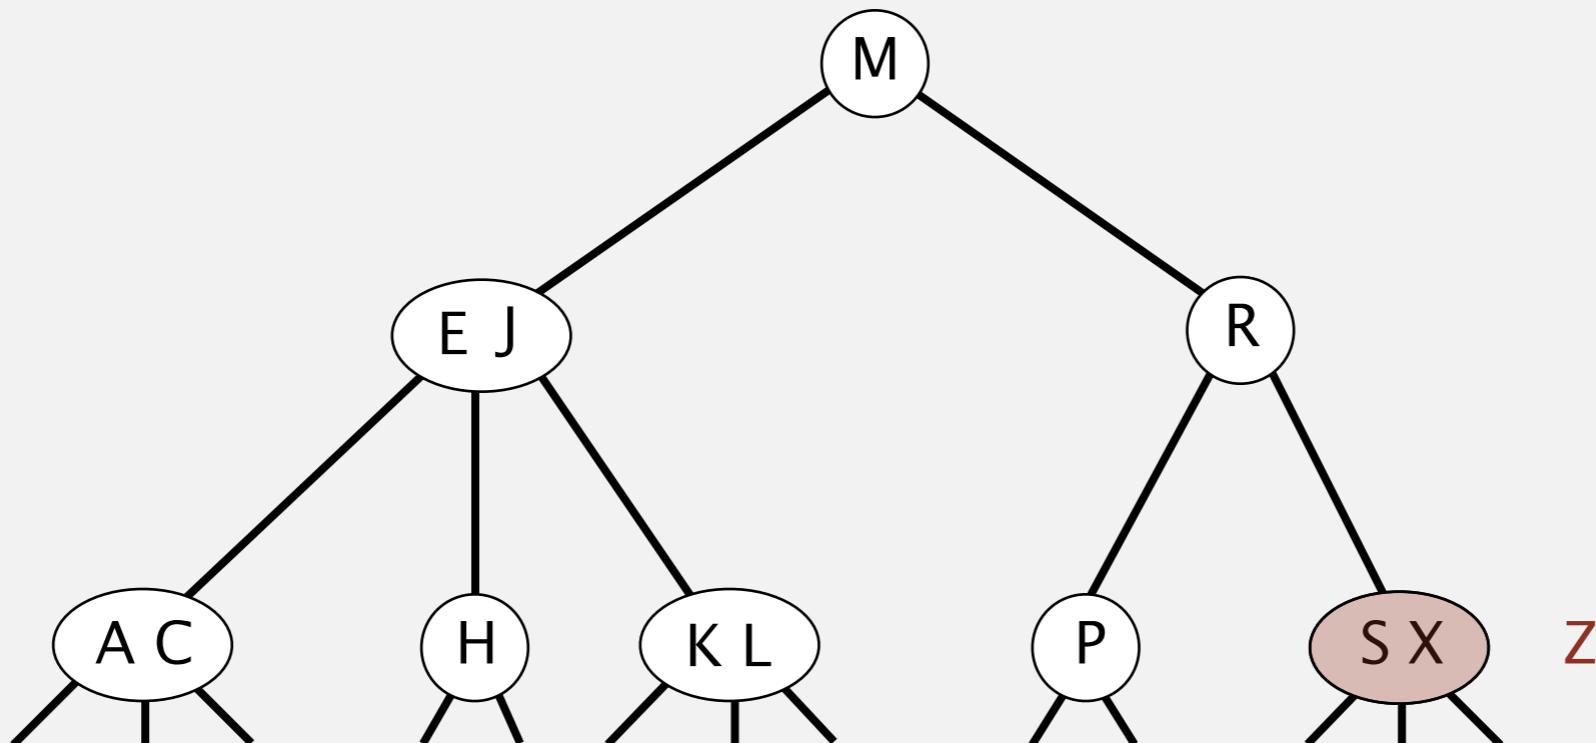
## 2-3 tree demo: insertion

---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.

insert Z



search ends here

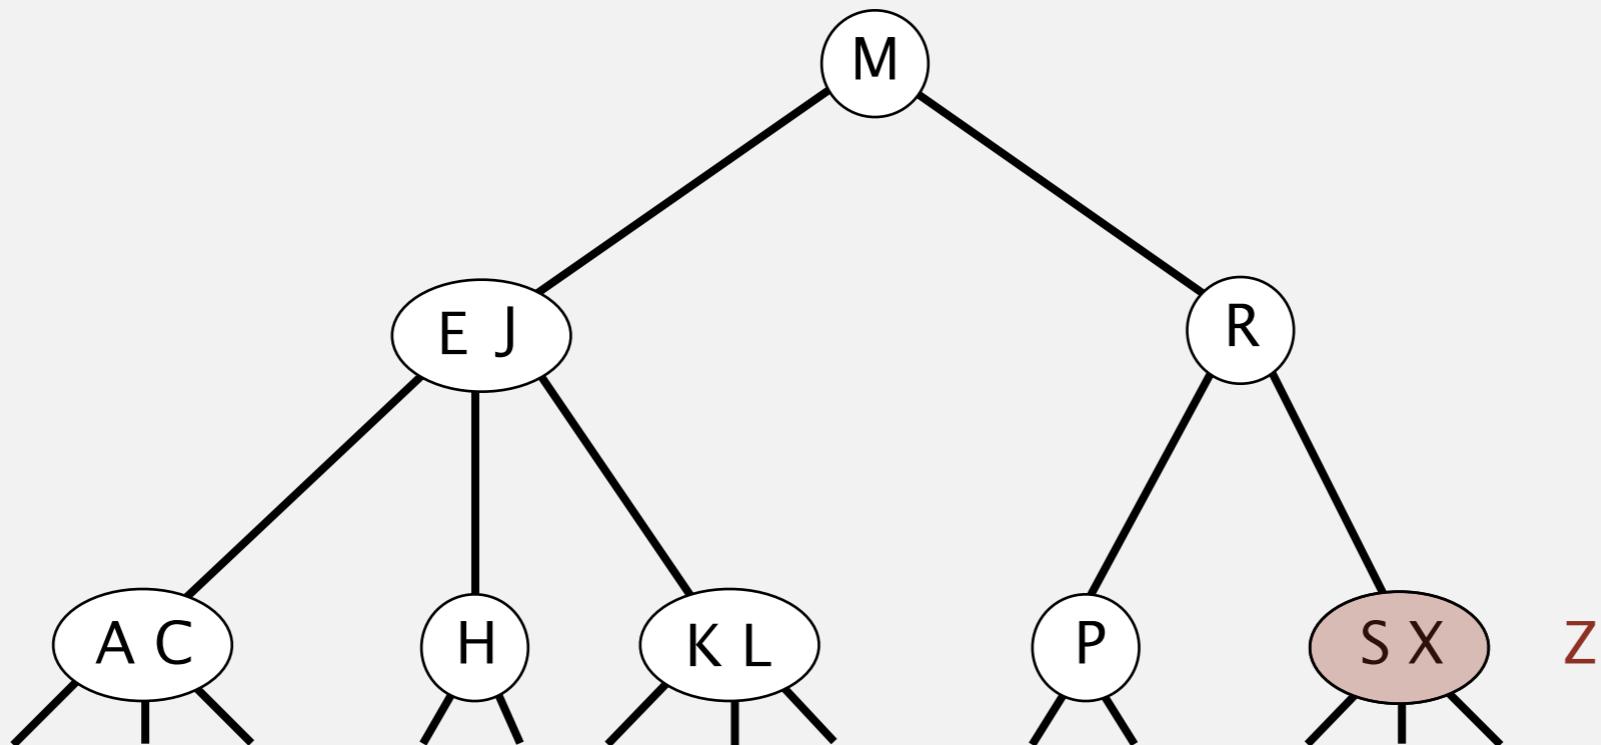
## 2-3 tree demo: insertion

---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.

insert Z



replace 3-node with  
temporary 4-node containing Z

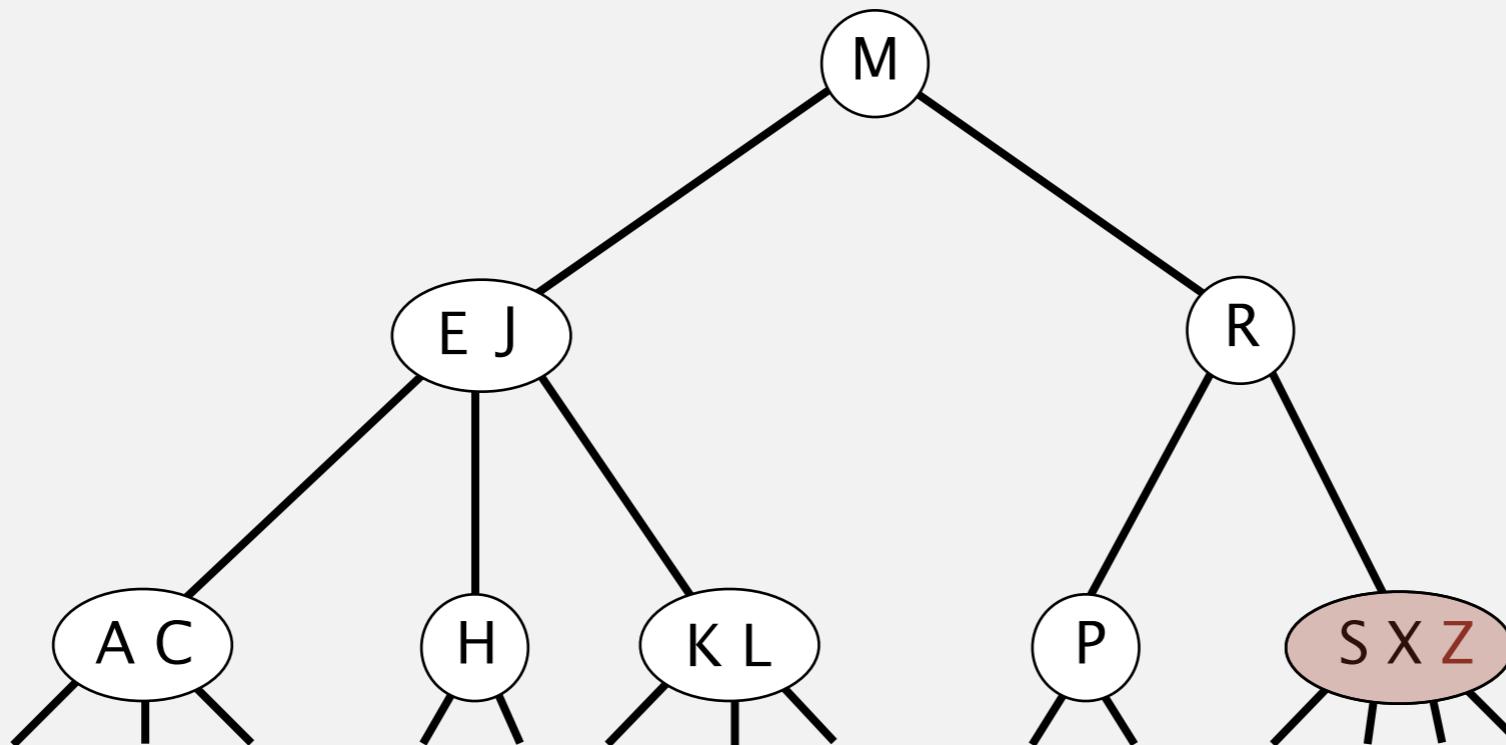
## 2-3 tree demo: insertion

---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.

insert Z



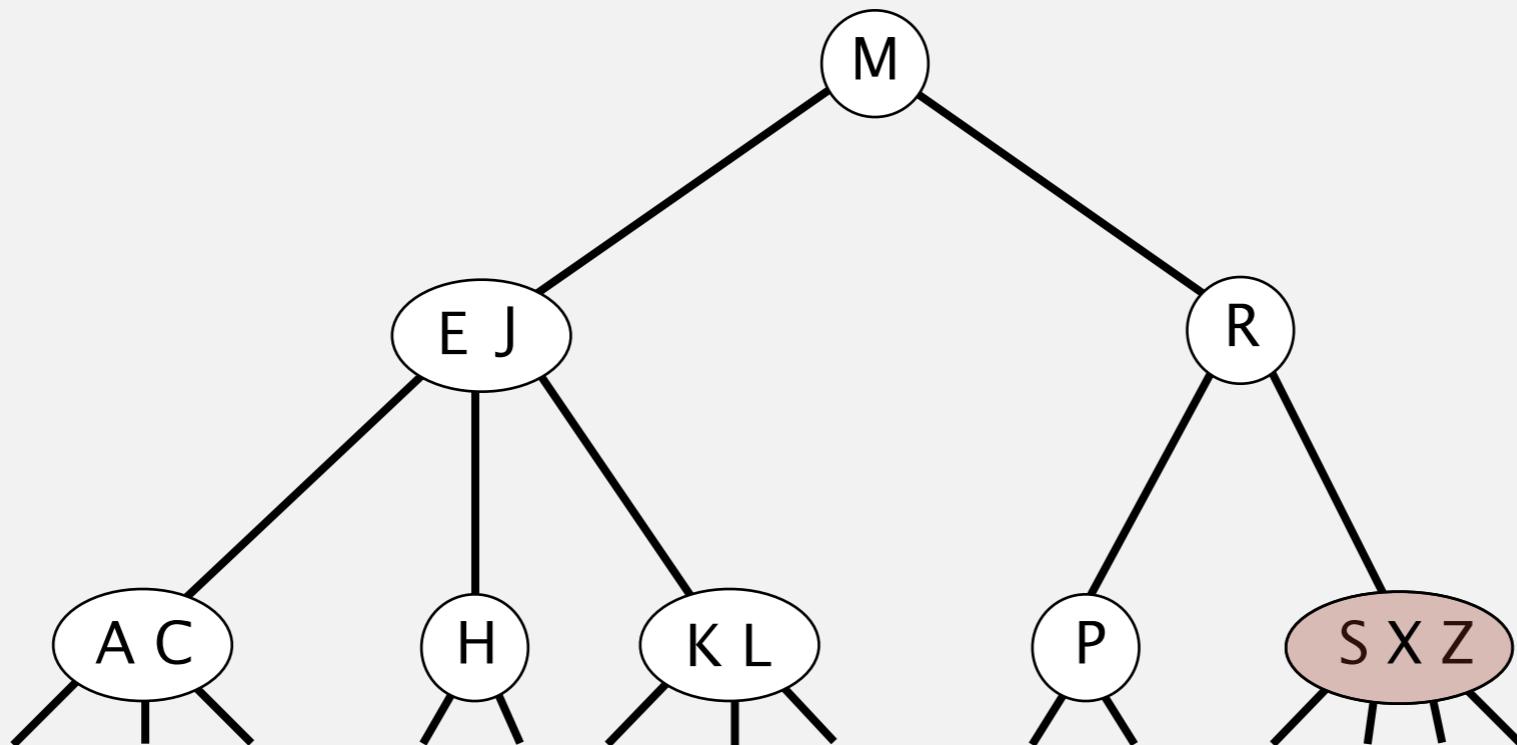
## 2-3 tree demo: insertion

---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.

insert Z



split 4-node into two 2-nodes  
(pass middle key to parent)

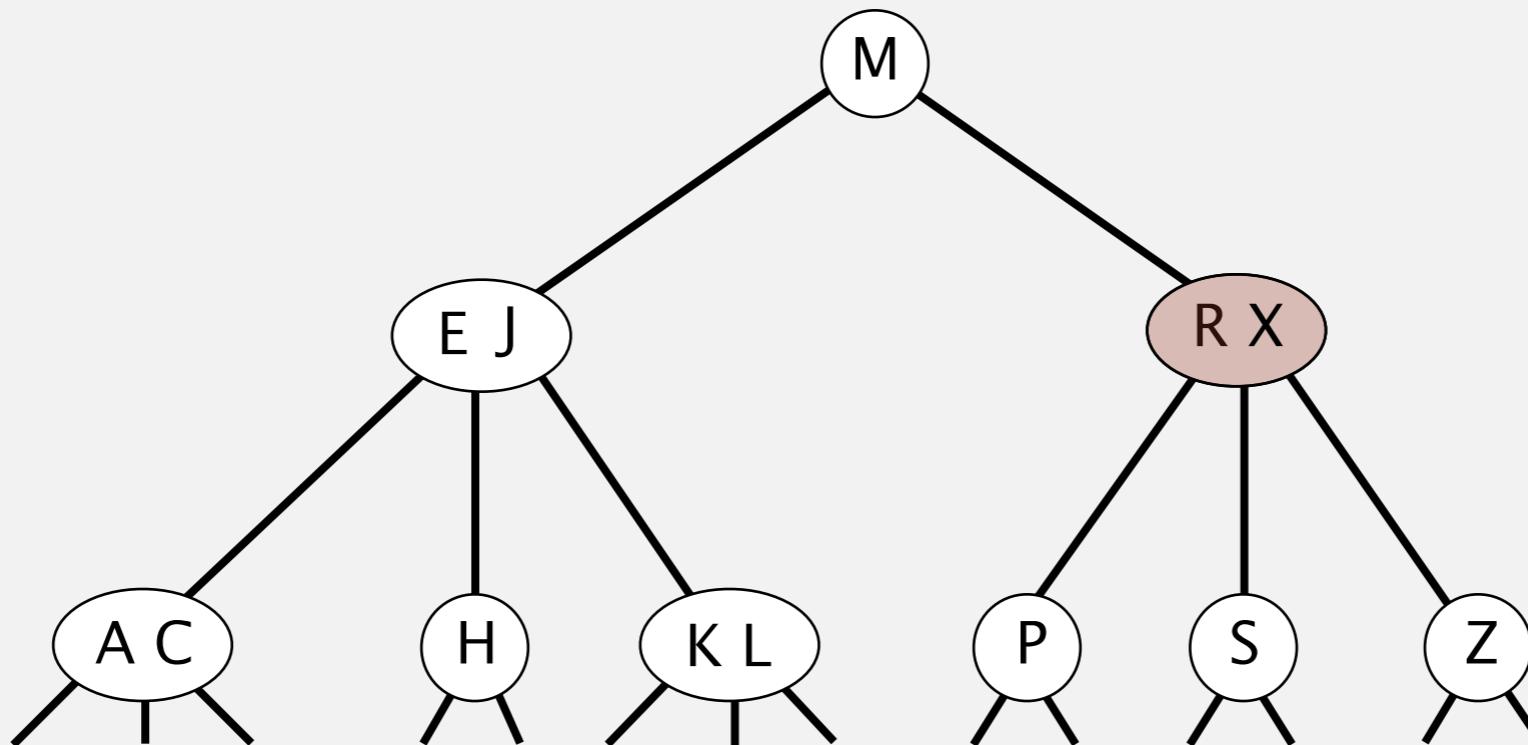
## 2-3 tree demo: insertion

---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.

insert Z



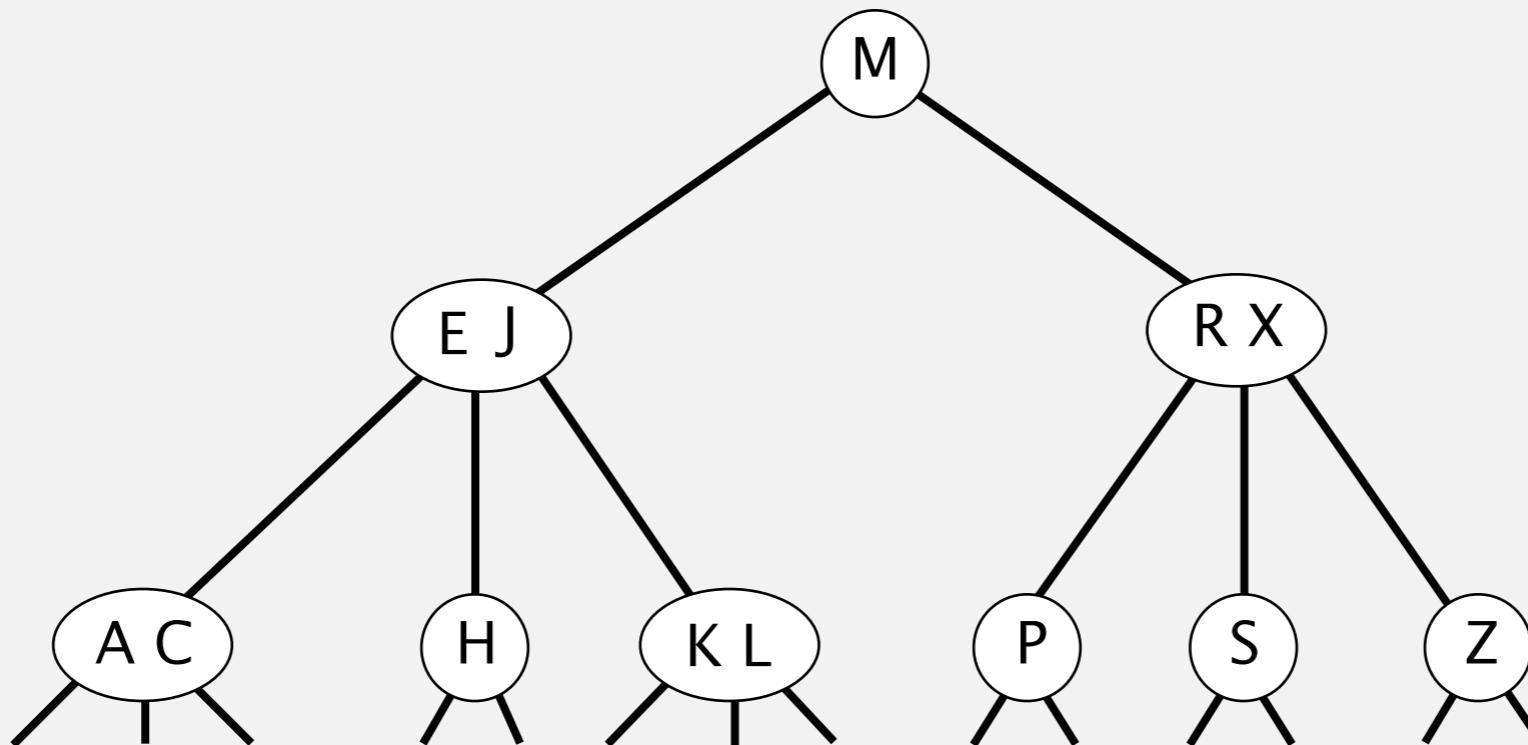
## 2-3 tree demo: insertion

---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.

insert Z



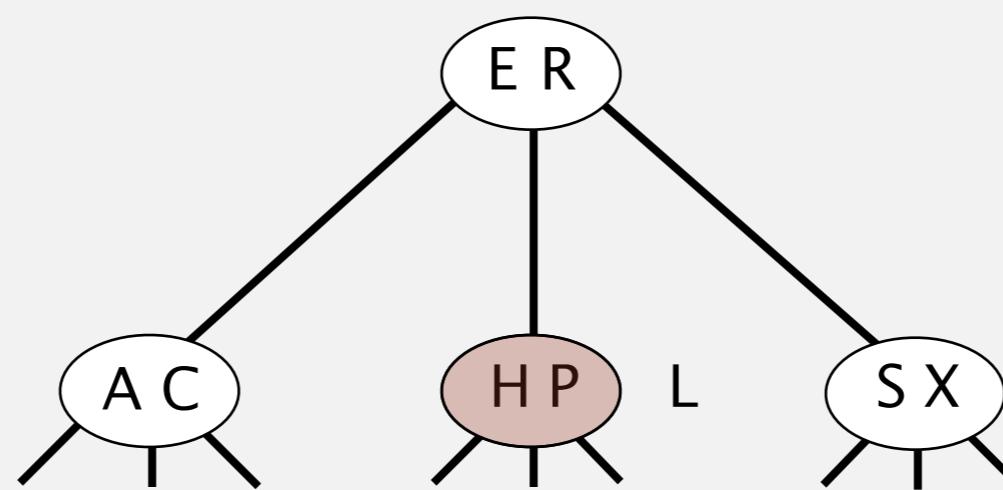
## 2-3 tree demo: insertion

---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.
- Repeat up the tree, as necessary.
- If you reach the root and it's a 4-node, split it into three 2-nodes.

insert L



convert 3-node into 4-node

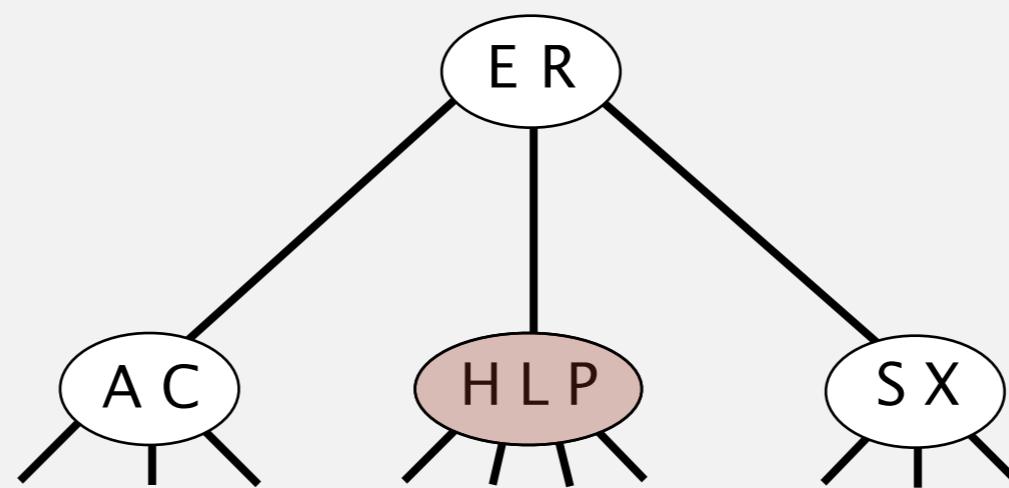
## 2-3 tree demo: insertion

---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.
- Repeat up the tree, as necessary.
- If you reach the root and it's a 4-node, split it into three 2-nodes.

insert L



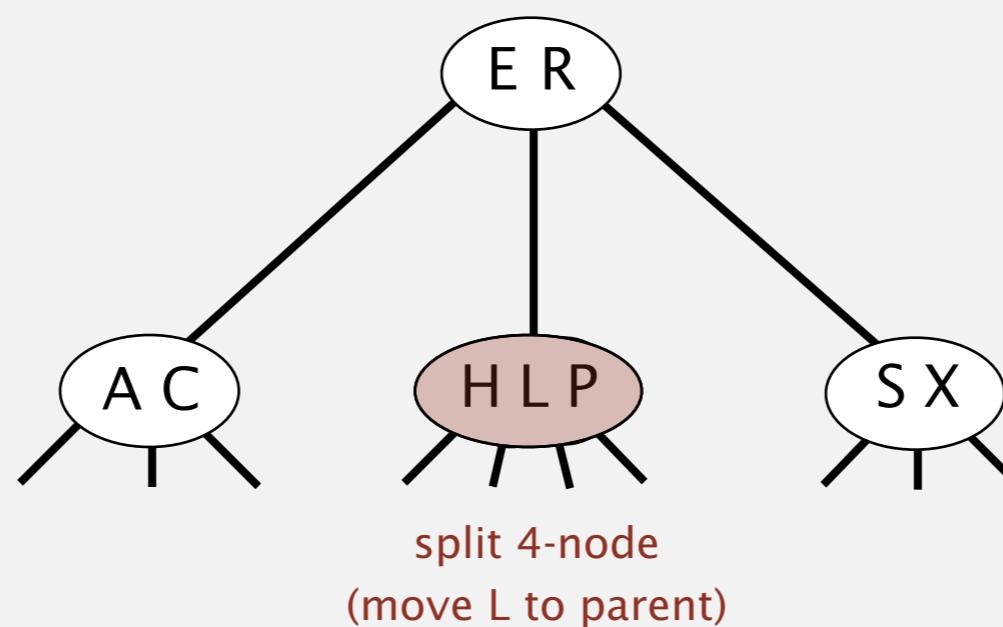
## 2-3 tree demo: insertion

---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.
- Repeat up the tree, as necessary.
- If you reach the root and it's a 4-node, split it into three 2-nodes.

insert L



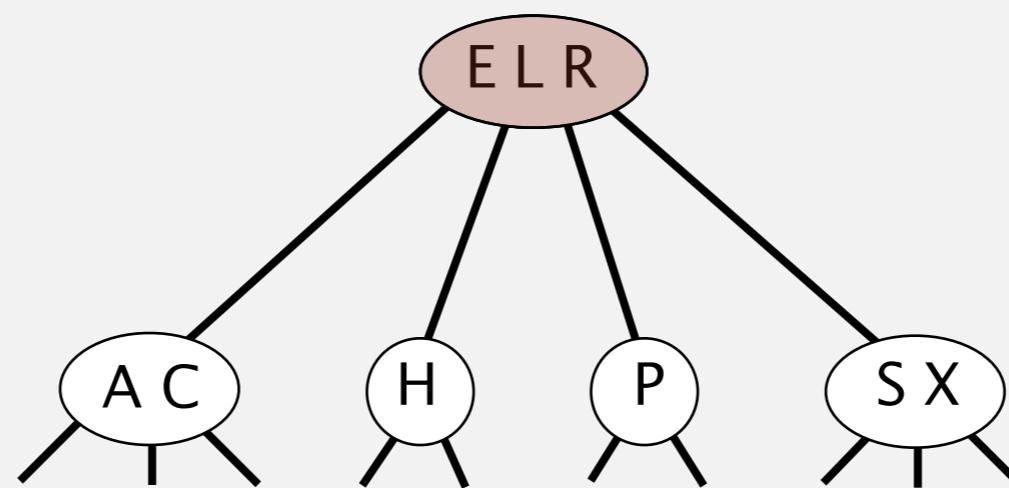
## 2-3 tree demo: insertion

---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.
- Repeat up the tree, as necessary.
- If you reach the root and it's a 4-node, split it into three 2-nodes.

insert L



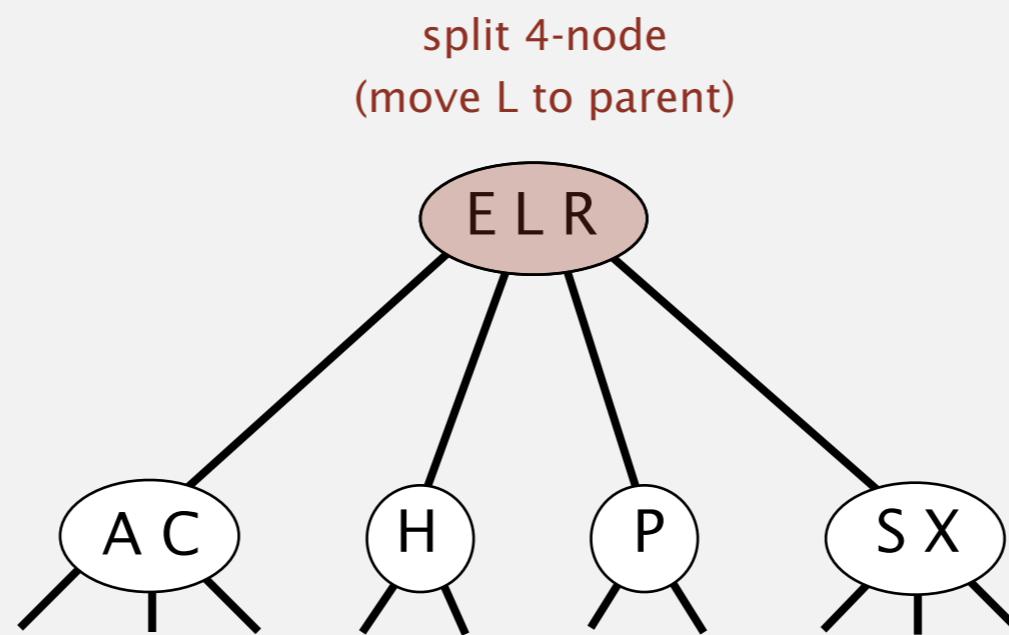
## 2-3 tree demo: insertion

---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.
- Repeat up the tree, as necessary.
- If you reach the root and it's a 4-node, split it into three 2-nodes.

insert L



## 2-3 tree demo: insertion

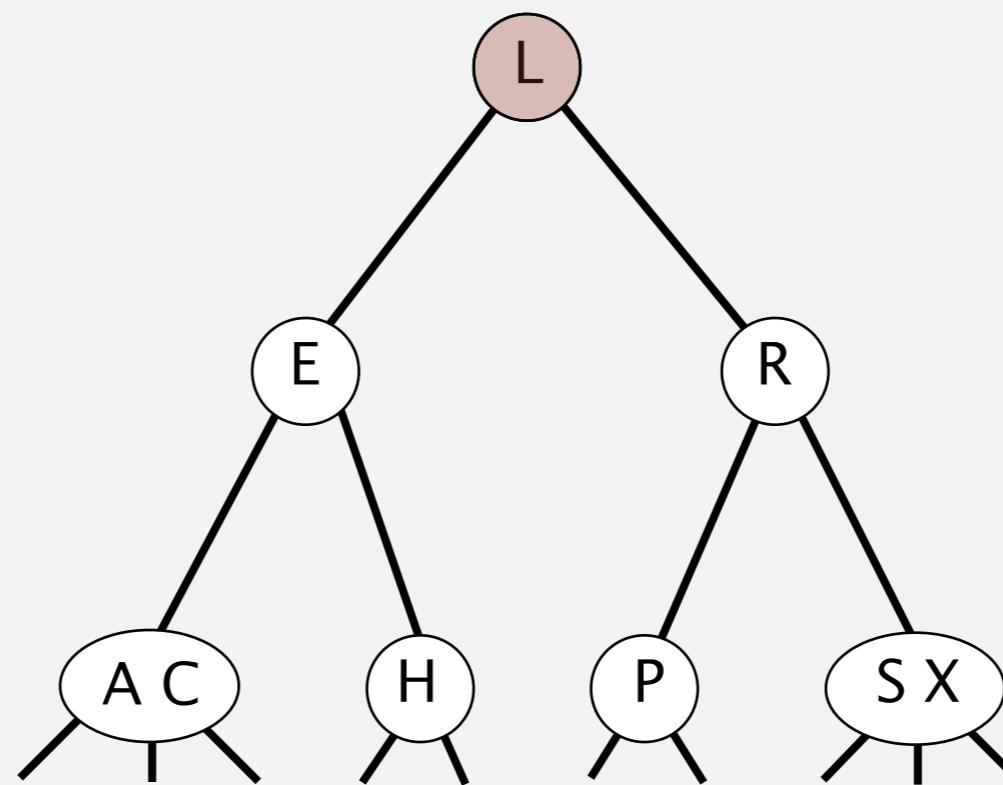
---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.
- Repeat up the tree, as necessary.
- If you reach the root and it's a 4-node, split it into three 2-nodes.

height of tree increases by 1

insert L



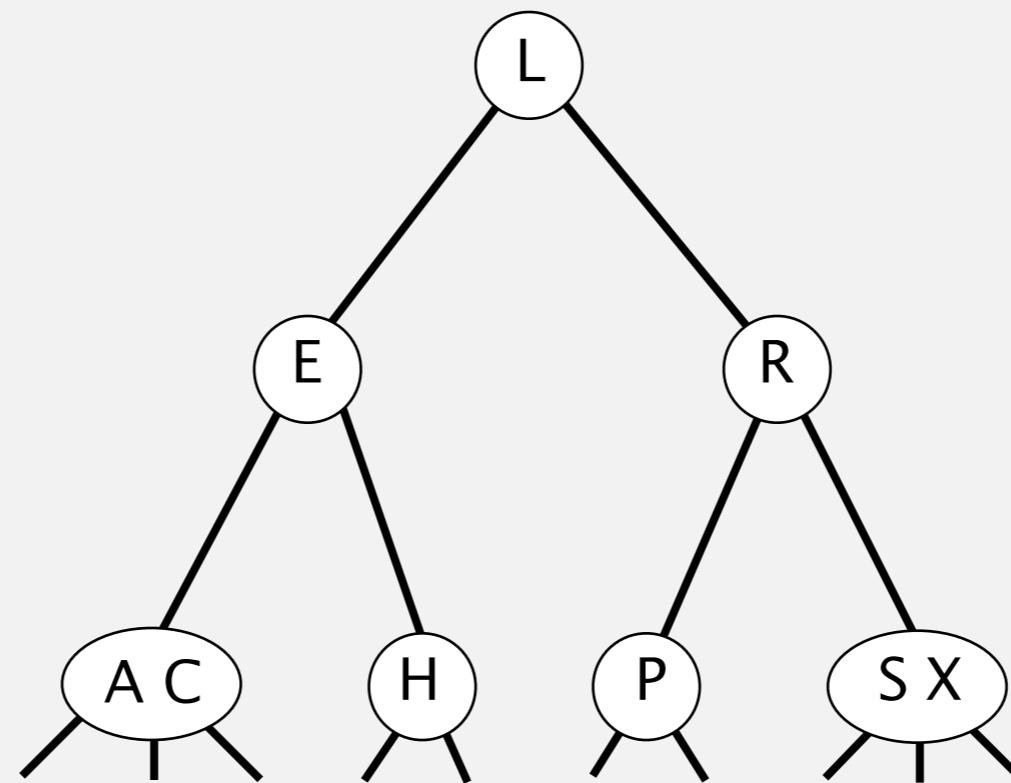
## 2-3 tree demo: insertion

---

Insert into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.
- Repeat up the tree, as necessary.
- If you reach the root and it's a 4-node, split it into three 2-nodes.

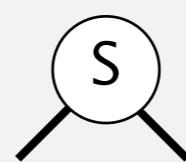
insert L



# 2-3 tree construction demo

---

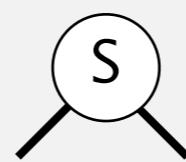
insert S



## 2-3 tree demo: construction

---

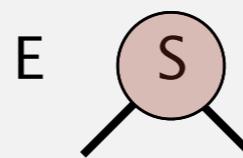
2-3 tree



## 2-3 tree demo: construction

---

insert E

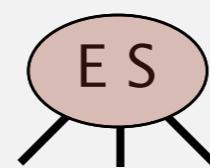


convert 2-node into 3-node

## 2-3 tree demo: construction

---

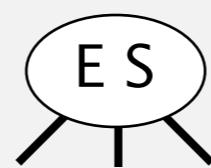
insert E



## 2-3 tree demo: construction

---

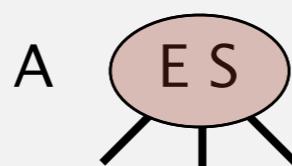
2-3 tree



## 2-3 tree demo: construction

---

insert A



convert 3-node into 4-node

## 2-3 tree demo: construction

---

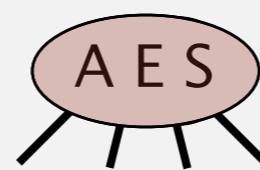
insert A



## 2-3 tree demo: construction

---

insert A

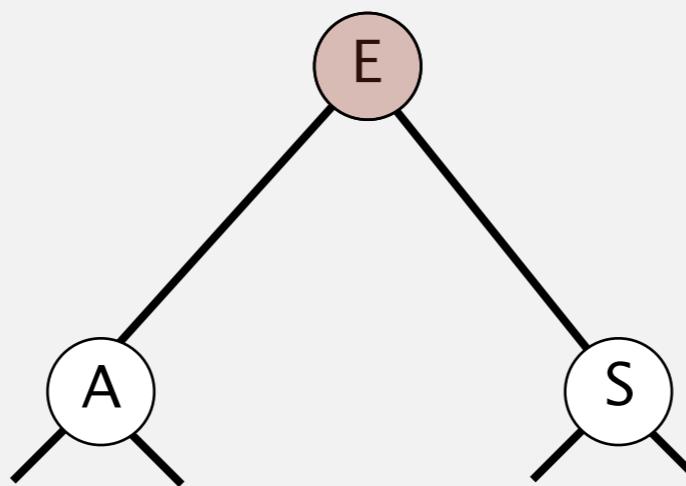


split 4-node  
(move E to parent)

## 2-3 tree demo: construction

---

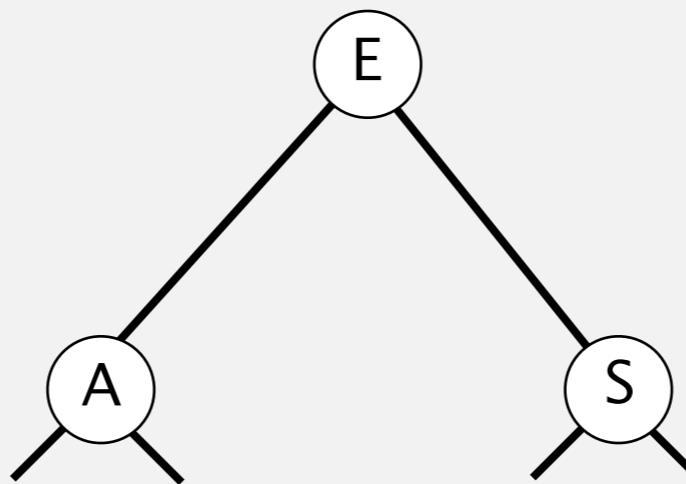
insert A



## 2-3 tree demo: construction

---

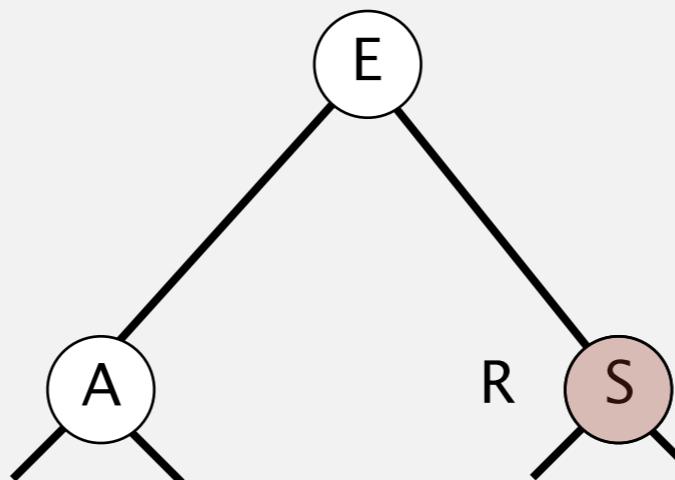
2-3 tree



## 2-3 tree demo: construction

---

insert R

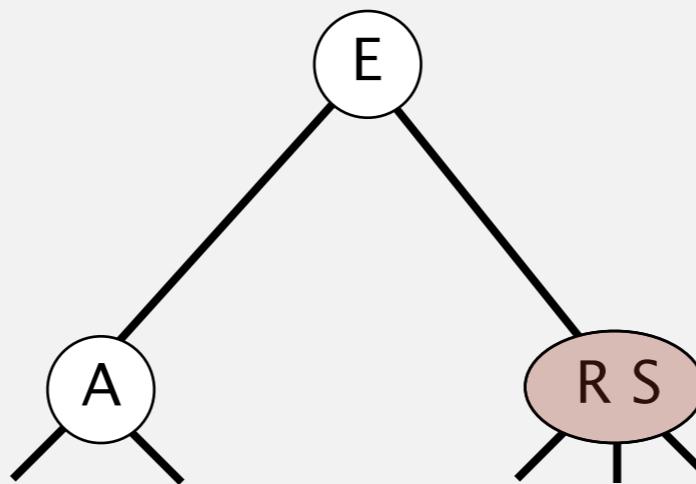


convert 2-node into 3-node

## 2-3 tree demo: construction

---

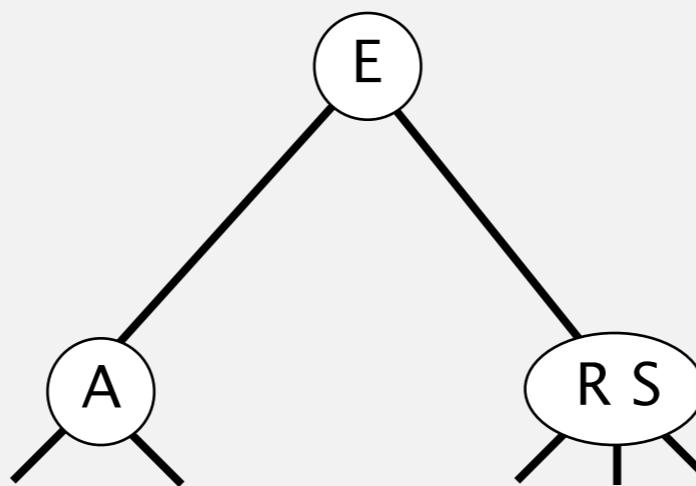
insert R



## 2-3 tree demo: construction

---

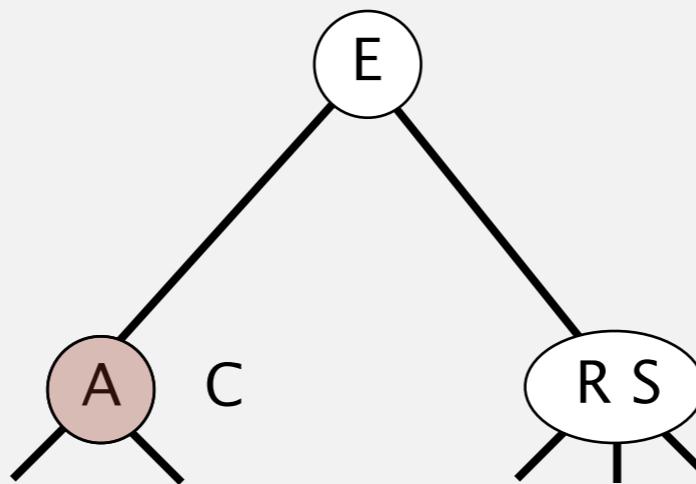
2-3 tree



## 2-3 tree demo: construction

---

insert C

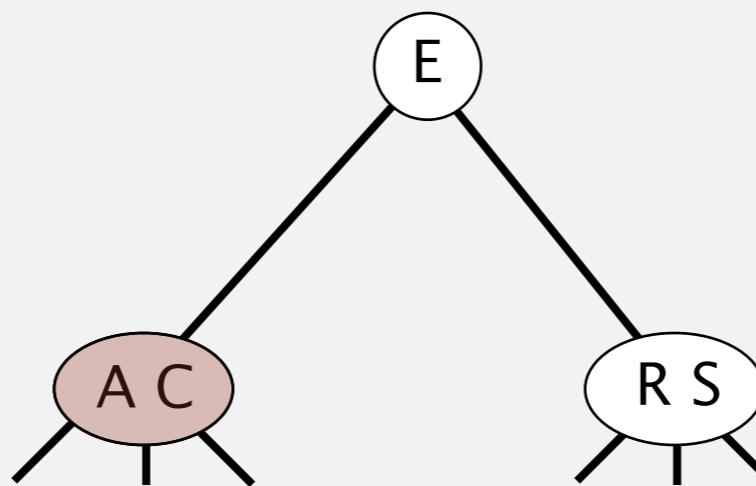


convert 2-node into 3-node

## 2-3 tree demo: construction

---

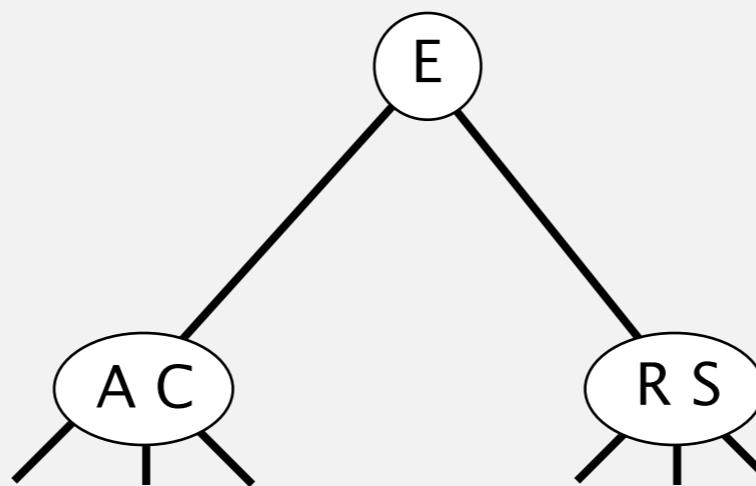
insert C



## 2-3 tree demo: construction

---

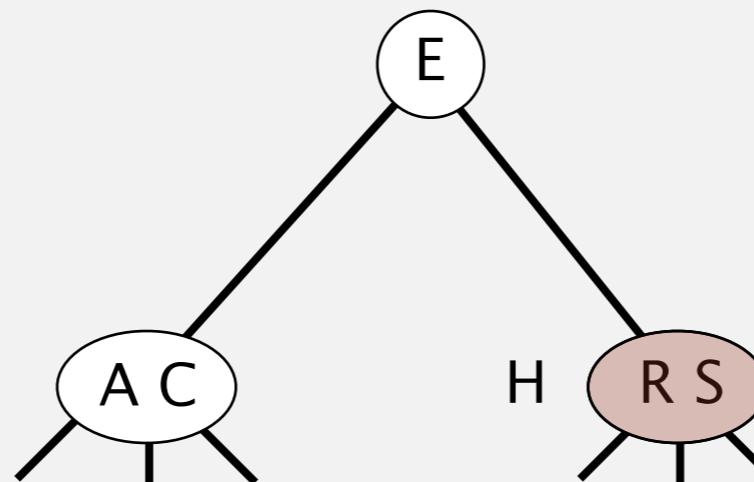
2-3 tree



## 2-3 tree demo: construction

---

insert H

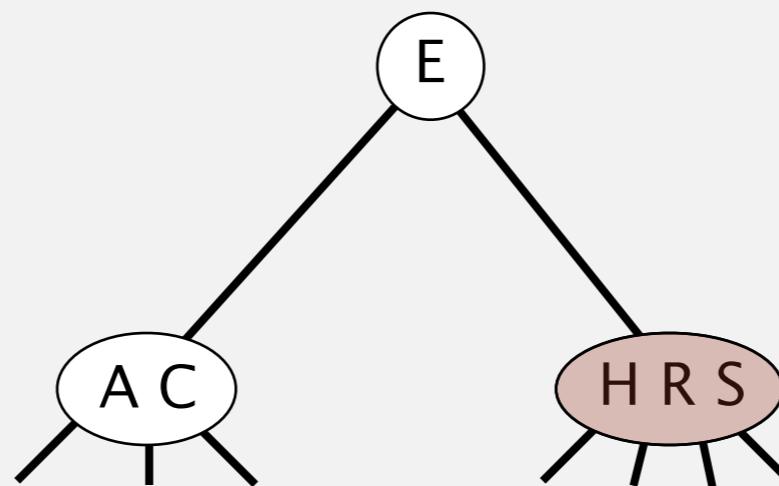


convert 3-node into 4-node

## 2-3 tree demo: construction

---

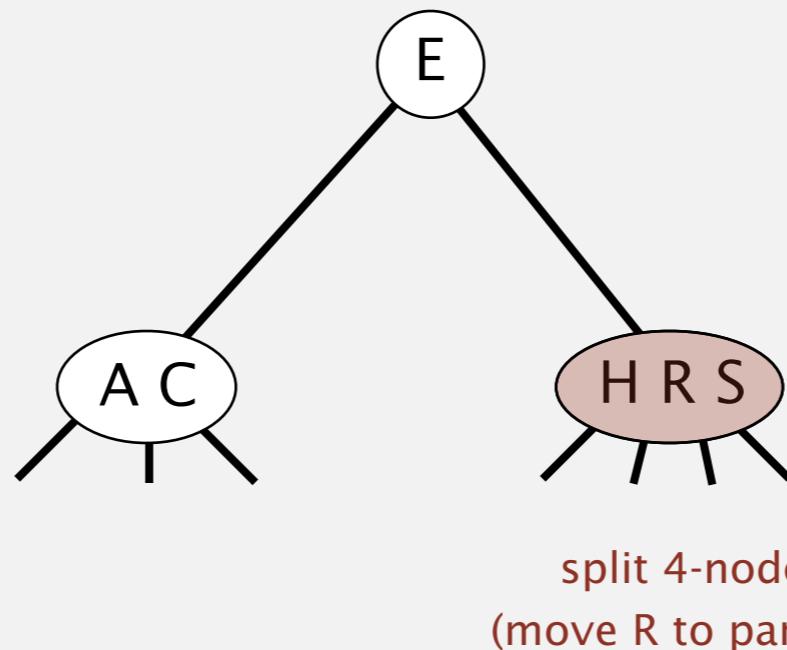
insert H



## 2-3 tree demo: construction

---

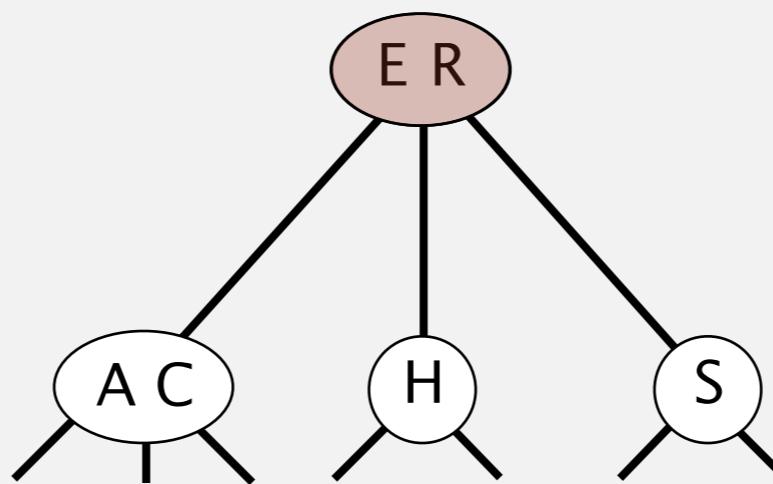
insert H



## 2-3 tree demo: construction

---

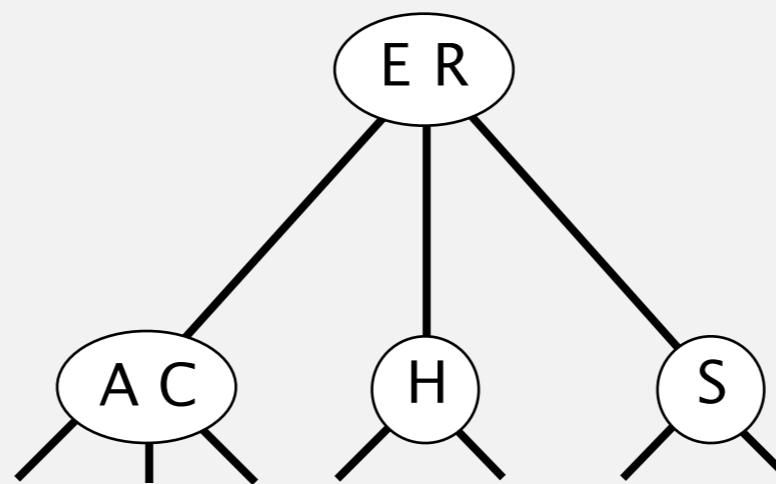
insert H



## 2-3 tree demo: construction

---

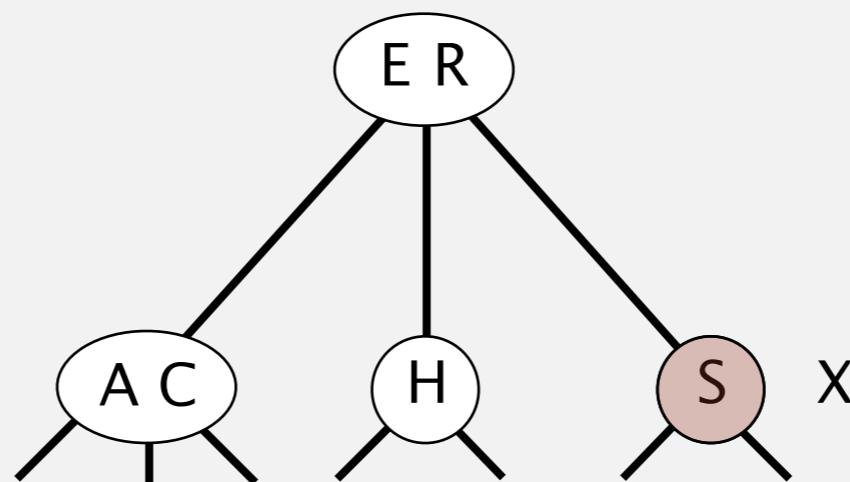
2-3 tree



## 2-3 tree demo: construction

---

insert X

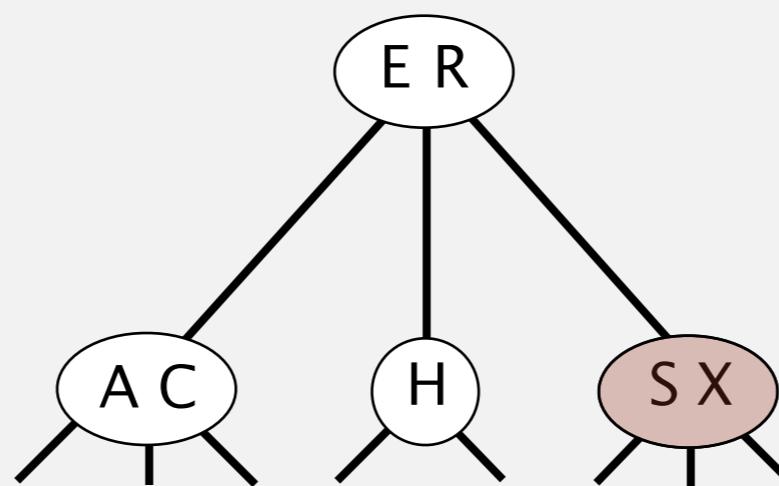


convert 2-node into 3-node

## 2-3 tree demo: construction

---

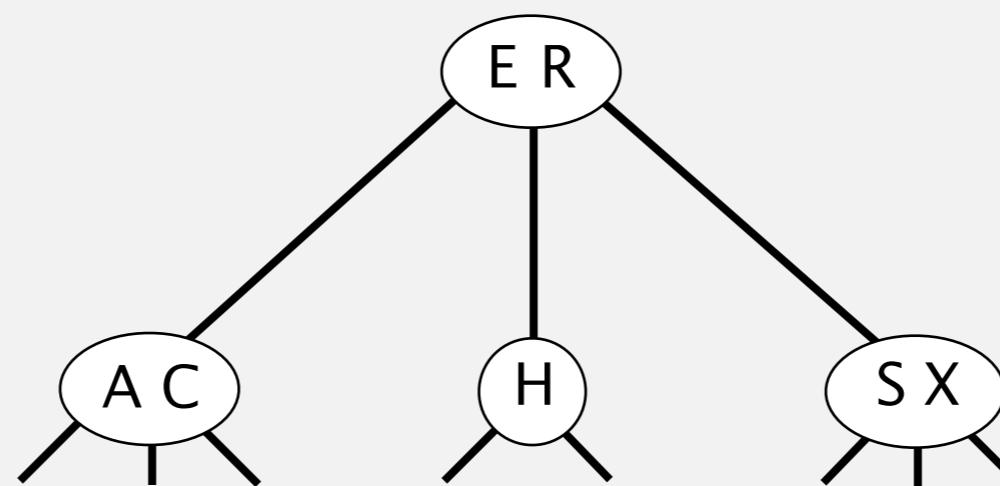
insert X



## 2-3 tree demo: construction

---

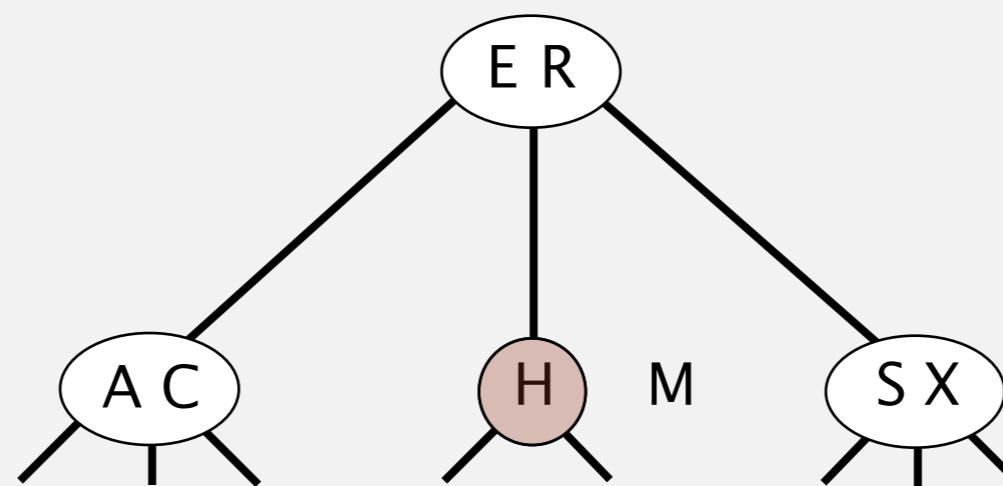
2-3 tree



## 2-3 tree demo: construction

---

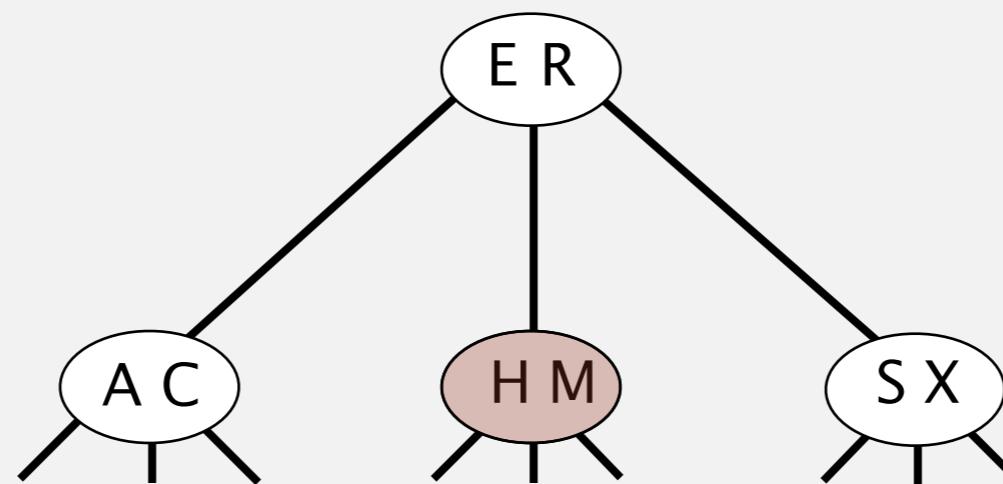
insert M



## 2-3 tree demo: construction

---

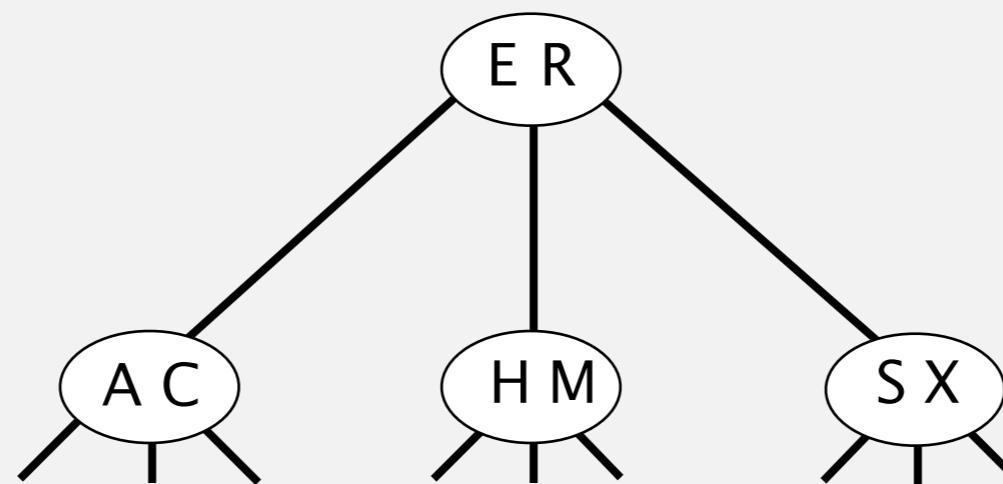
insert M



## 2-3 tree demo: construction

---

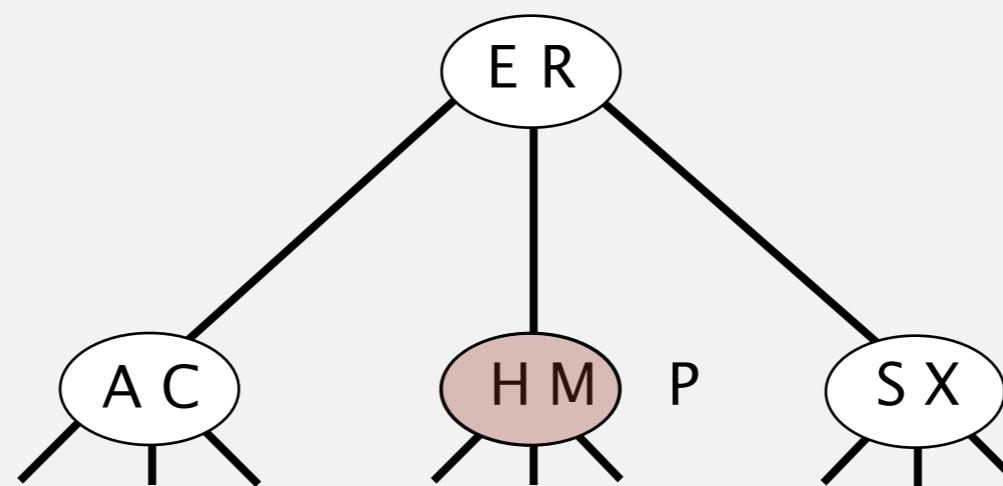
2-3 tree



## 2-3 tree demo: construction

---

insert P

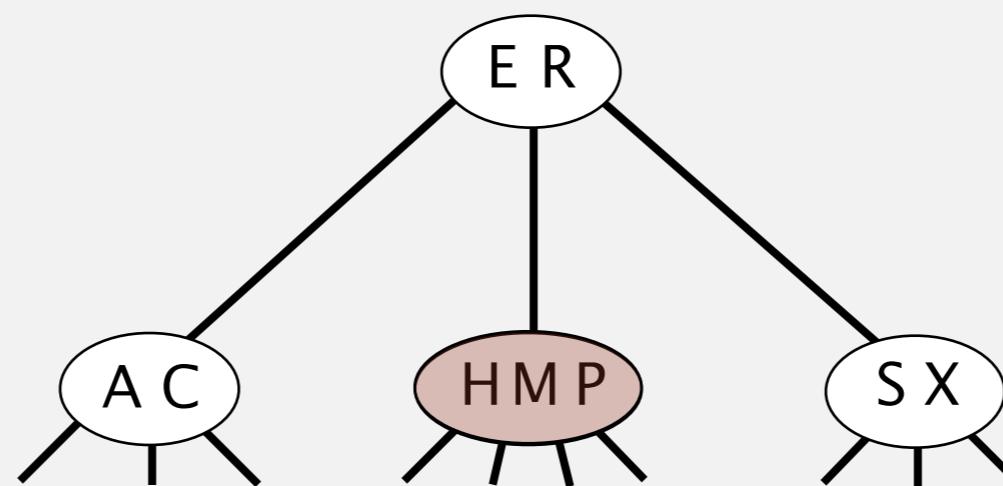


convert 3-node into 4-node

## 2-3 tree demo: construction

---

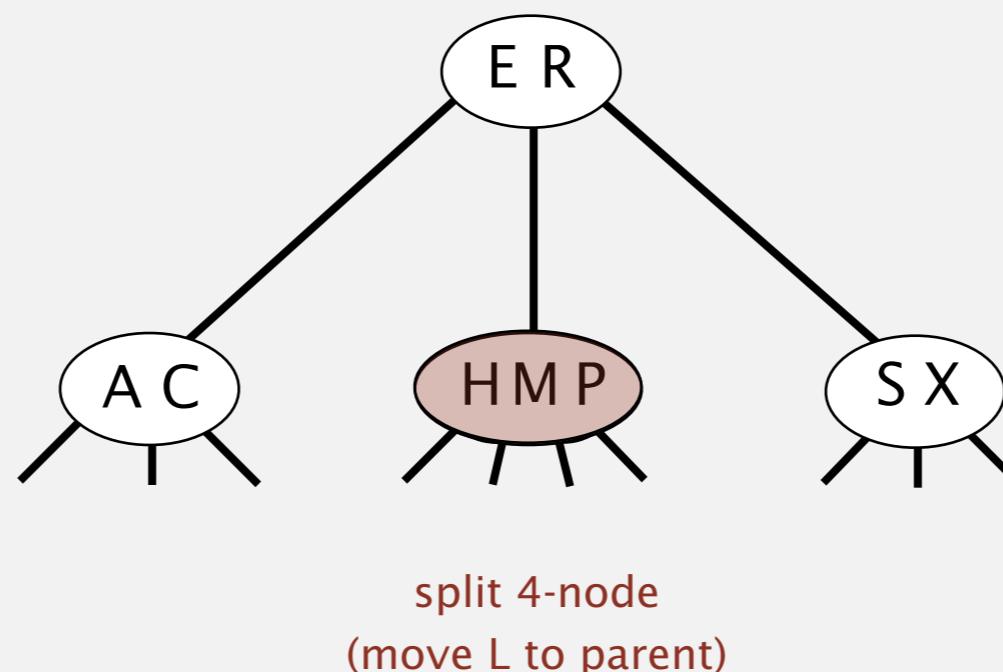
insert P



## 2-3 tree demo: construction

---

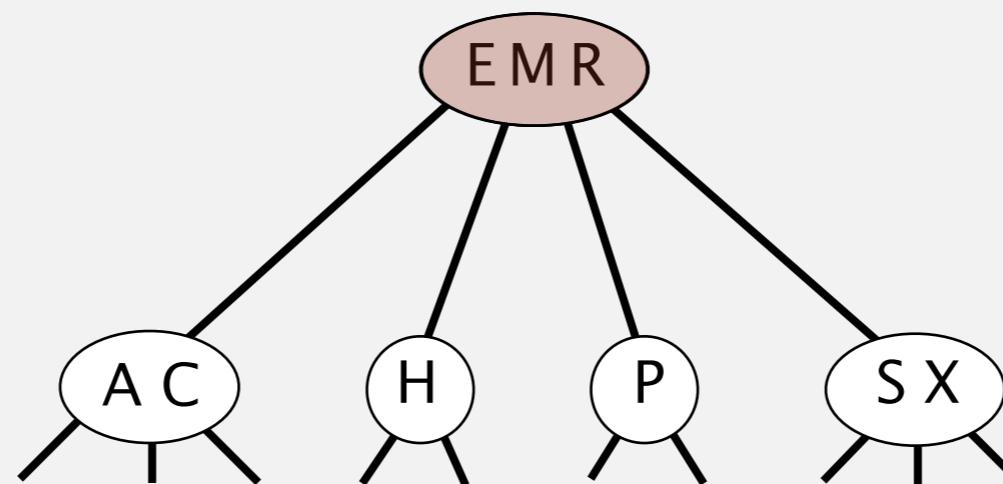
insert P



## 2-3 tree demo: construction

---

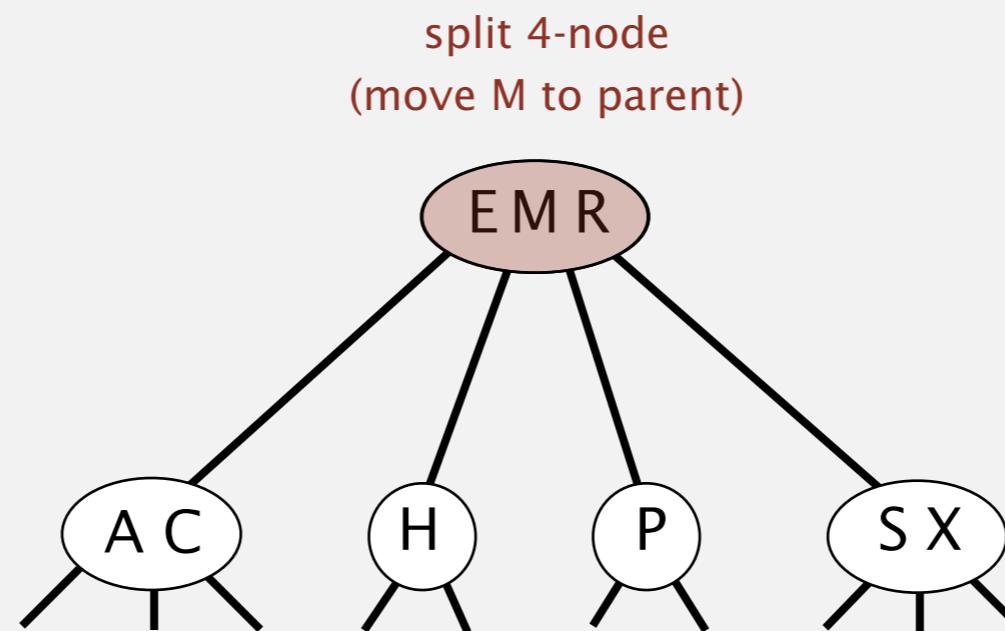
insert P



## 2-3 tree demo: construction

---

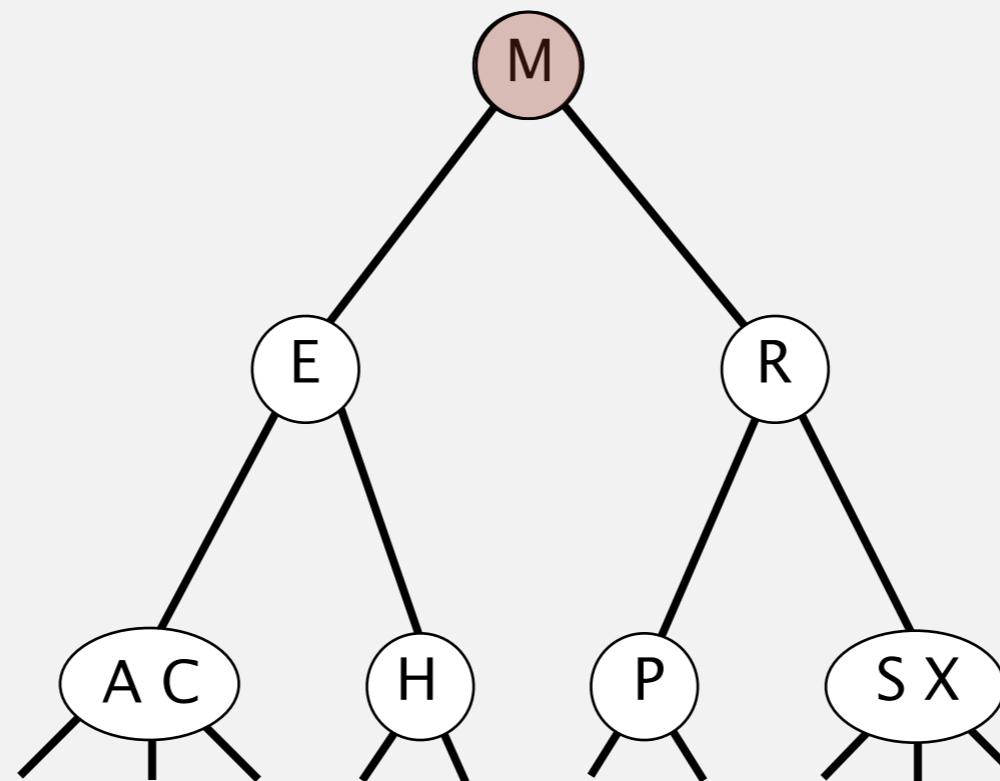
insert P



## 2-3 tree demo: construction

---

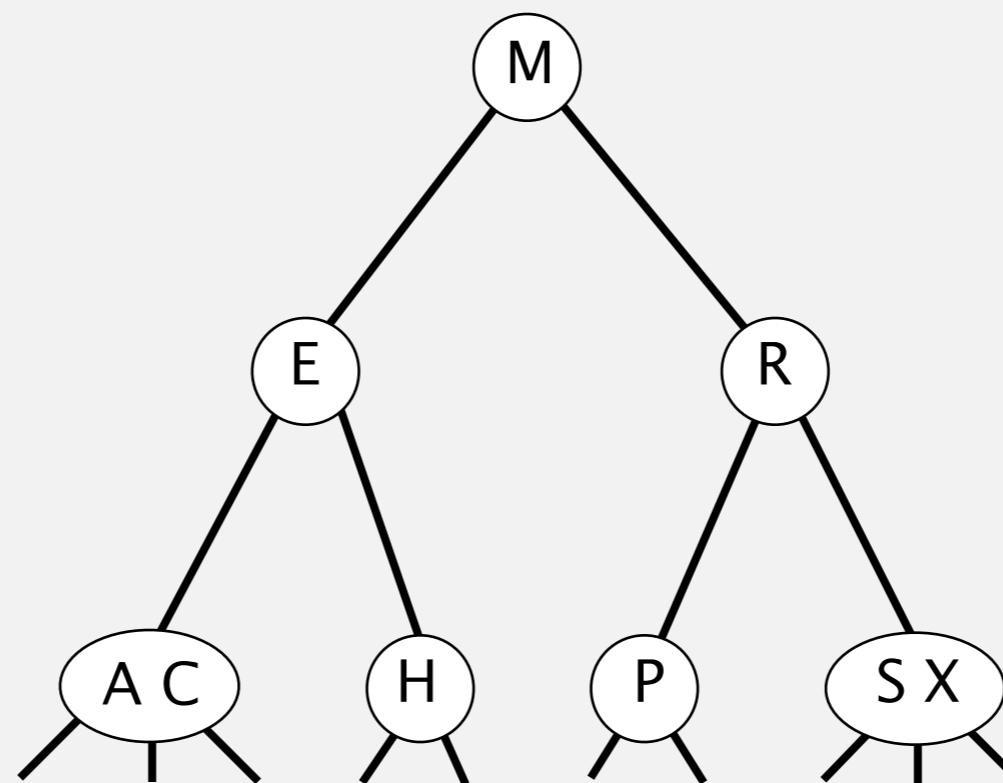
insert P



## 2-3 tree demo: construction

---

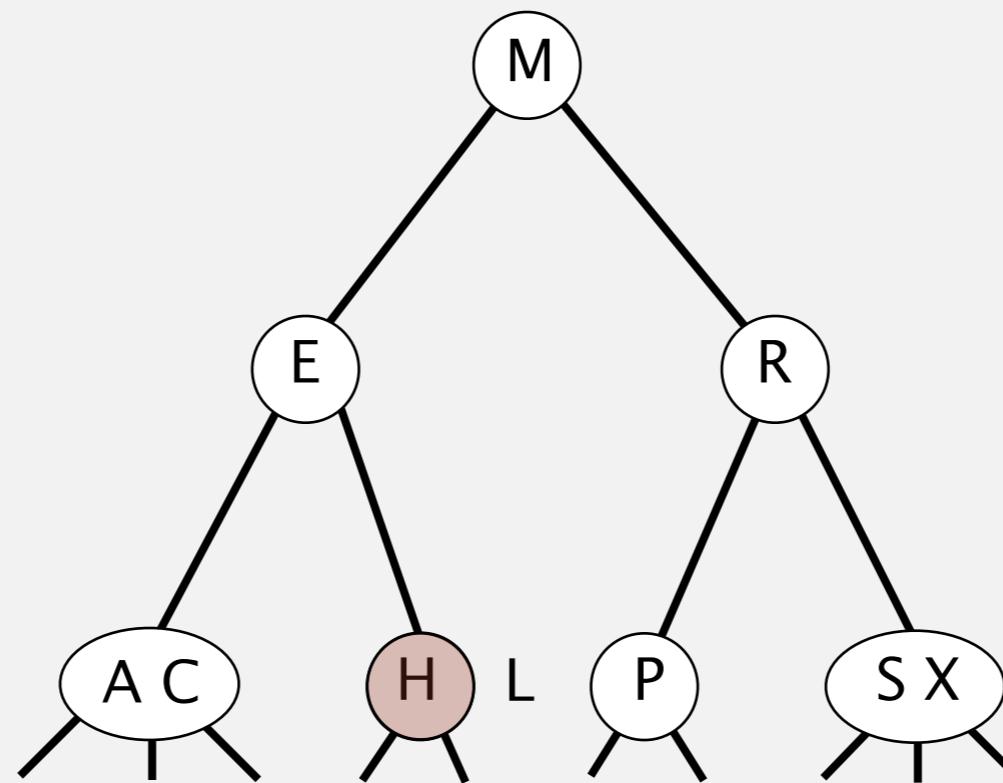
2-3 tree



## 2-3 tree demo: construction

---

insert L

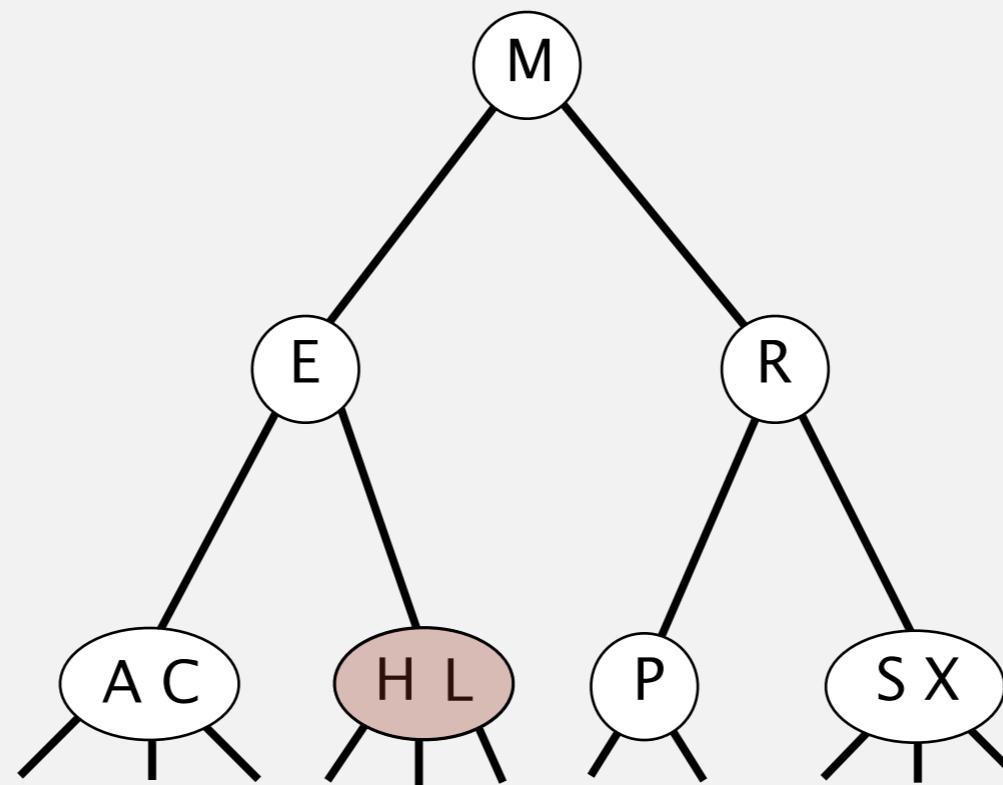


convert 2-node into 3-node

## 2-3 tree demo: construction

---

insert L

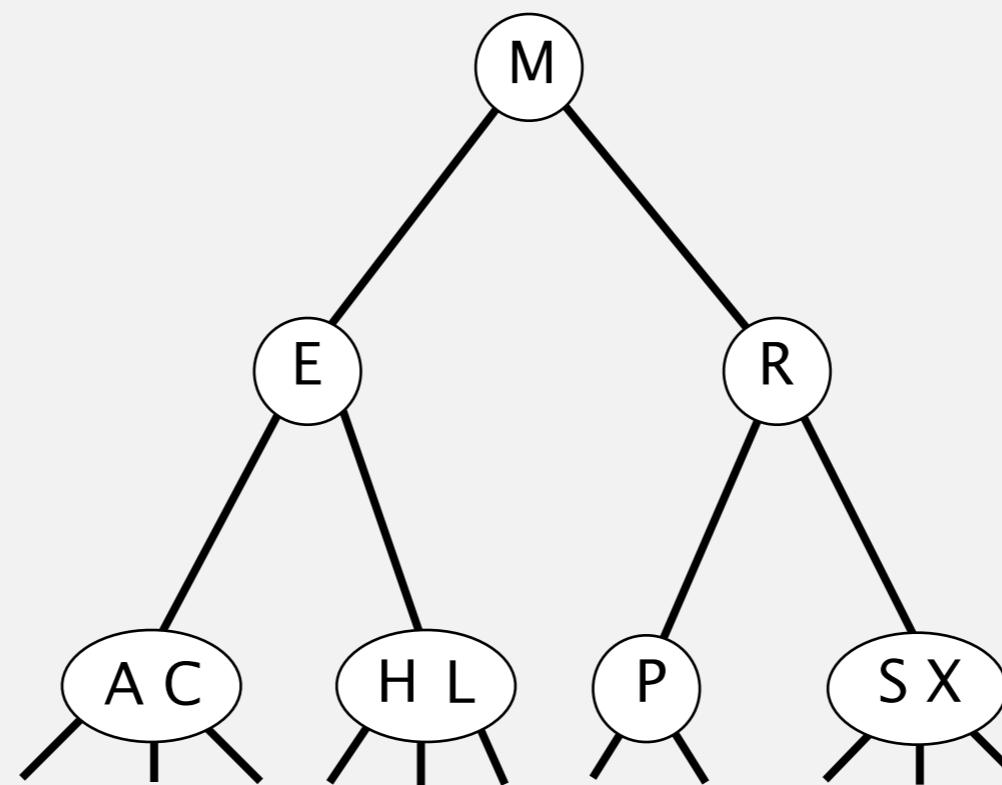


convert 2-node into 3-node

## 2-3 tree demo: construction

---

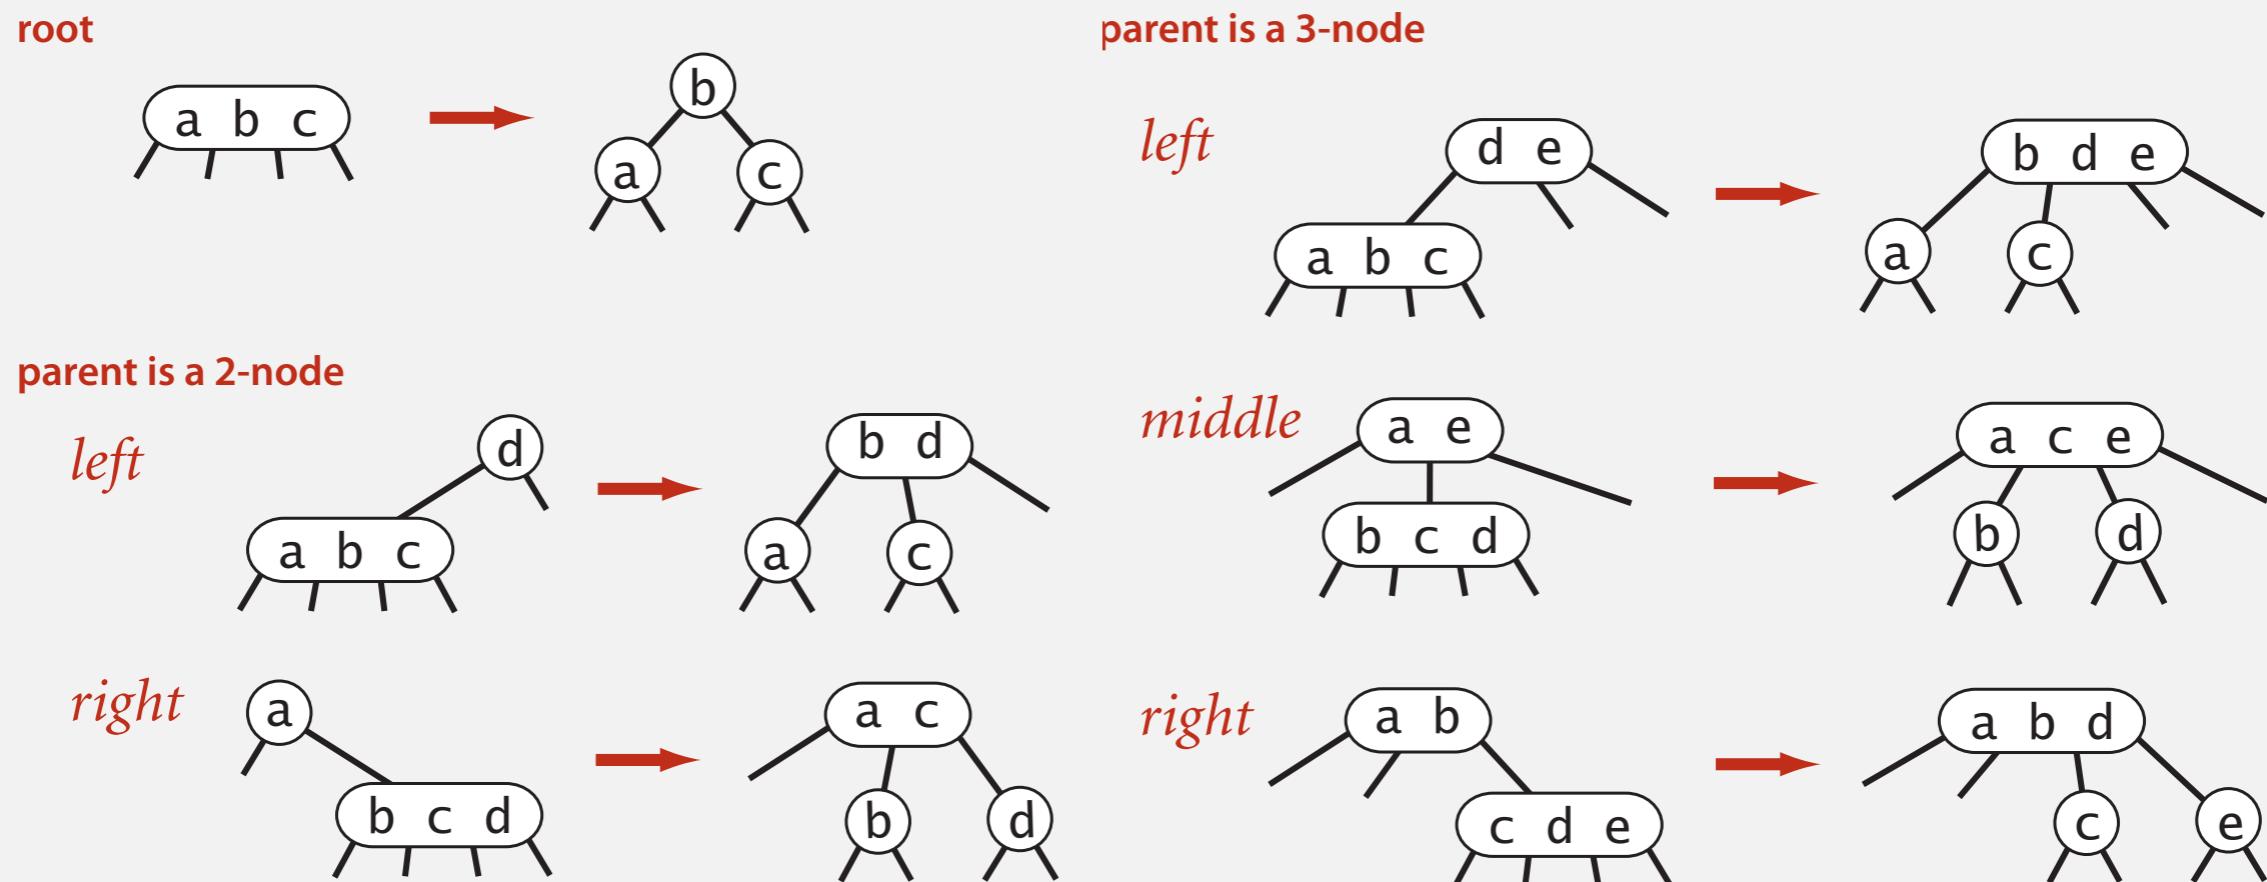
2-3 tree



## 2-3 tree: global properties

Invariants. Maintains symmetric order and perfect balance.

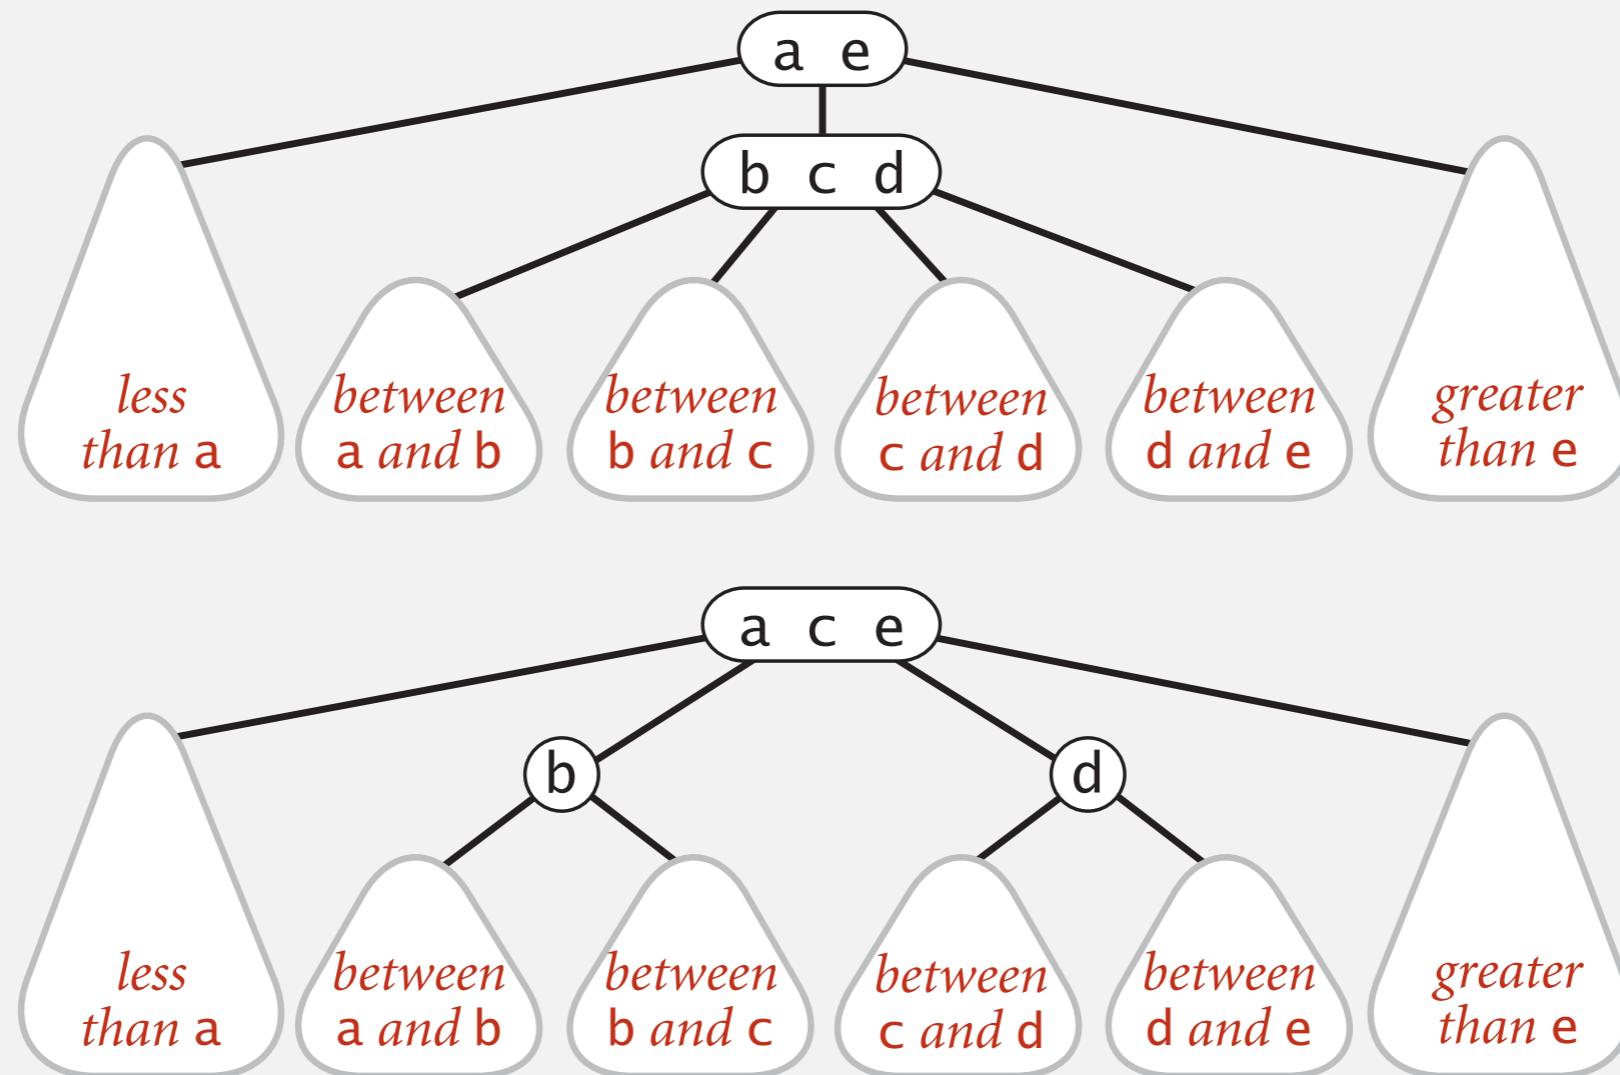
Pf. Each transformation maintains symmetric order and perfect balance.



## 2-3 tree: performance

---

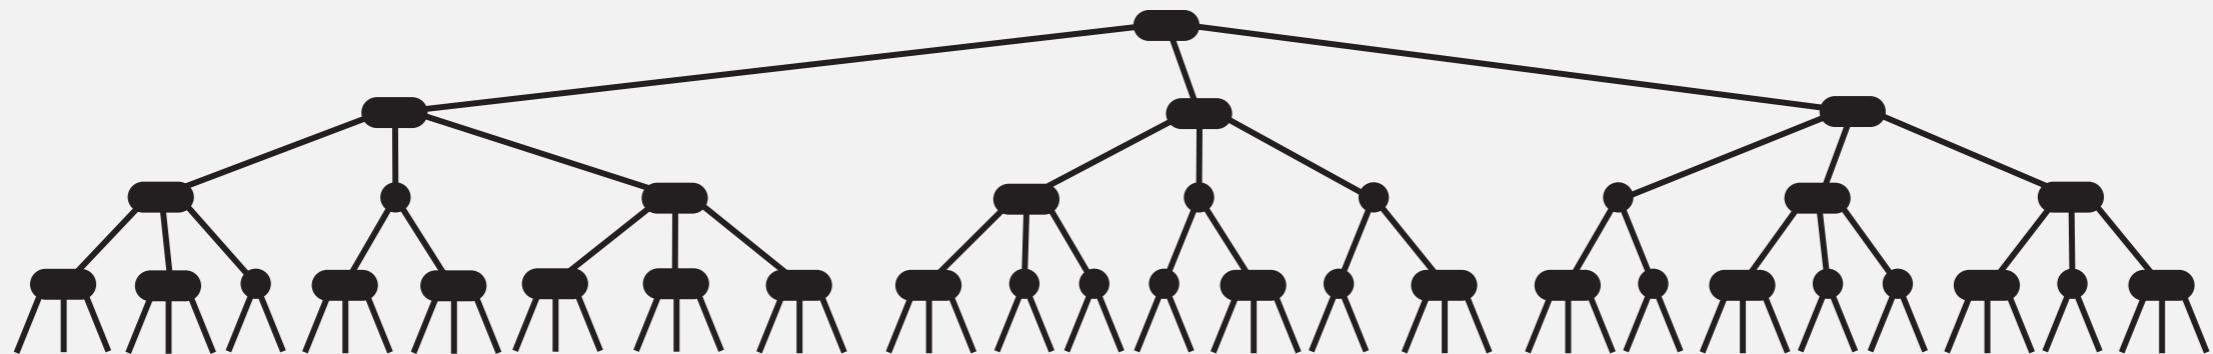
Splitting a 4-node is a **local** transformation: constant number of operations.



## 2–3 tree: performance

---

Perfect balance. Every path from root to null link has same length.



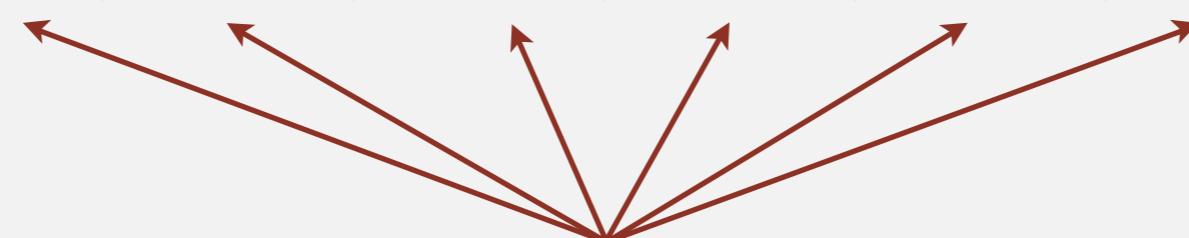
Tree height.

- Worst case:  $\lg N$ . [all 2-nodes]
- Best case:  $\log_3 N \approx .631 \lg N$ . [all 3-nodes]
- Between 12 and 20 for a million nodes.
- Between 18 and 30 for a billion nodes.

Bottom line. Guaranteed logarithmic performance for search and insert.

# ST implementations: summary

implementation	guarantee			average case			ordered ops?	key interface
	search	insert	delete	search hit	insert	delete		
<b>sequential search (unordered list)</b>	$N$	$N$	$N$	$N$	$N$	$N$		<code>equals()</code>
<b>binary search (ordered array)</b>	$\log N$	$N$	$N$	$\log N$	$N$	$N$	✓	<code>compareTo()</code>
<b>BST</b>	$N$	$N$	$N$	$\log N$	$\log N$	$\sqrt{N}$	✓	<code>compareTo()</code>
<b>2-3 tree</b>	$\log N$	$\log N$	$\log N$	$\log N$	$\log N$	$\log N$	✓	<code>compareTo()</code>



but hidden constant  $c$  is large  
(depends upon implementation)

## 2-3 tree: implementation?

---

Direct implementation is complicated, because:

- Maintaining multiple node types is cumbersome.
- Need multiple compares to move down tree.
- Need to move back up the tree to split 4-nodes.
- Large number of cases for splitting.

**fantasy code**

```
public void put(Key key, Value val)
{
    Node x = root;
    while (x.getTheCorrectChild(key) != null)
    {
        x = x.getTheCorrectChildKey();
        if (x.is4Node()) x.split();
    }
    if (x.is2Node()) x.make3Node(key, val);
    else if (x.is3Node()) x.make4Node(key, val);
}
```

Bottom line. Could do it, but there's a better way.

# Algorithms

ROBERT SEDGEWICK | KEVIN WAYNE

<http://algs4.cs.princeton.edu>

## 3.3 BALANCED SEARCH TREES

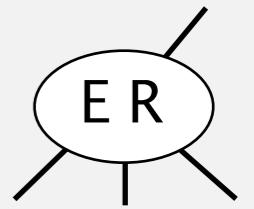
---

- ▶ *2-3 search trees*
- ▶ *red-black BSTs*

# How to implement 2–3 trees with binary trees?

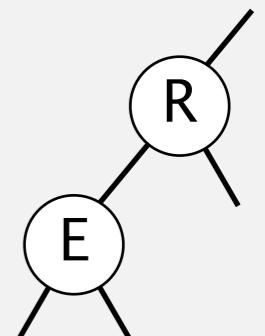
---

Challenge. How to represent a 3 node?



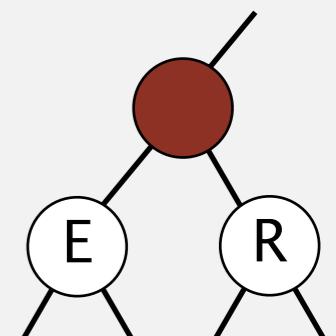
Approach 1. Regular BST.

- No way to tell a 3-node from a 2-node.
- Cannot map from BST back to 2–3 tree.



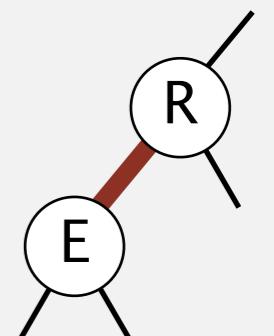
Approach 2. Regular BST with red "glue" nodes.

- Wastes space, wasted link.
- Code probably messy.



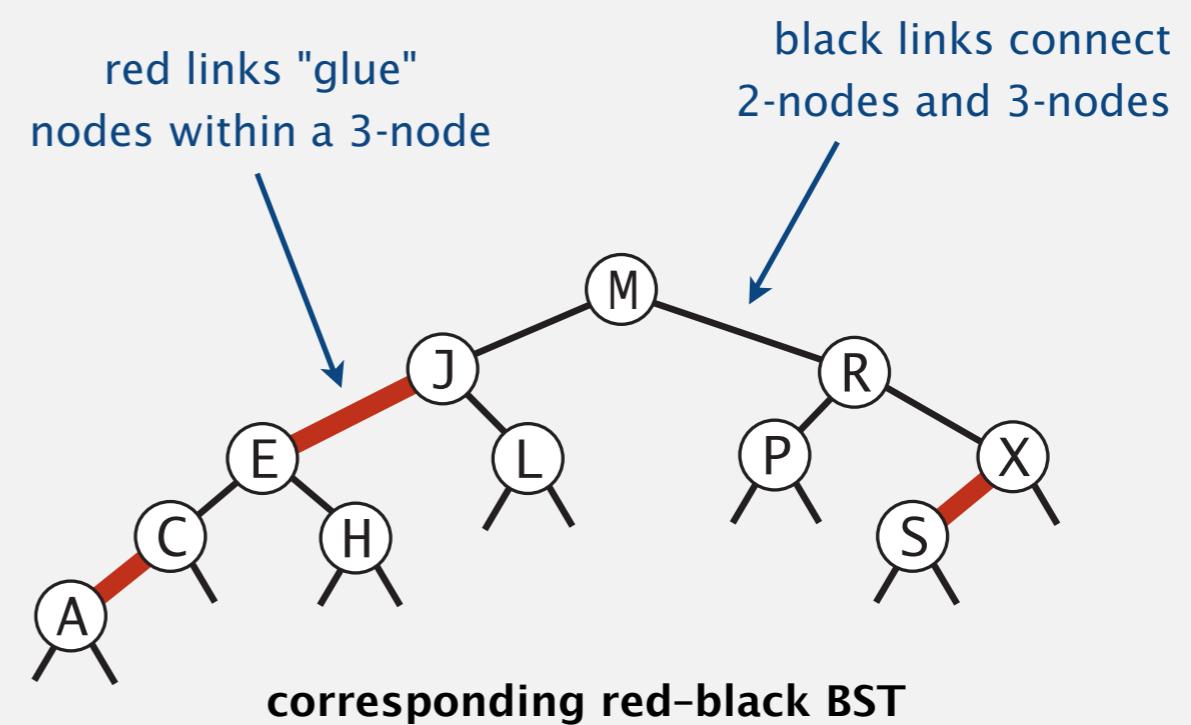
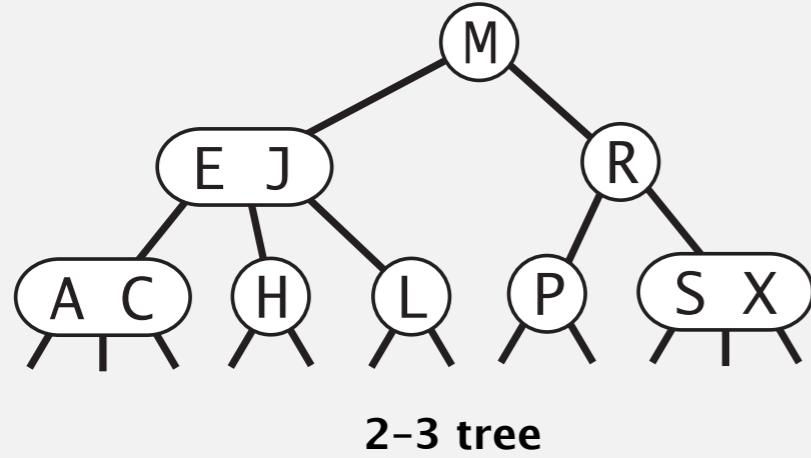
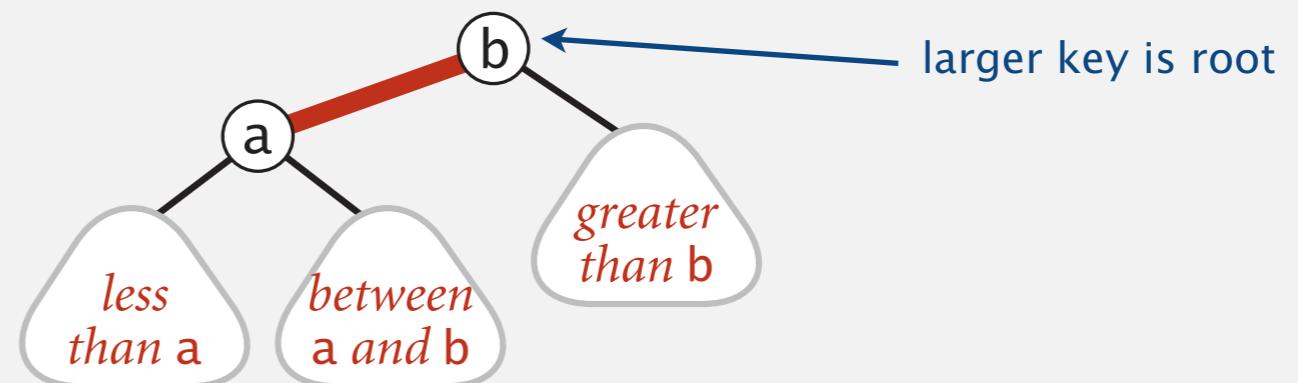
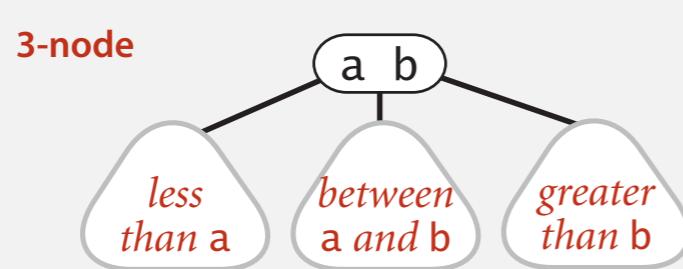
Approach 3. Regular BST with red "glue" links.

- Widely used in practice.
- Arbitrary restriction: red links lean left.



# Left-leaning red-black BSTs (Guibas-Sedgewick 1979 and Sedgewick 2007)

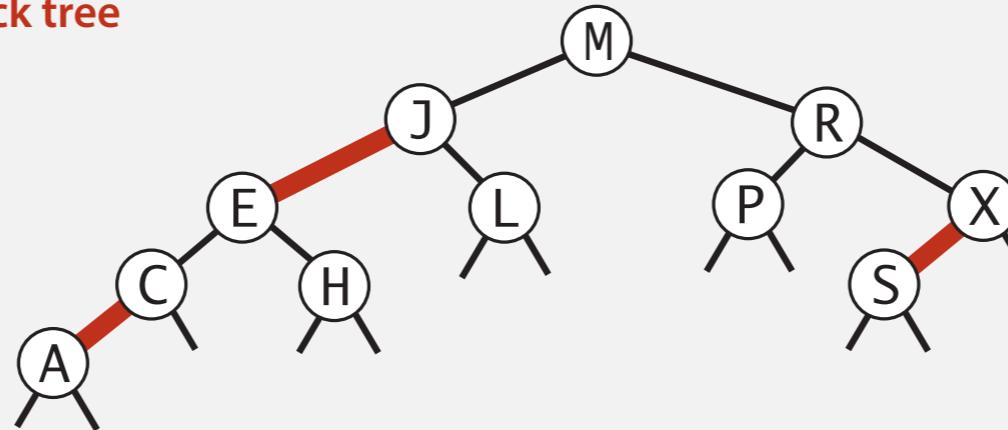
1. Represent 2–3 tree as a BST.
2. Use "internal" left-leaning links as "glue" for 3-nodes.



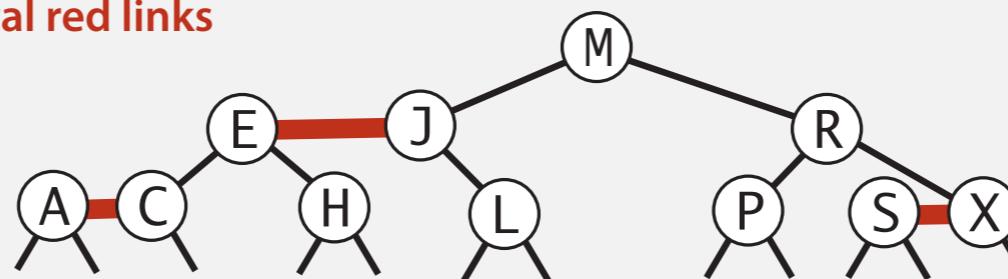
# Left-leaning red-black BSTs: 1–1 correspondence with 2–3 trees

Key property. 1–1 correspondence between 2–3 and LLRB.

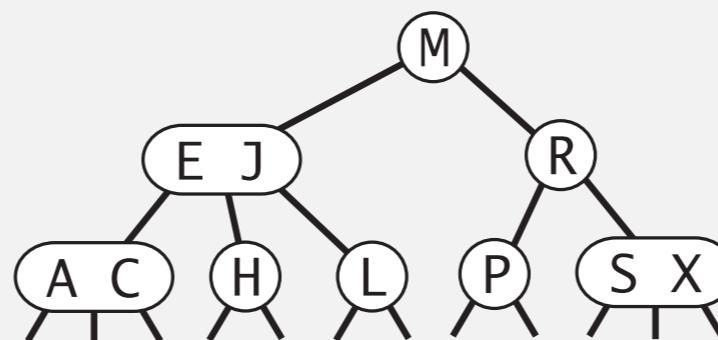
red–black tree



horizontal red links



2-3 tree



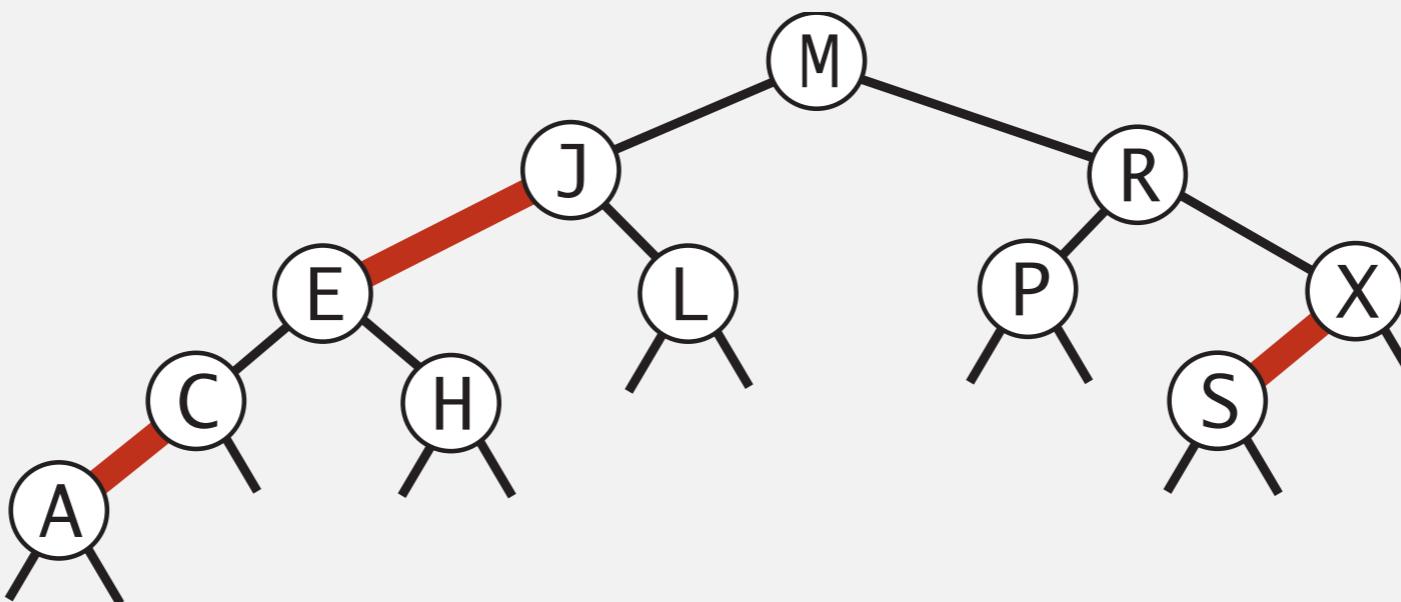
## An equivalent definition

---

A BST such that:

- No node has two red links connected to it.
- Every path from root to null link has the same number of black links.
- Red links lean left.

"perfect black balance"

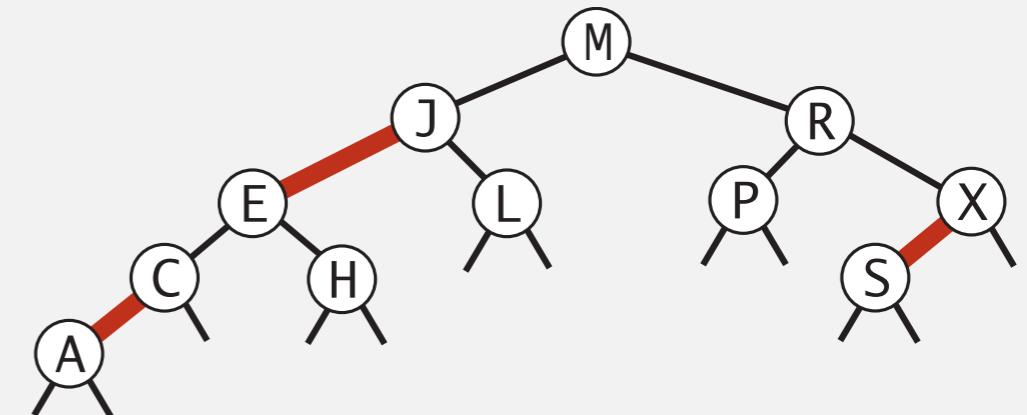


# Search implementation for red-black BSTs

Observation. Search is the same as for elementary BST (ignore color).

but runs faster because  
of better balance

```
public Value get(Key key)
{
    Node x = root;
    while (x != null)
    {
        int cmp = key.compareTo(x.key);
        if      (cmp < 0) x = x.left;
        else if (cmp > 0) x = x.right;
        else if (cmp == 0) return x.val;
    }
    return null;
}
```



Remark. Most other ops (e.g., floor, iteration, selection) are also identical.

# Red-black BST representation

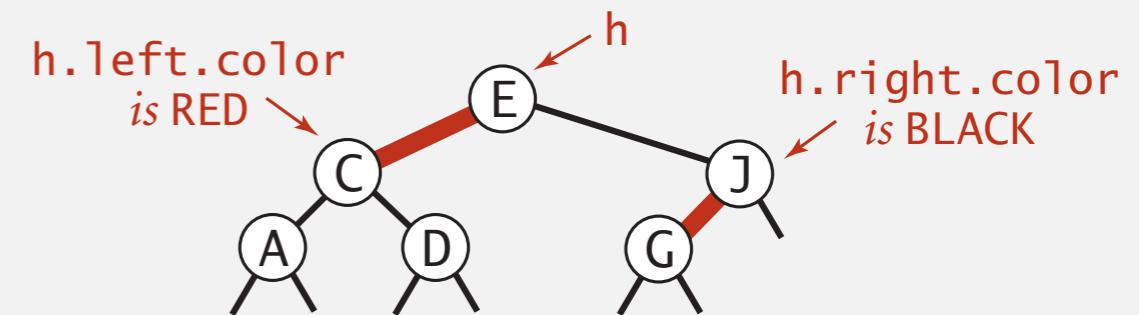
Each node is pointed to by precisely one link (from its parent)  $\Rightarrow$   
can encode color of links in nodes.

```
private static final boolean RED = true;
private static final boolean BLACK = false;
```

```
private class Node
{
    Key key;
    Value val;
    Node left, right;
    boolean color; // color of parent link
}
```

```
private boolean isRed(Node x)
{
    if (x == null) return false;
    return x.color == RED;
}
```

null links are black



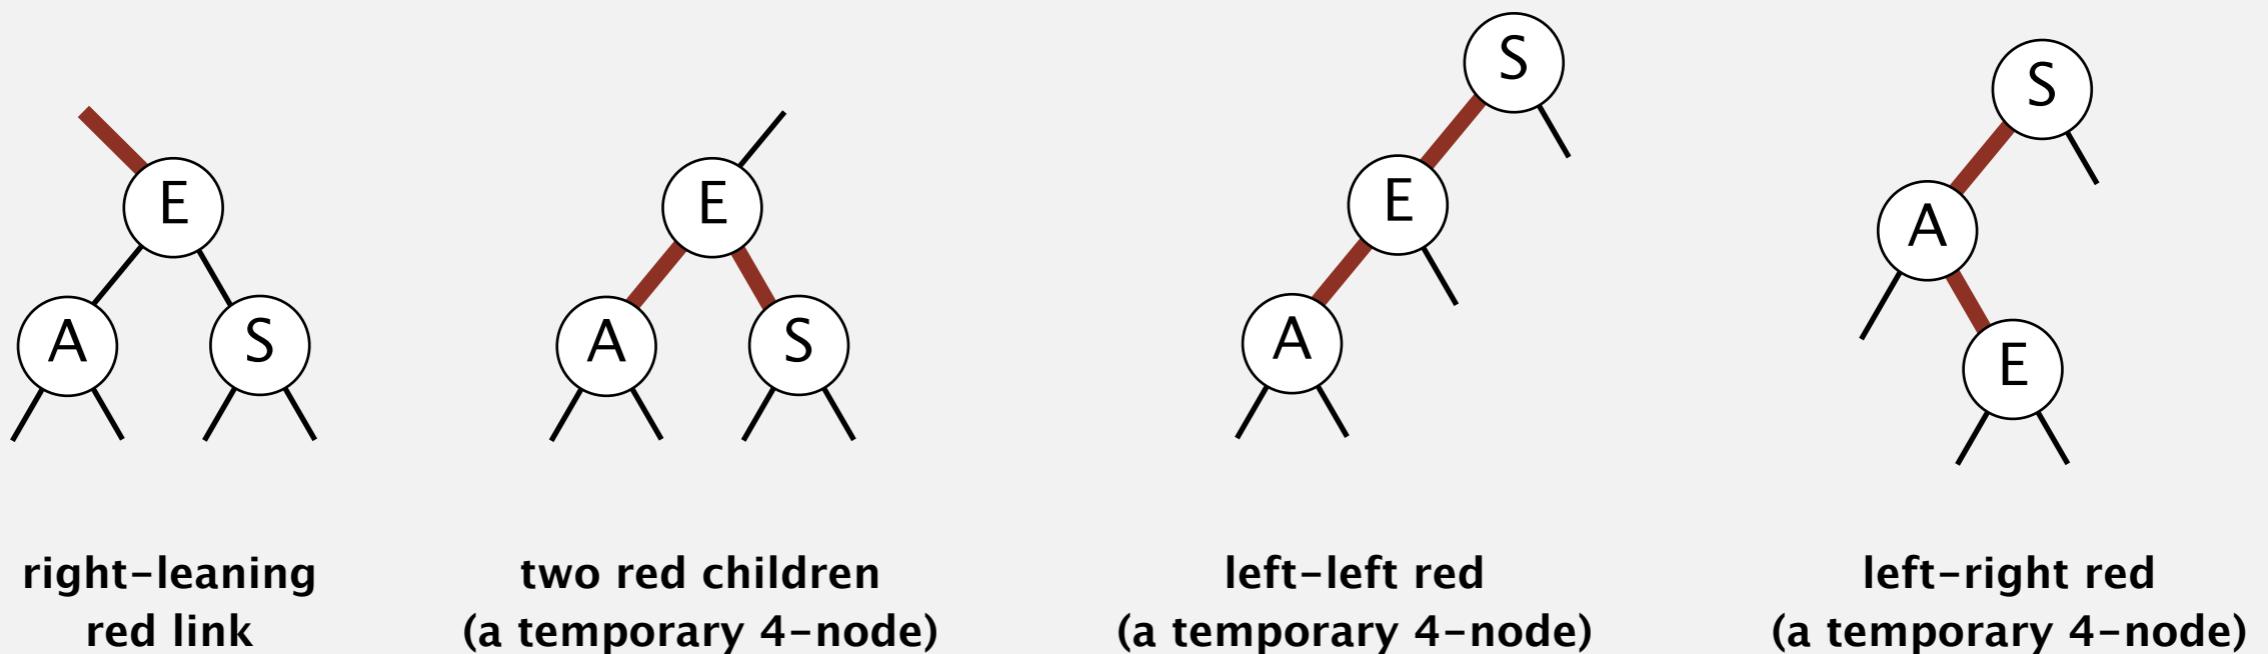
# Insertion into a LLRB tree: overview

---

Basic strategy. Maintain 1–1 correspondence with 2–3 trees.

During internal operations, maintain:

- Symmetric order.
- Perfect black balance.  
[ but not necessarily color invariants ]



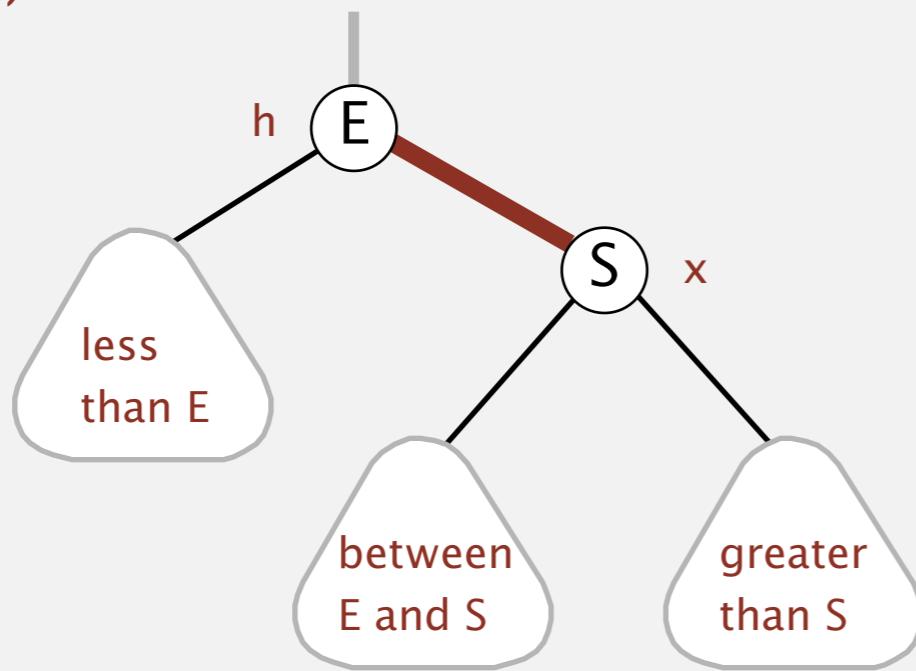
How? Apply elementary red–black BST operations: rotation and color flip.

# Elementary red-black BST operations

Left rotation. Orient a (temporarily) right-leaning red link to lean left.

rotate E left

(before)



```
private Node rotateLeft(Node h)
{
    assert isRed(h.right);
    Node x = h.right;
    h.right = x.left;
    x.left = h;
    x.color = h.color;
    h.color = RED;
    return x;
}
```

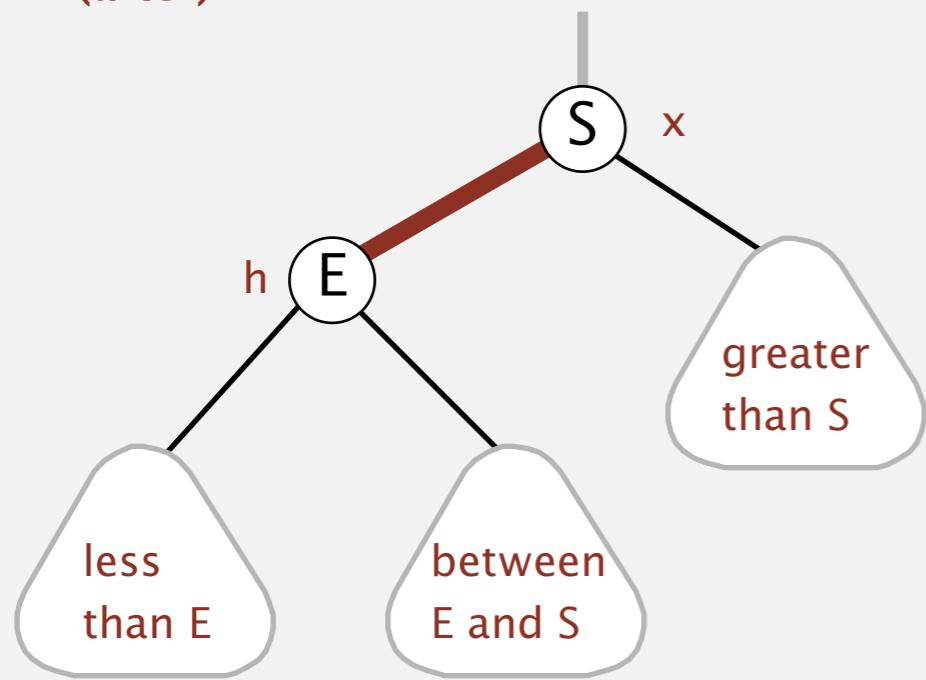
Invariants. Maintains symmetric order and perfect black balance.

# Elementary red-black BST operations

Left rotation. Orient a (temporarily) right-leaning red link to lean left.

rotate E left

(after)



```
private Node rotateLeft(Node h)
{
    assert isRed(h.right);
    Node x = h.right;
    h.right = x.left;
    x.left = h;
    x.color = h.color;
    h.color = RED;
    return x;
}
```

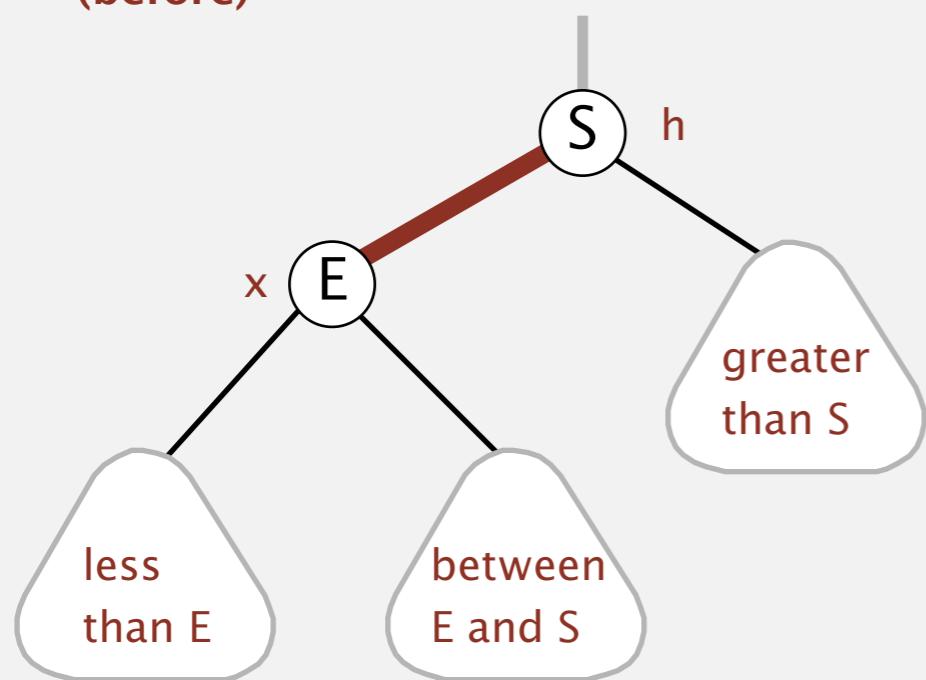
Invariants. Maintains symmetric order and perfect black balance.

# Elementary red-black BST operations

Right rotation. Orient a left-leaning red link to (temporarily) lean right.

rotate S right

(before)



```
private Node rotateRight(Node h)  
{  
    assert isRed(h.left);  
    Node x = h.left;  
    h.left = x.right;  
    x.right = h;  
    x.color = h.color;  
    h.color = RED;  
    return x;  
}
```

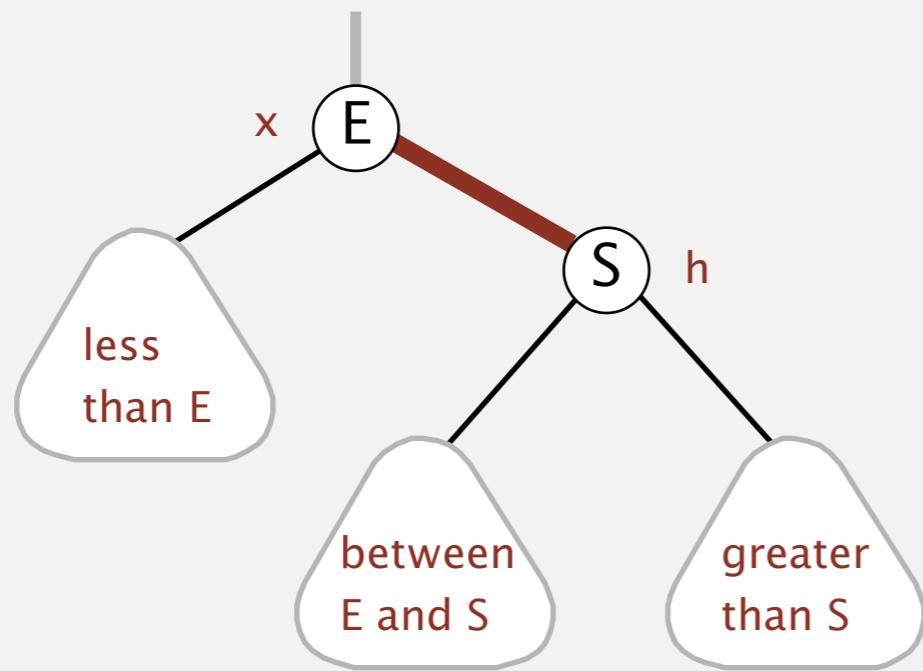
Invariants. Maintains symmetric order and perfect black balance.

# Elementary red-black BST operations

Right rotation. Orient a left-leaning red link to (temporarily) lean right.

rotate S right

(after)

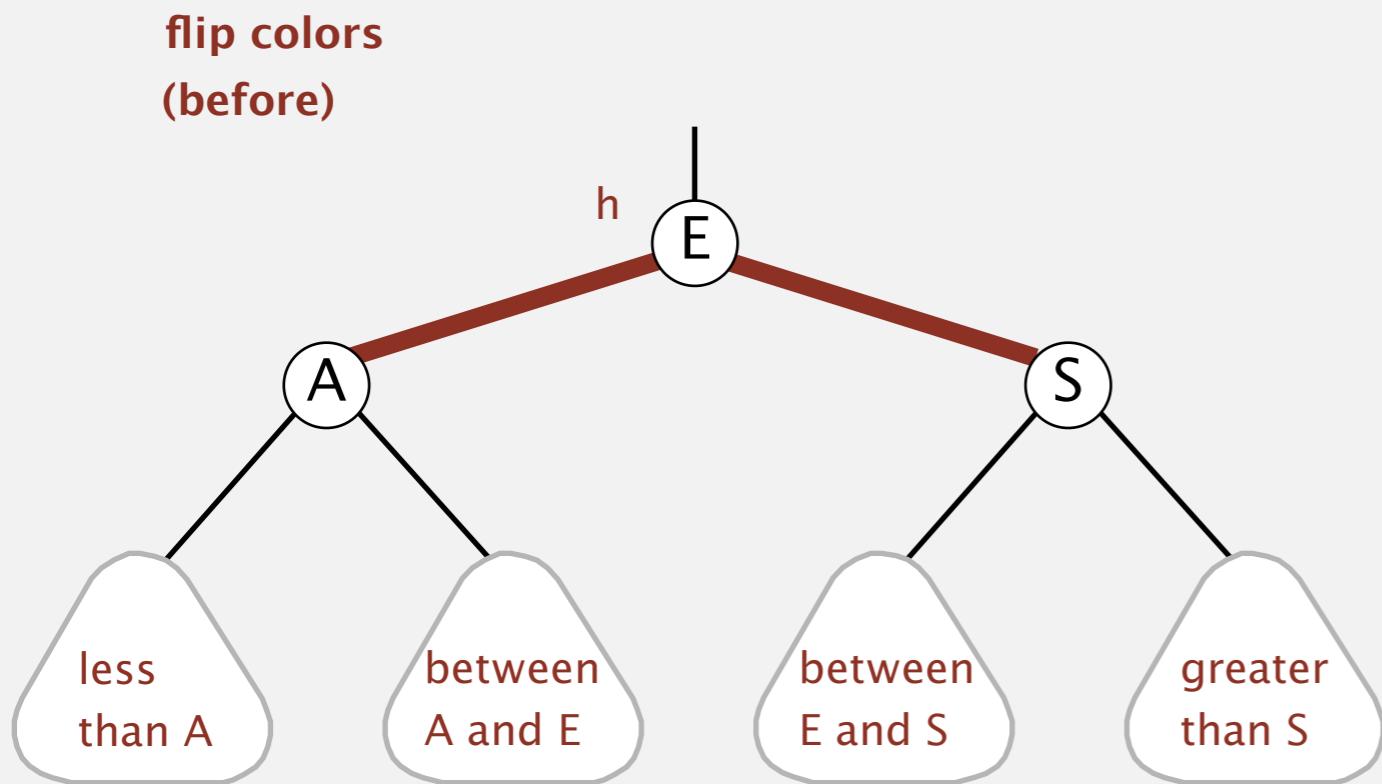


```
private Node rotateRight(Node h)
{
    assert isRed(h.left);
    Node x = h.left;
    h.left = x.right;
    x.right = h;
    x.color = h.color;
    h.color = RED;
    return x;
}
```

Invariants. Maintains symmetric order and perfect black balance.

# Elementary red-black BST operations

Color flip. Recolor to split a (temporary) 4-node.

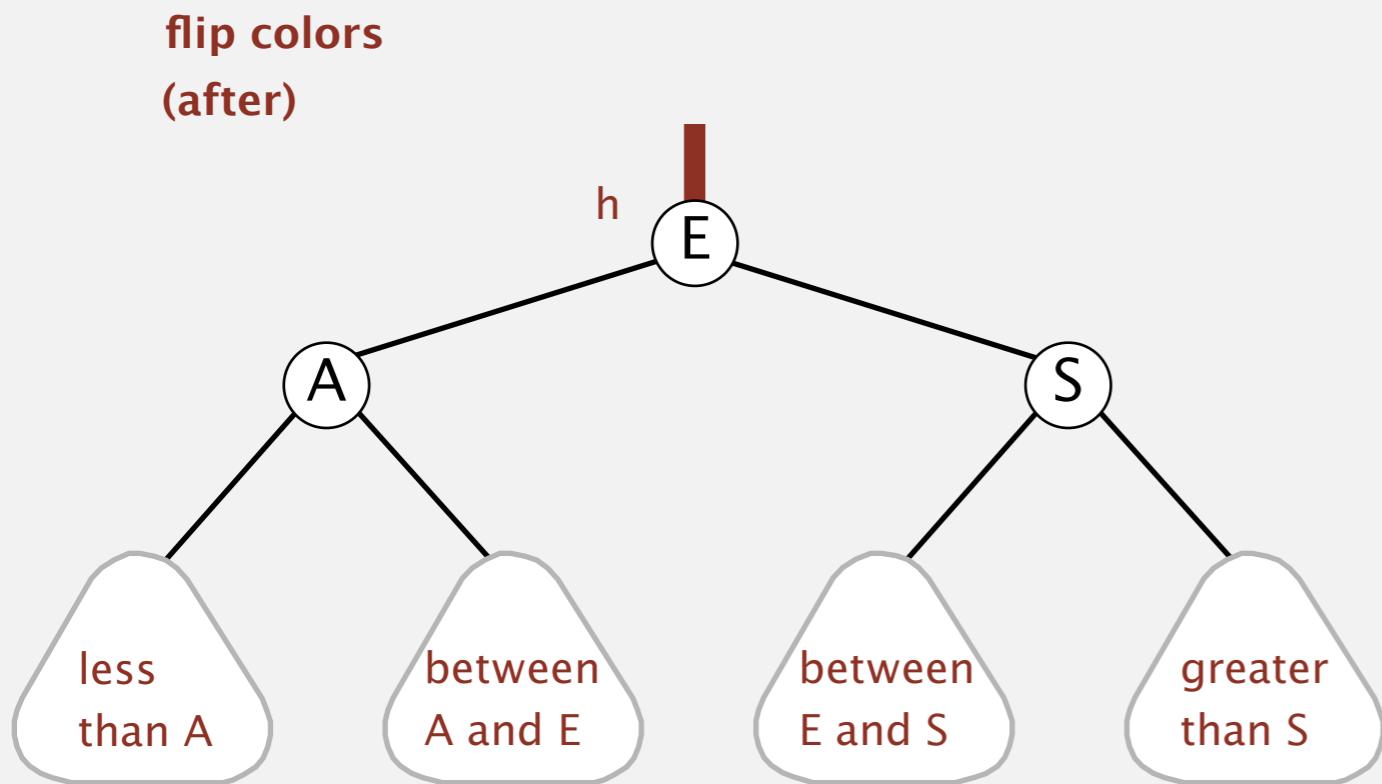


```
private void flipColors(Node h)
{
    assert !isRed(h);
    assert isRed(h.left);
    assert isRed(h.right);
    h.color = RED;
    h.left.color = BLACK;
    h.right.color = BLACK;
}
```

Invariants. Maintains symmetric order and perfect black balance.

# Elementary red-black BST operations

Color flip. Recolor to split a (temporary) 4-node.



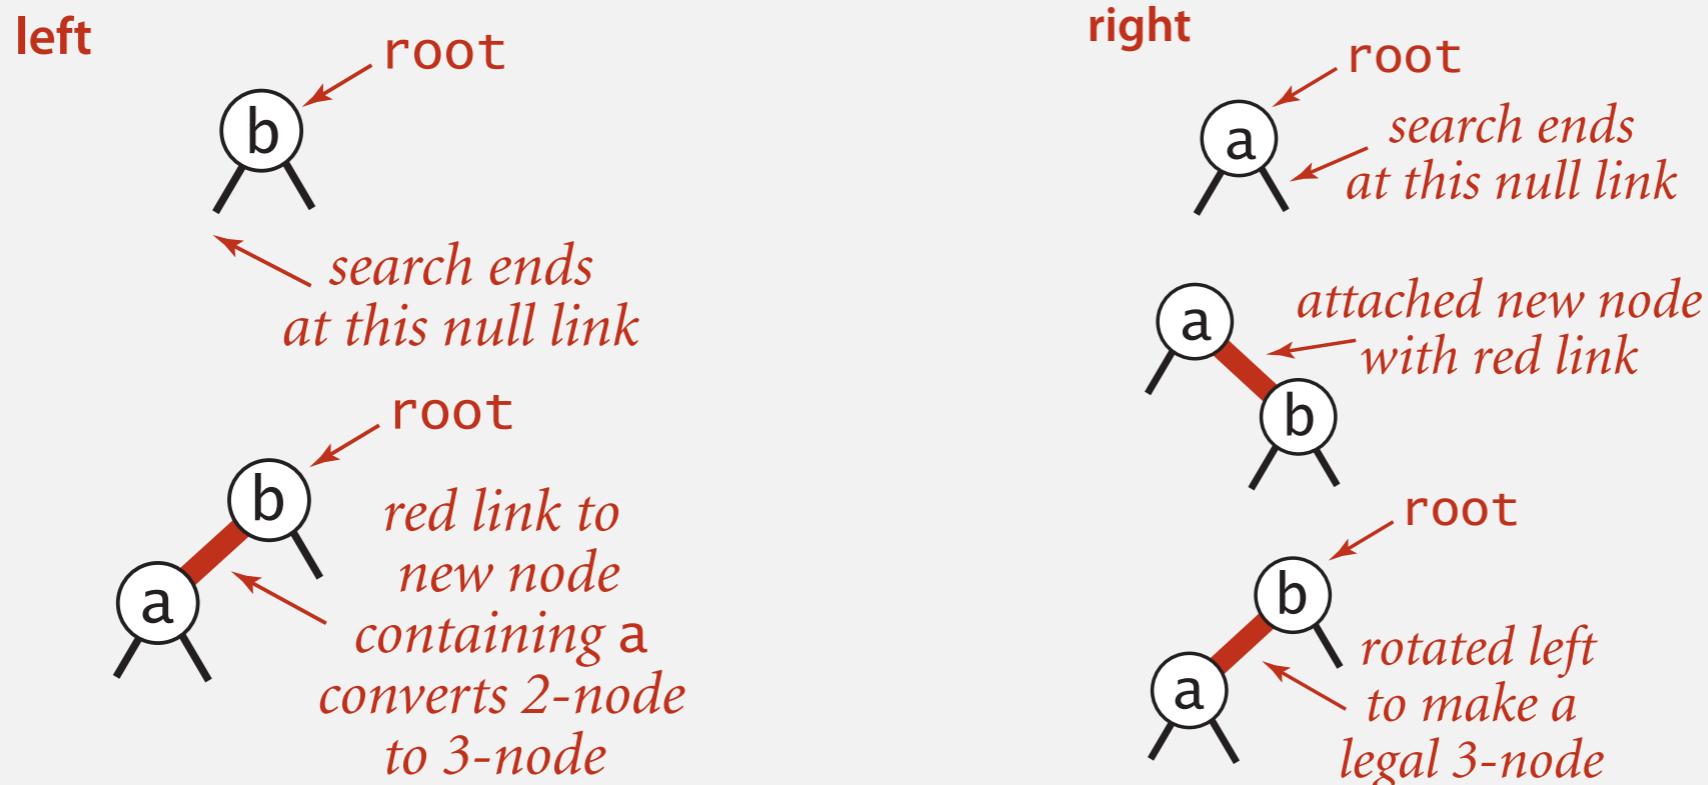
```
private void flipColors(Node h)
{
    assert !isRed(h);
    assert isRed(h.left);
    assert isRed(h.right);
    h.color = RED;
    h.left.color = BLACK;
    h.right.color = BLACK;
}
```

Invariants. Maintains symmetric order and perfect black balance.

# Insertion into a LLRB tree

---

Warmup 1. Insert into a tree with exactly 1 node.



# Insertion into a LLRB tree

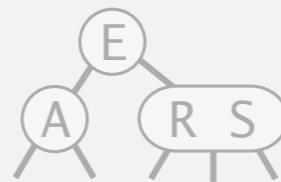
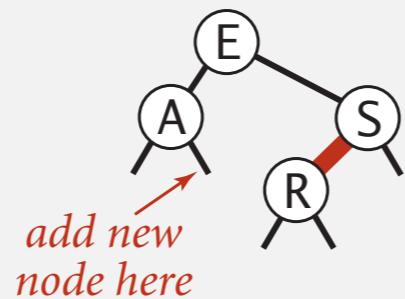
## Case 1. Insert into a 2-node at the bottom.

- Do standard BST insert; color new link red. ←
- If new red link is a right link, rotate left. ←

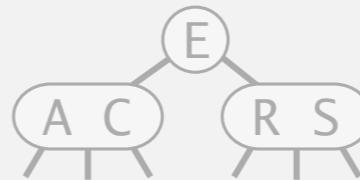
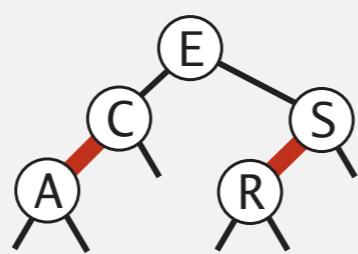
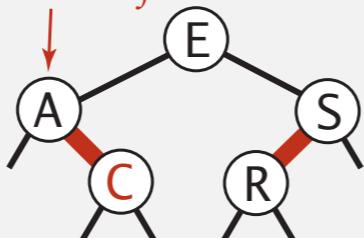
to maintain symmetric order  
and perfect black balance

to fix color invariants

insert C

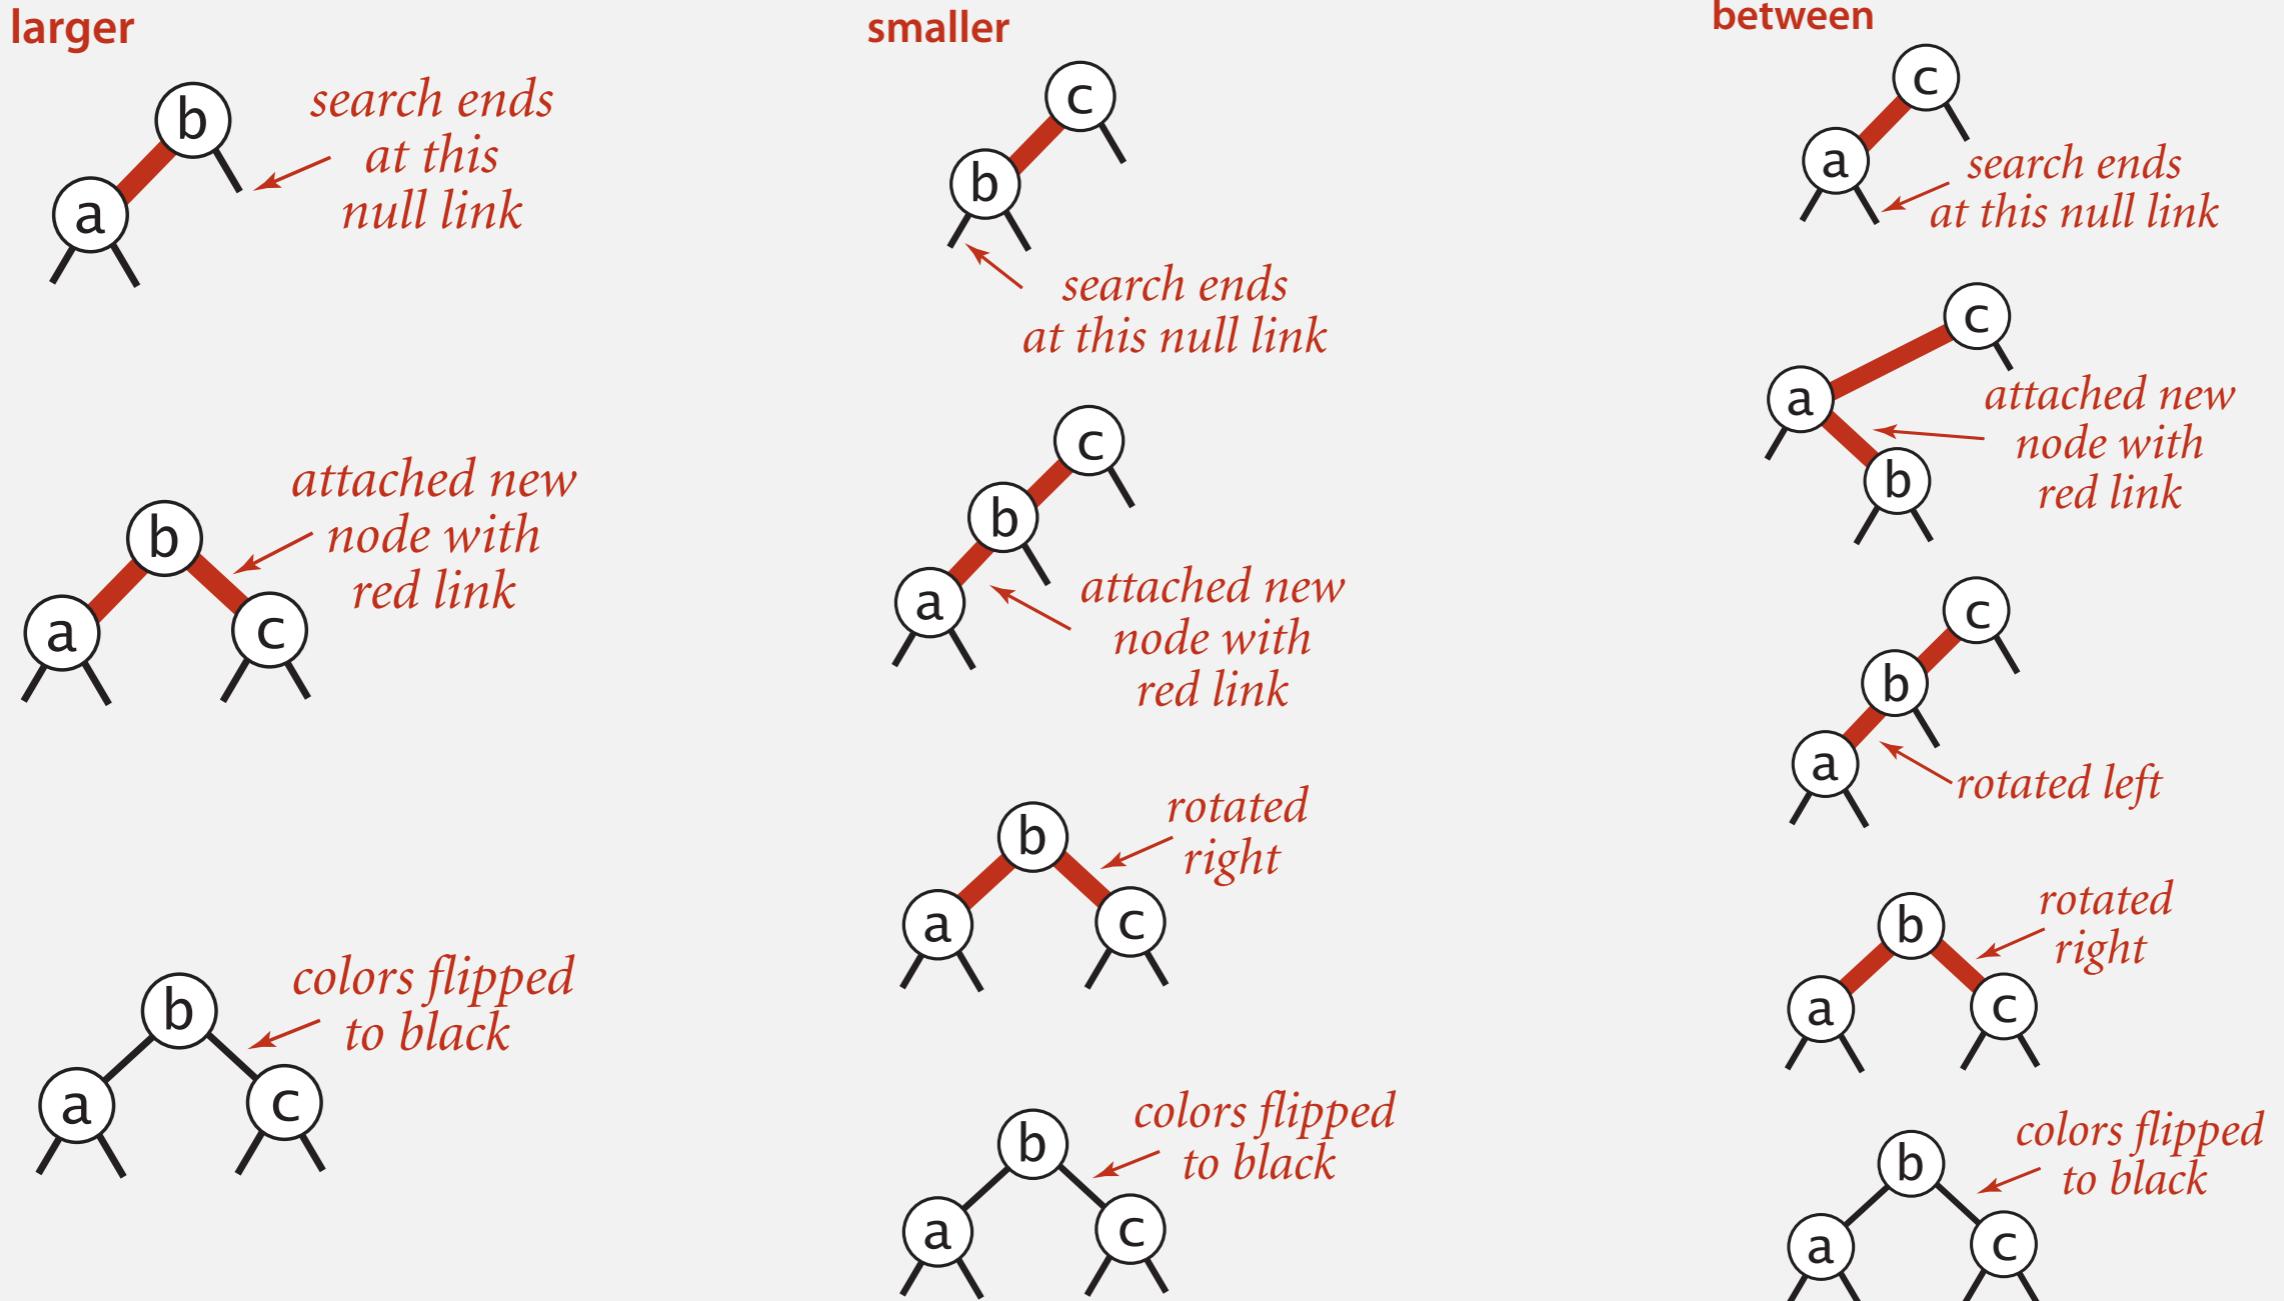


right link red  
so rotate left



# Insertion into a LLRB tree

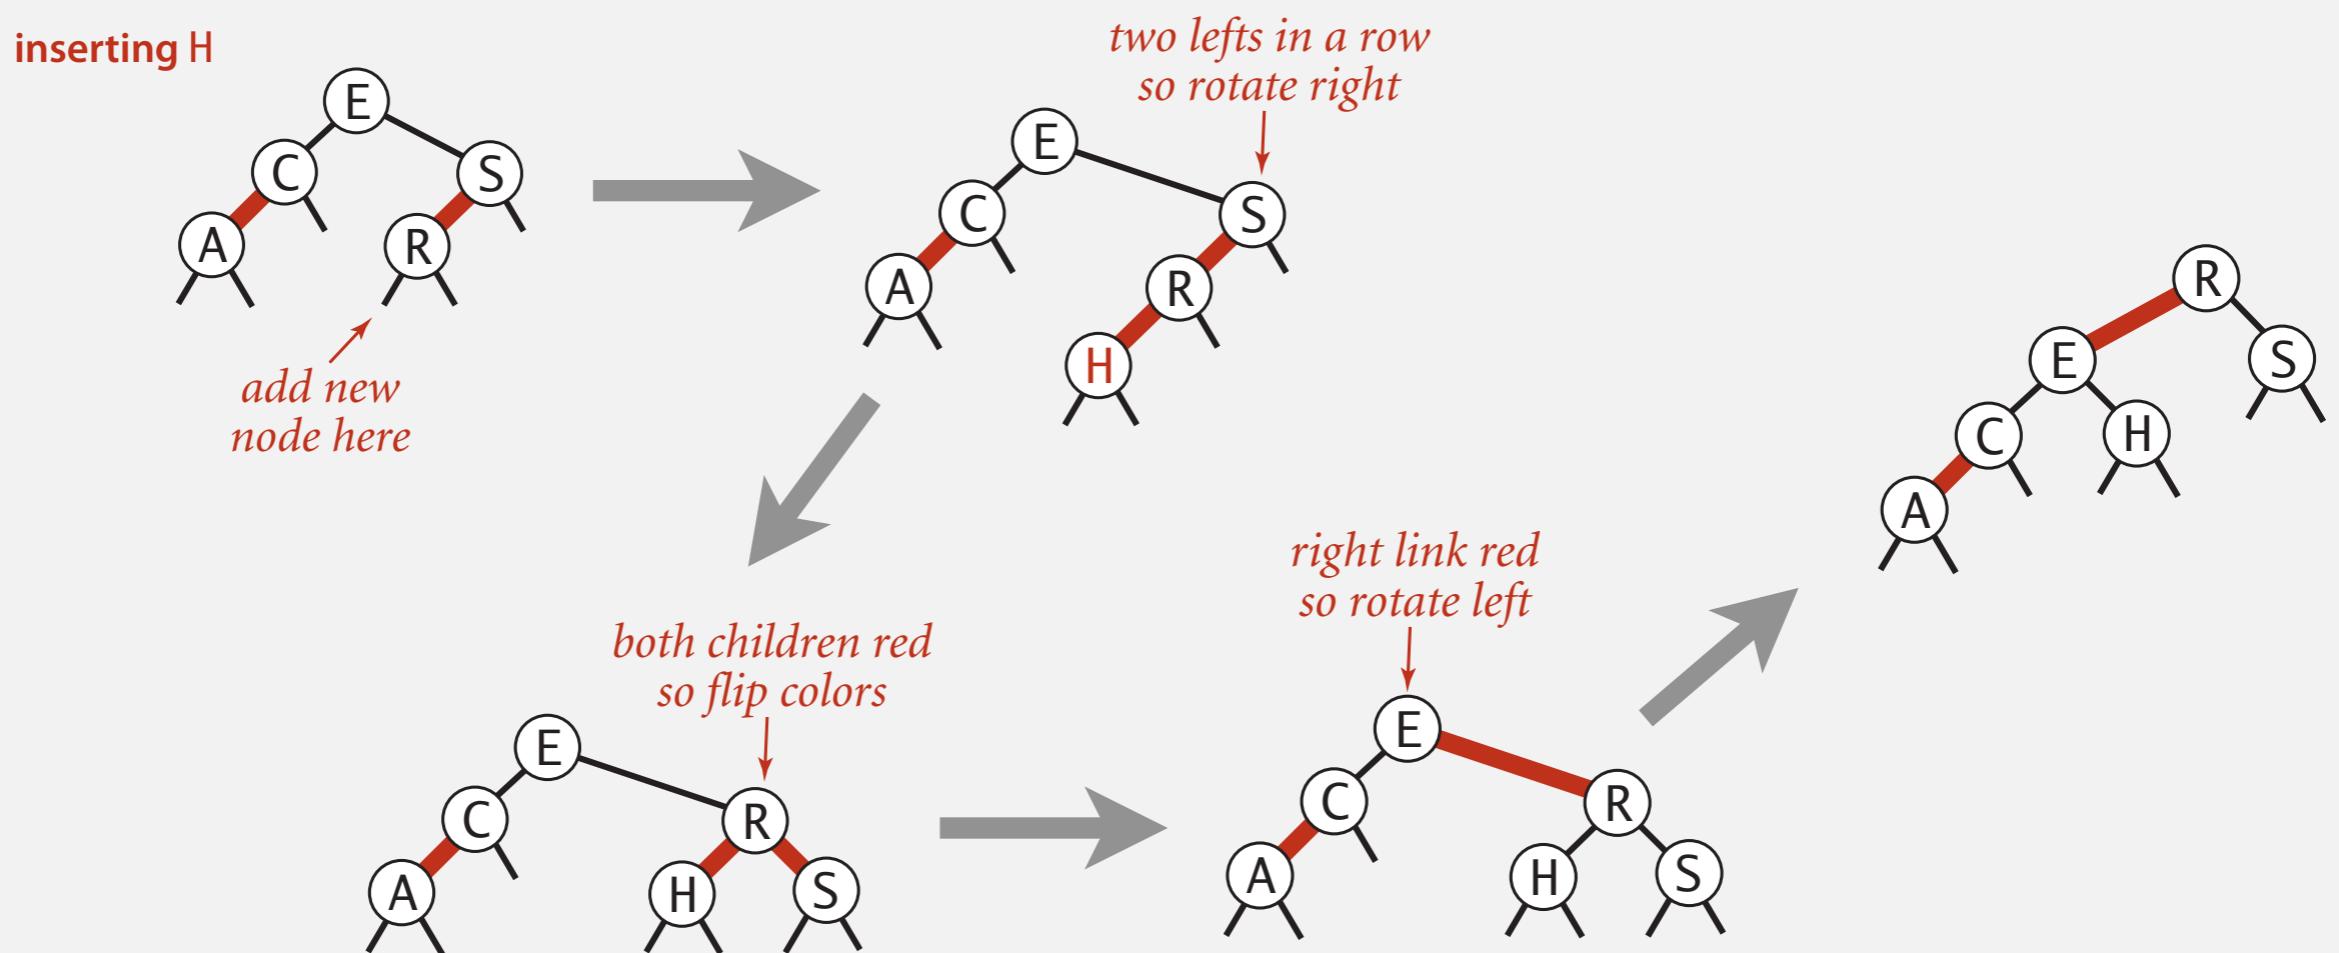
Warmup 2. Insert into a tree with exactly 2 nodes.



# Insertion into a LLRB tree

## Case 2. Insert into a 3-node at the bottom.

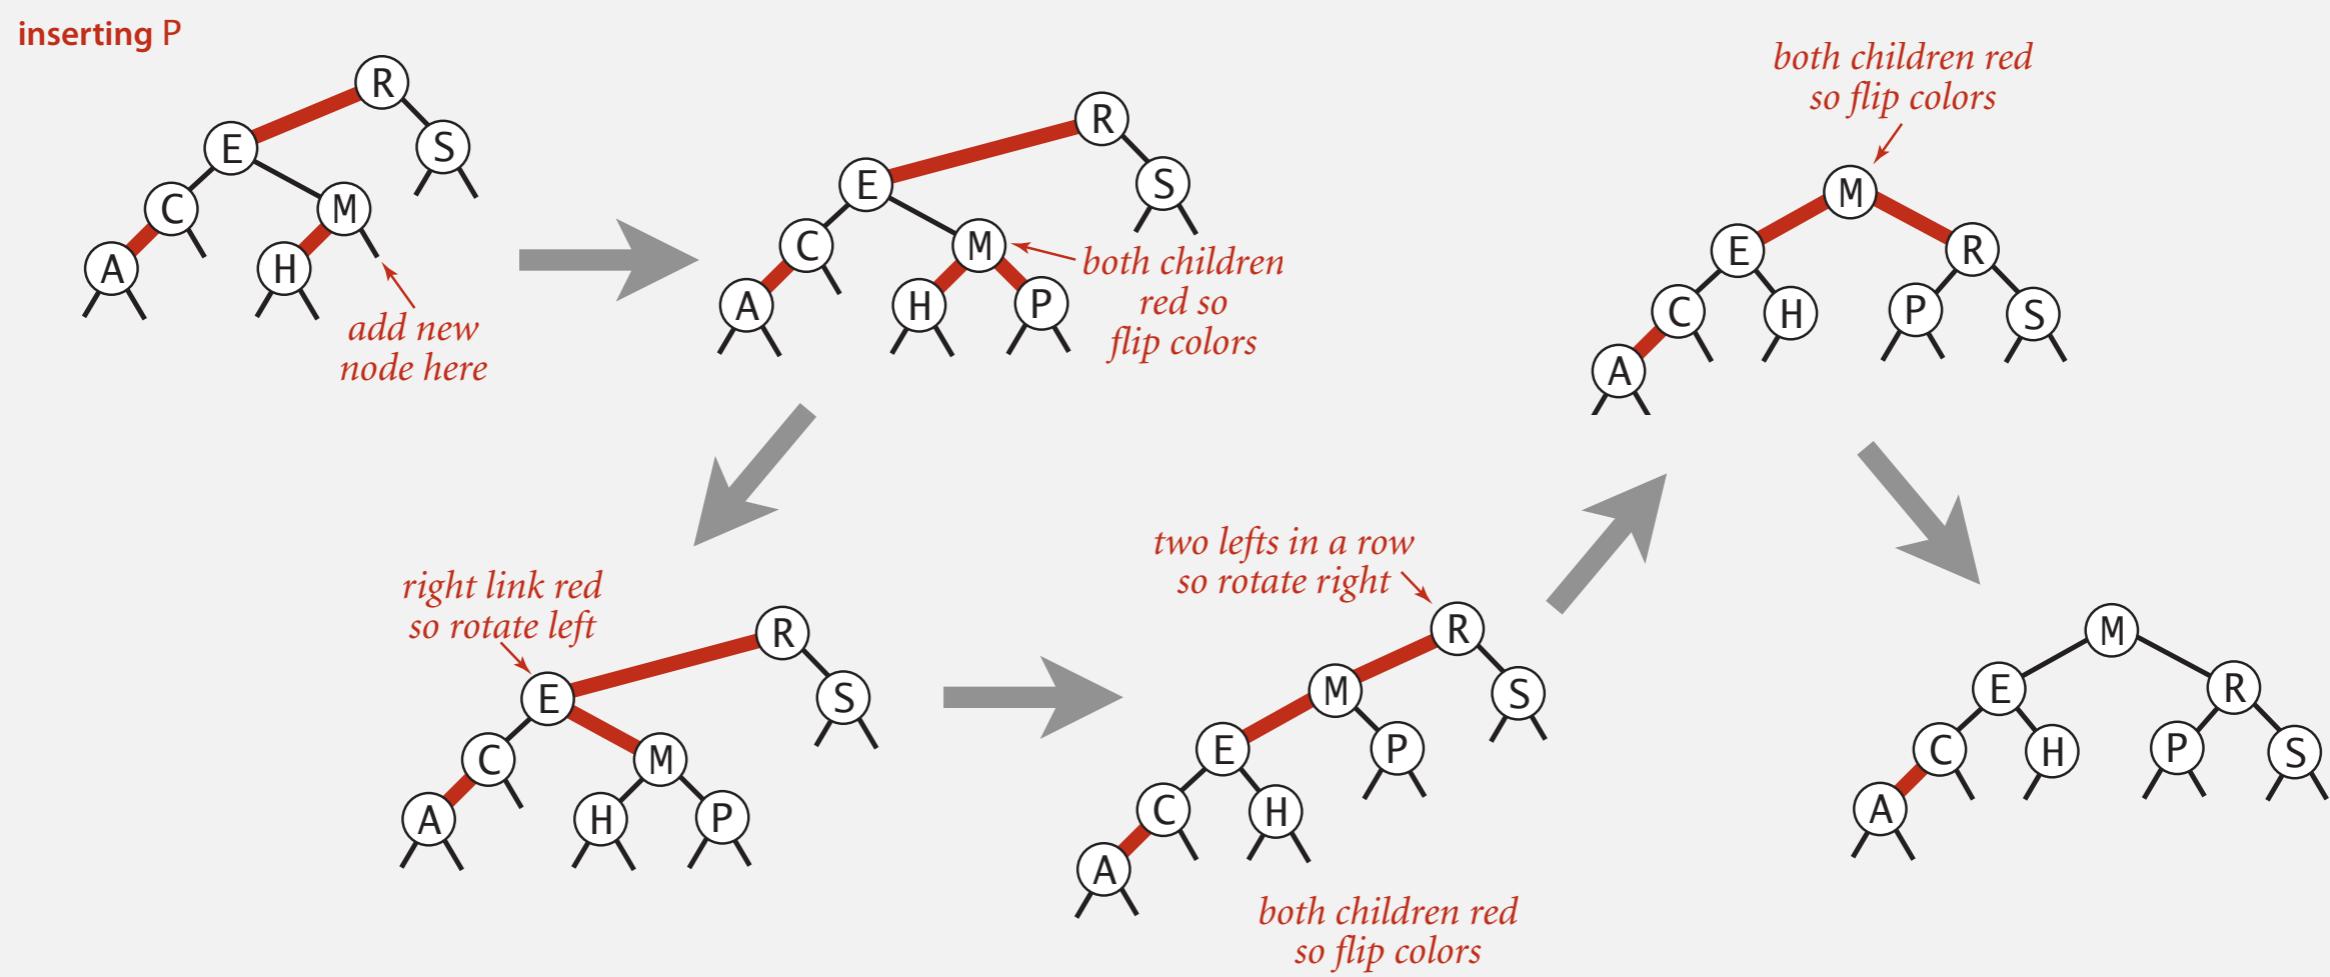
- Do standard BST insert; color new link red. ← to maintain symmetric order and perfect black balance
- Rotate to balance the 4-node (if needed).
- Flip colors to pass red link up one level. ← to fix color invariants
- Rotate to make lean left (if needed).



# Insertion into a LLRB tree: passing red links up the tree

## Case 2. Insert into a 3-node at the bottom.

- Do standard BST insert; color new link red. ← to maintain symmetric order and perfect black balance
- Rotate to balance the 4-node (if needed).
- Flip colors to pass red link up one level.
- Rotate to make lean left (if needed).
- Repeat case 1 or case 2 up the tree (if needed). ← to fix color invariants



# Red-black BST construction demo

---

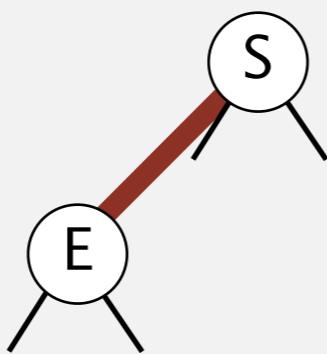
insert S



# Red-black BST construction demo

---

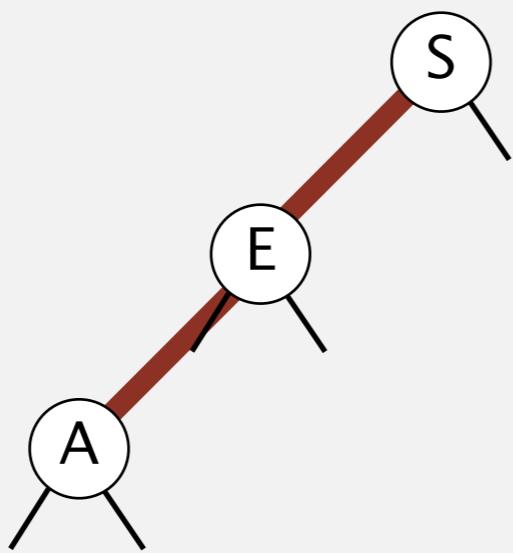
insert E



# Red-black BST construction demo

---

insert A

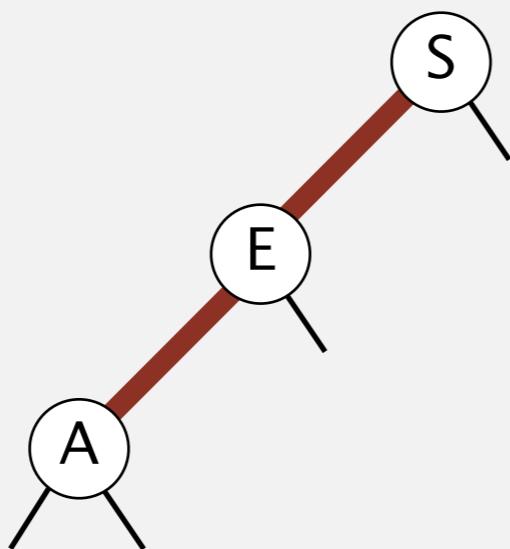


# Red-black BST construction demo

---

insert A

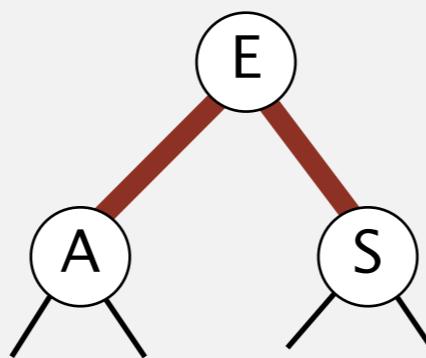
two left reds in a row  
(rotate S right)



# Red-black BST construction demo

---

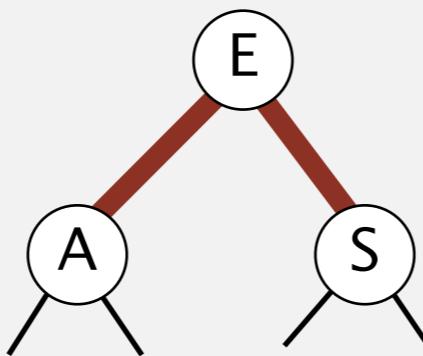
both children red  
(flip colors)



# Red-black BST construction demo

---

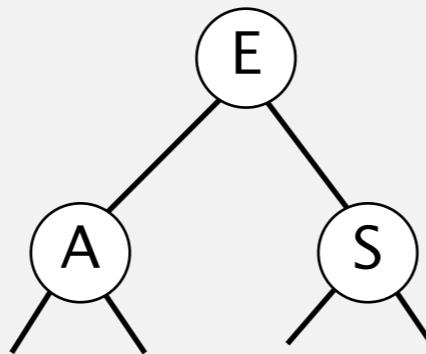
both children red  
(flip colors)



# Red-black BST construction demo

---

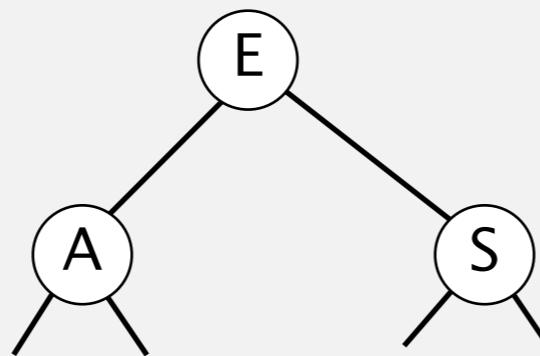
red-black BST



# Red-black BST construction demo

---

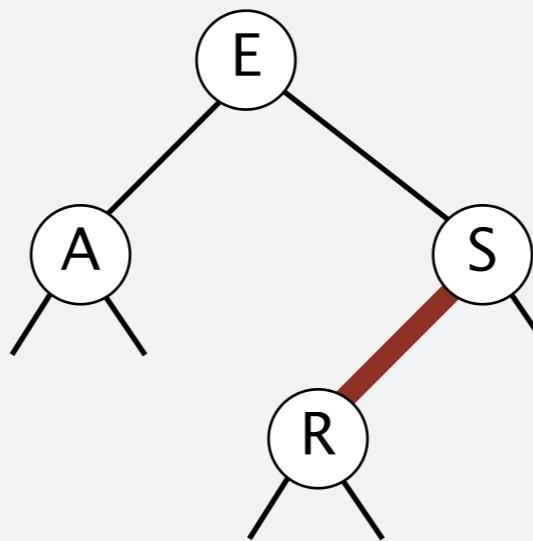
red-black BST



# Red-black BST construction demo

---

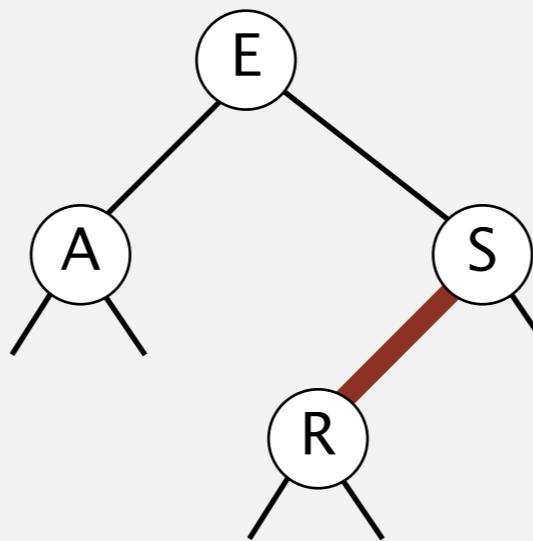
insert R



# Red-black BST construction demo

---

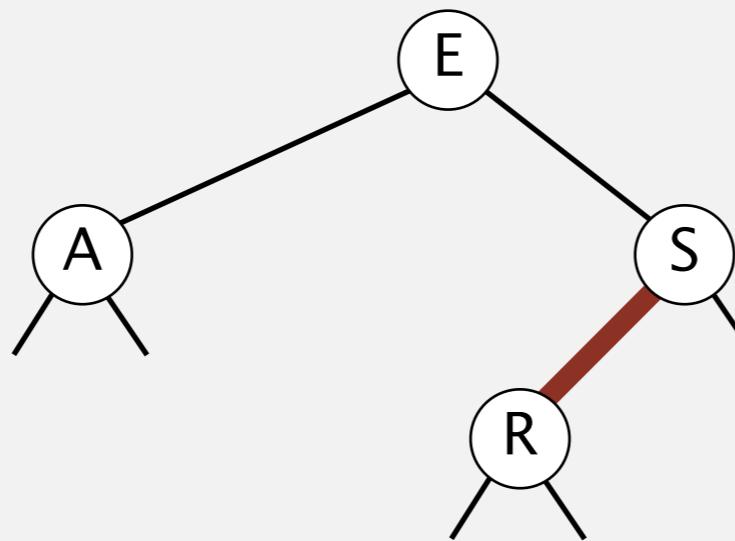
red-black BST



# Red-black BST construction demo

---

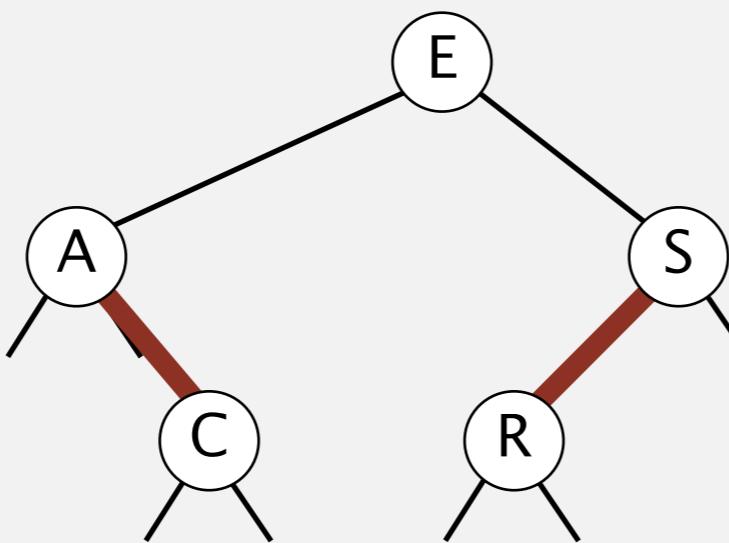
red-black BST



# Red-black BST construction demo

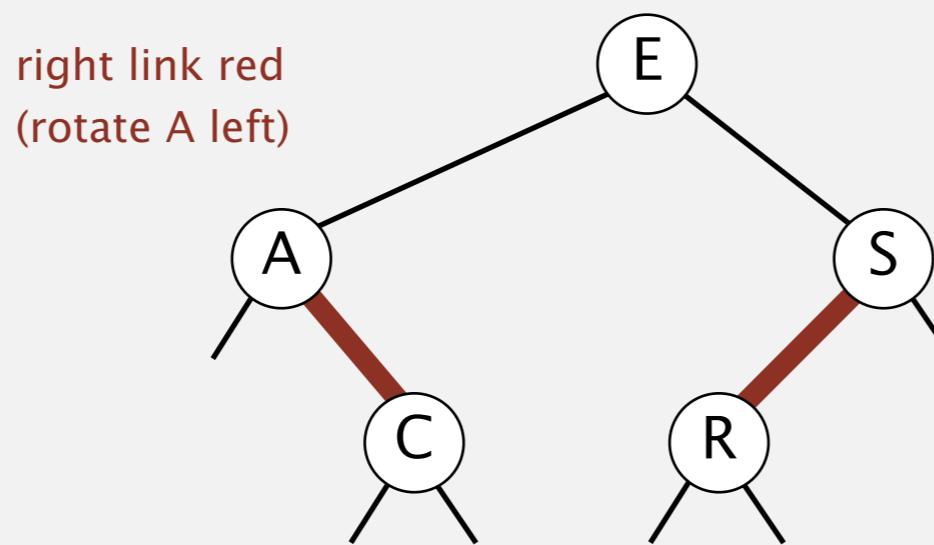
---

insert C



# Red-black BST construction demo

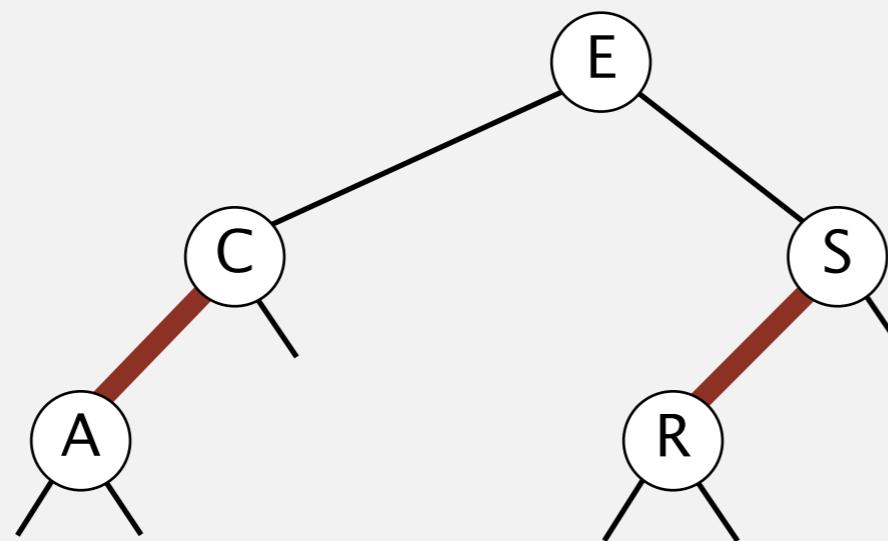
---



# Red-black BST construction demo

---

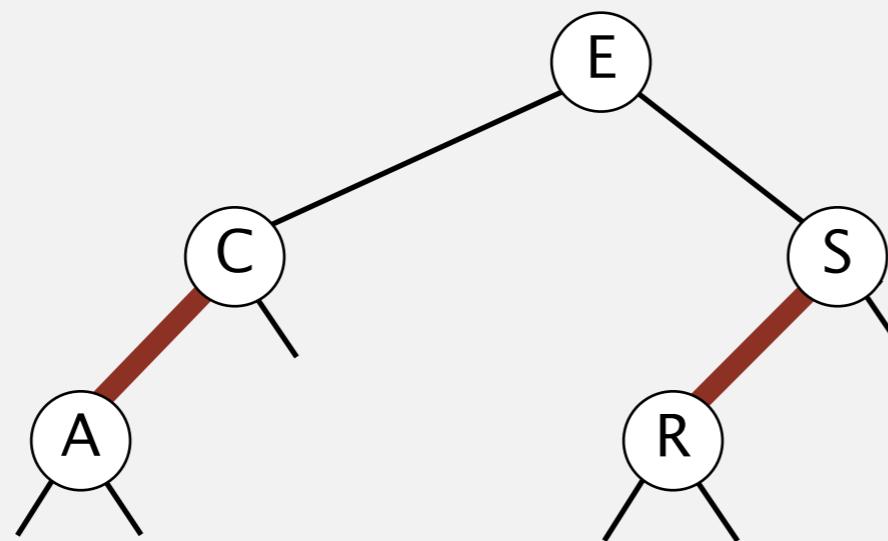
red-black BST



# Red-black BST construction demo

---

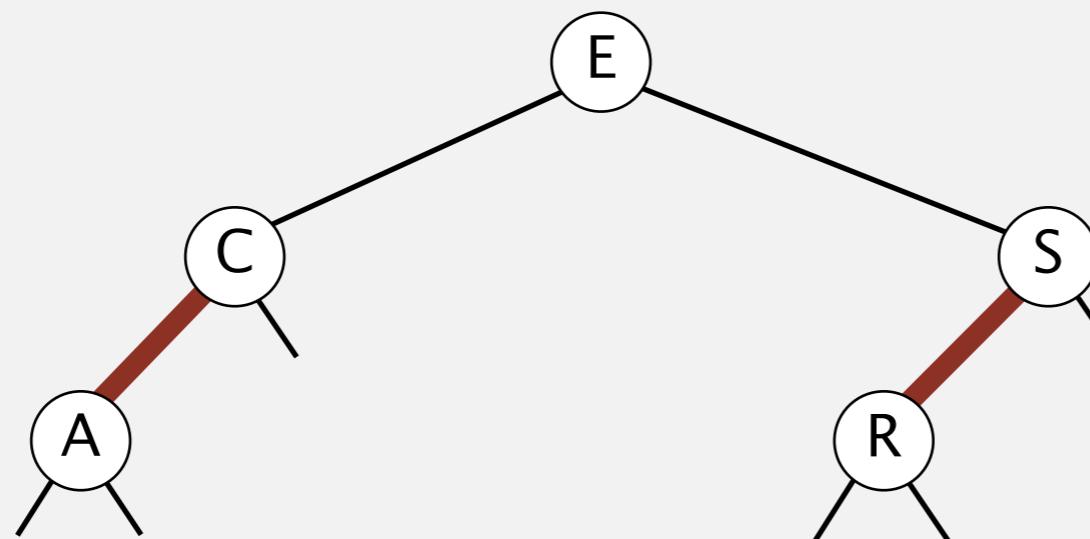
red-black BST



# Red-black BST construction demo

---

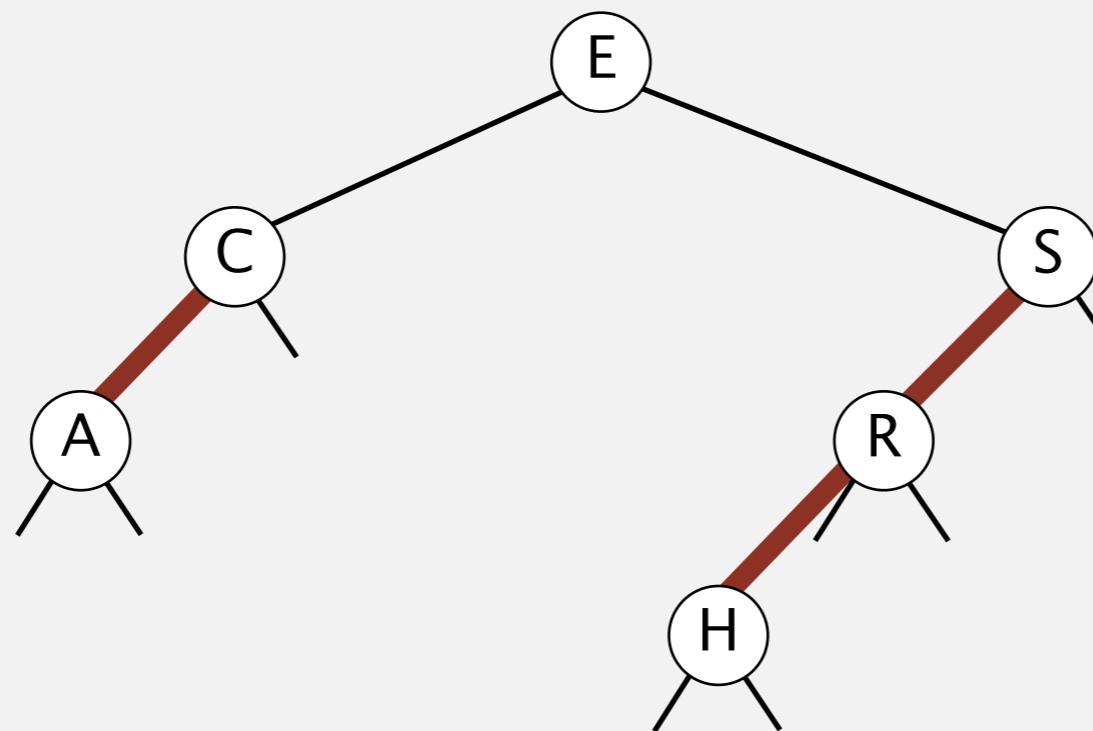
red-black BST



# Red-black BST construction demo

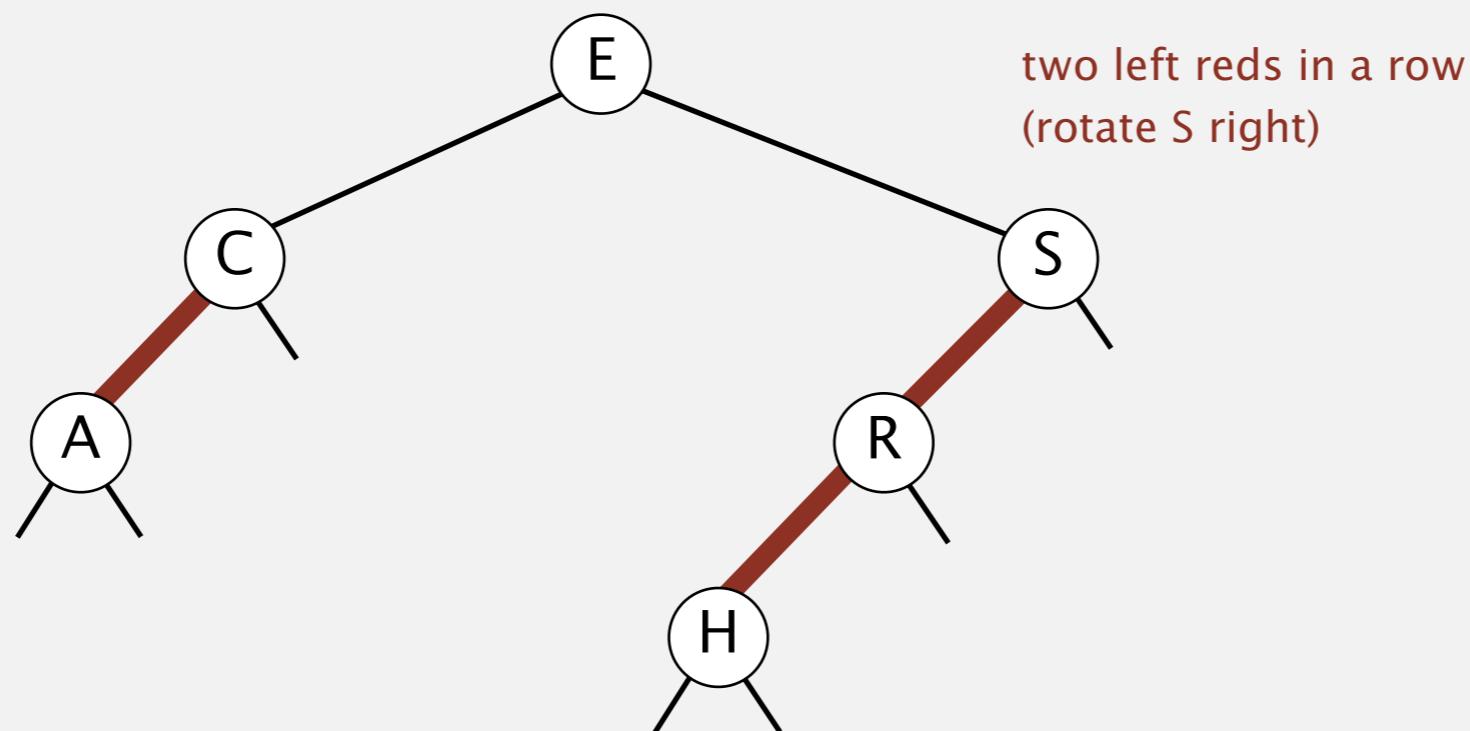
---

insert H



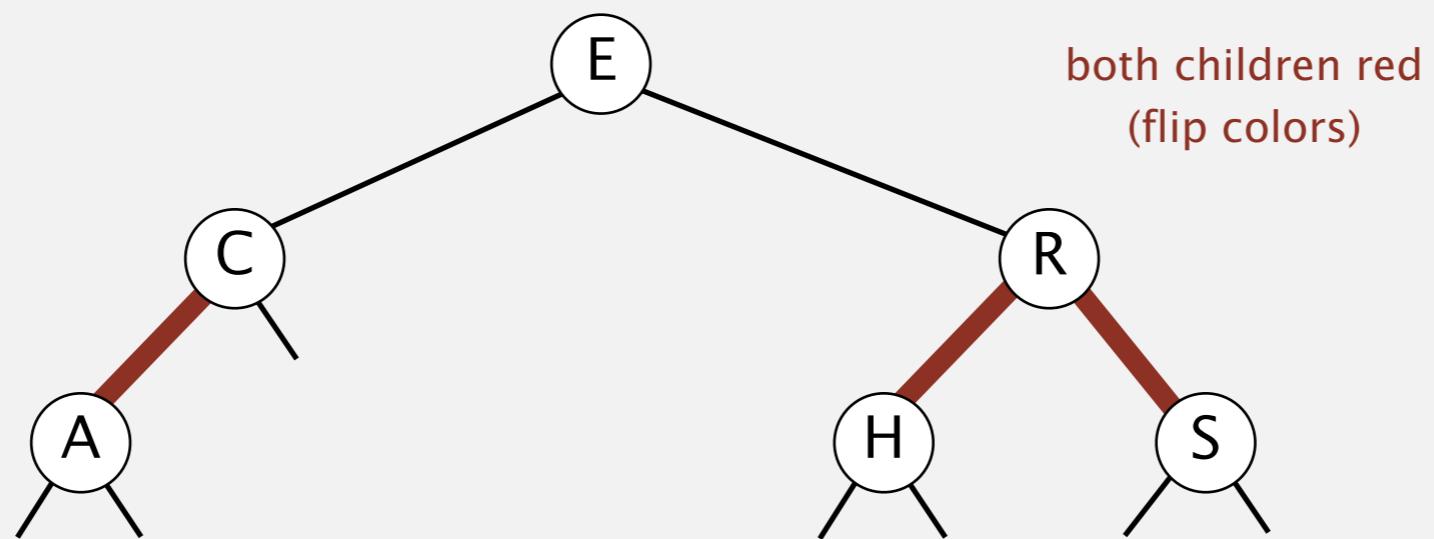
# Red-black BST construction demo

---



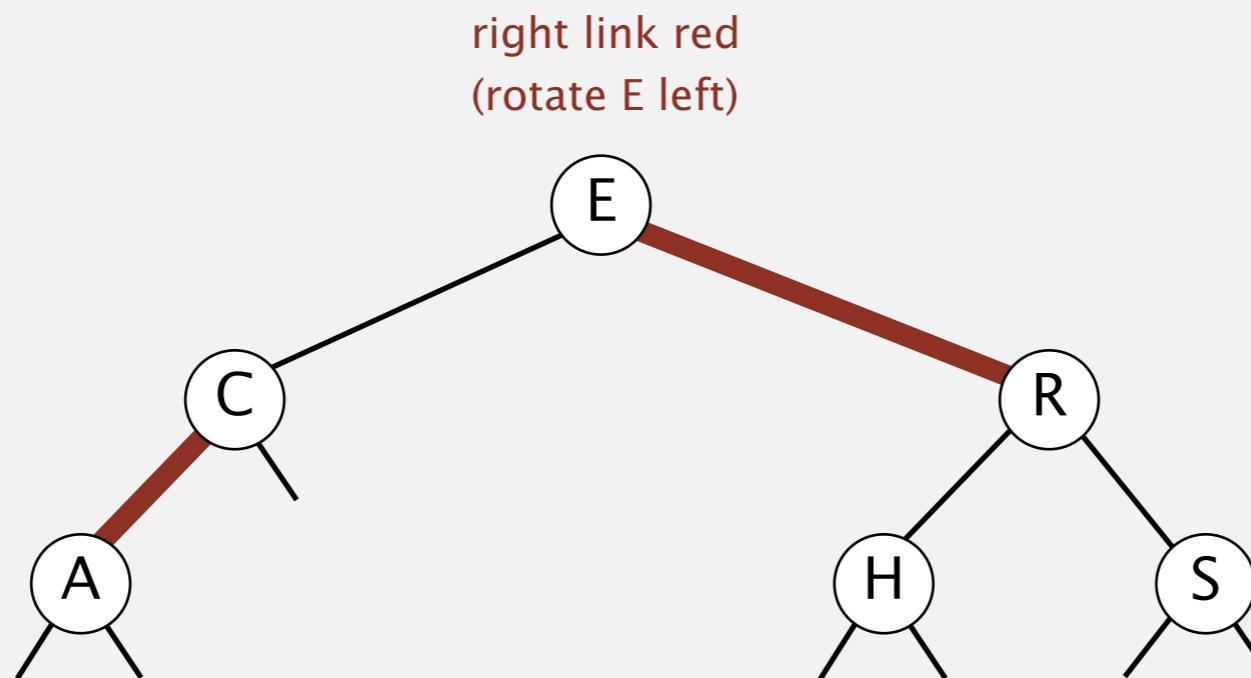
# Red-black BST construction demo

---



# Red-black BST construction demo

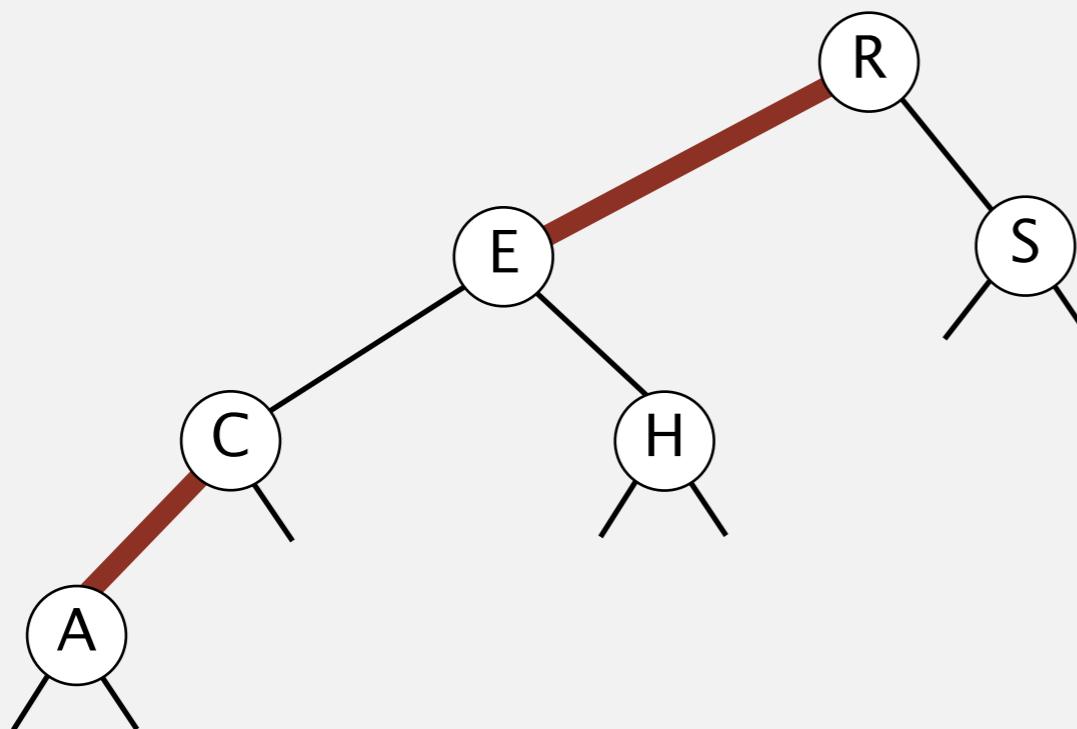
---



# Red-black BST construction demo

---

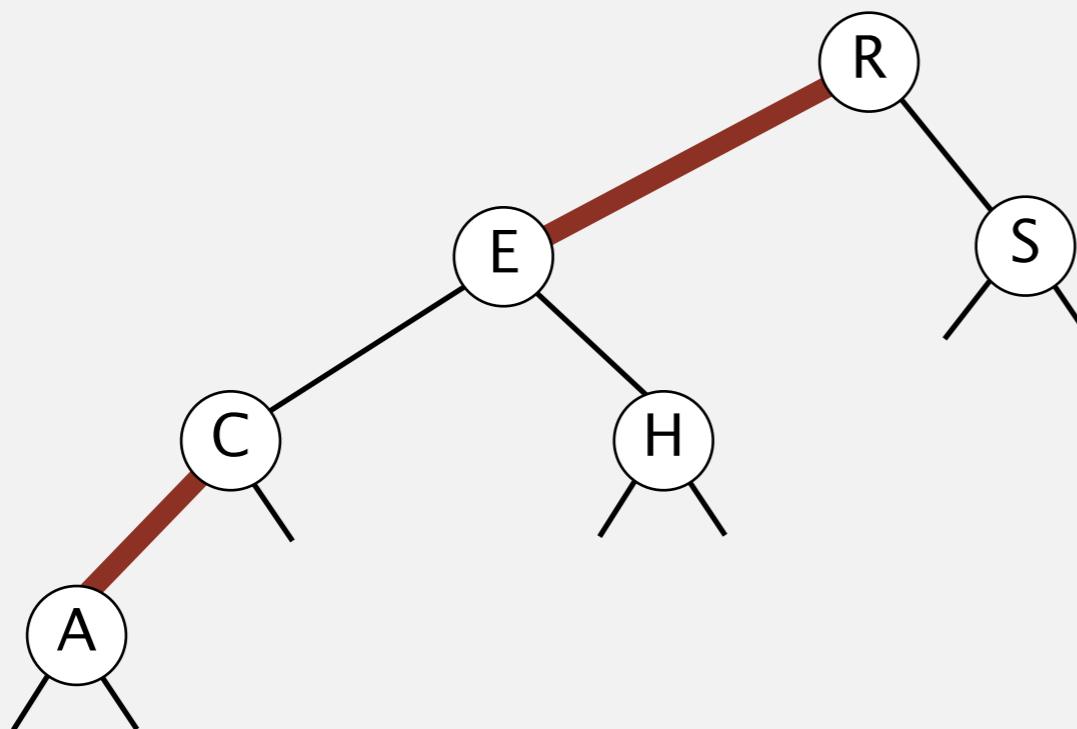
red-black BST



# Red-black BST construction demo

---

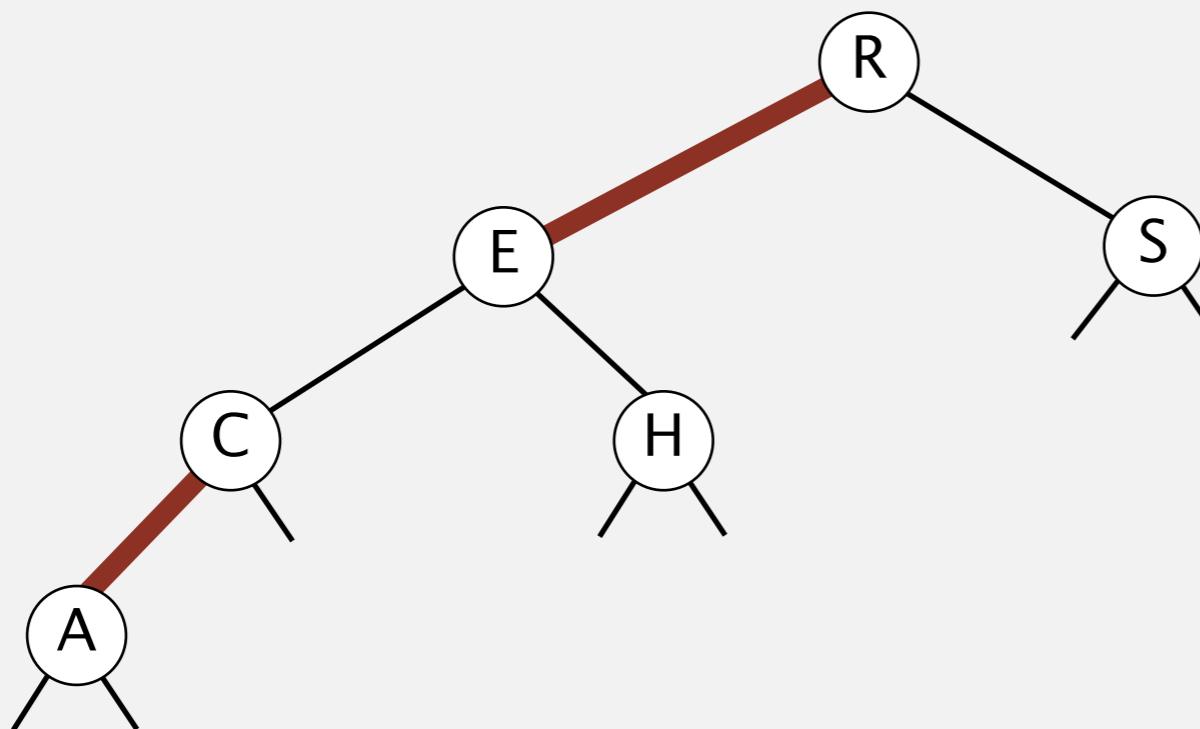
red-black BST



# Red-black BST construction demo

---

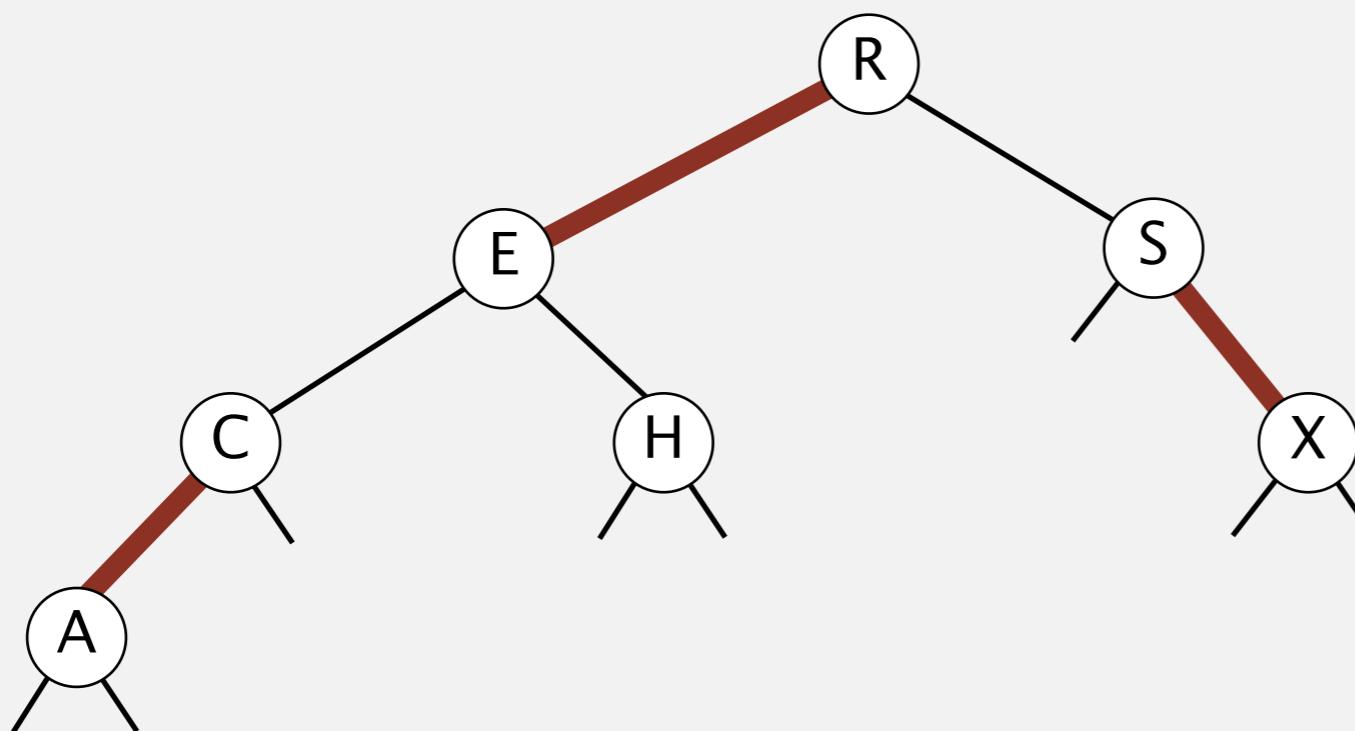
red-black BST



# Red-black BST construction demo

---

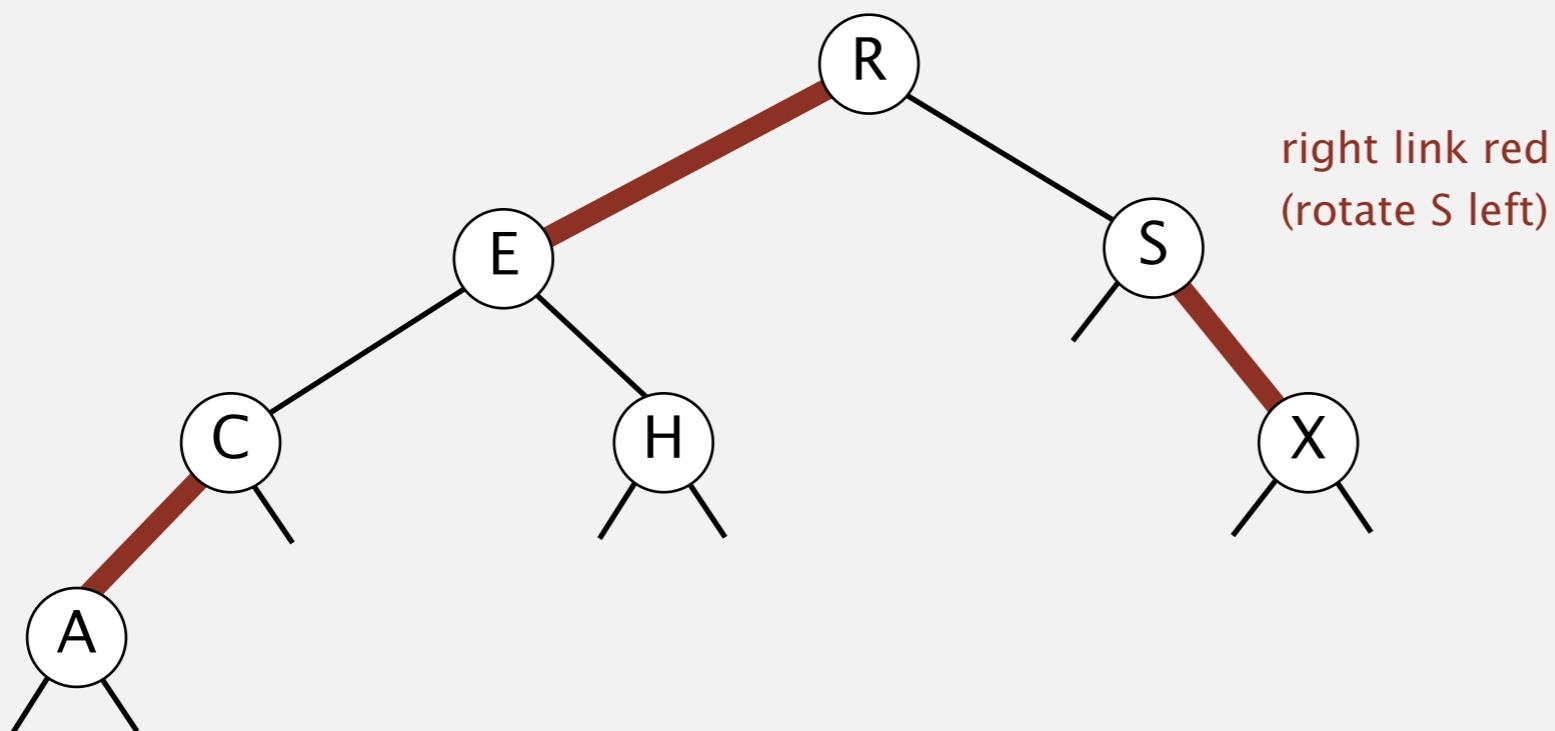
insert X



# Red-black BST construction demo

---

insert X

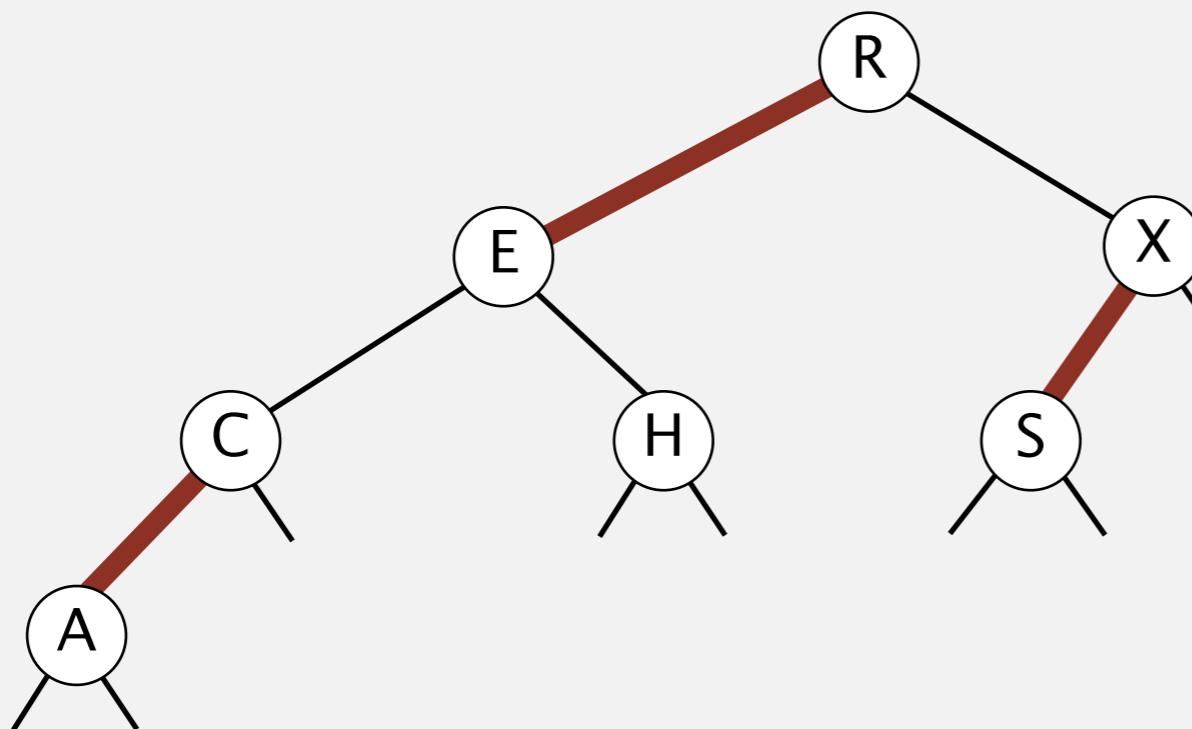


right link red  
(rotate S left)

# Red-black BST construction demo

---

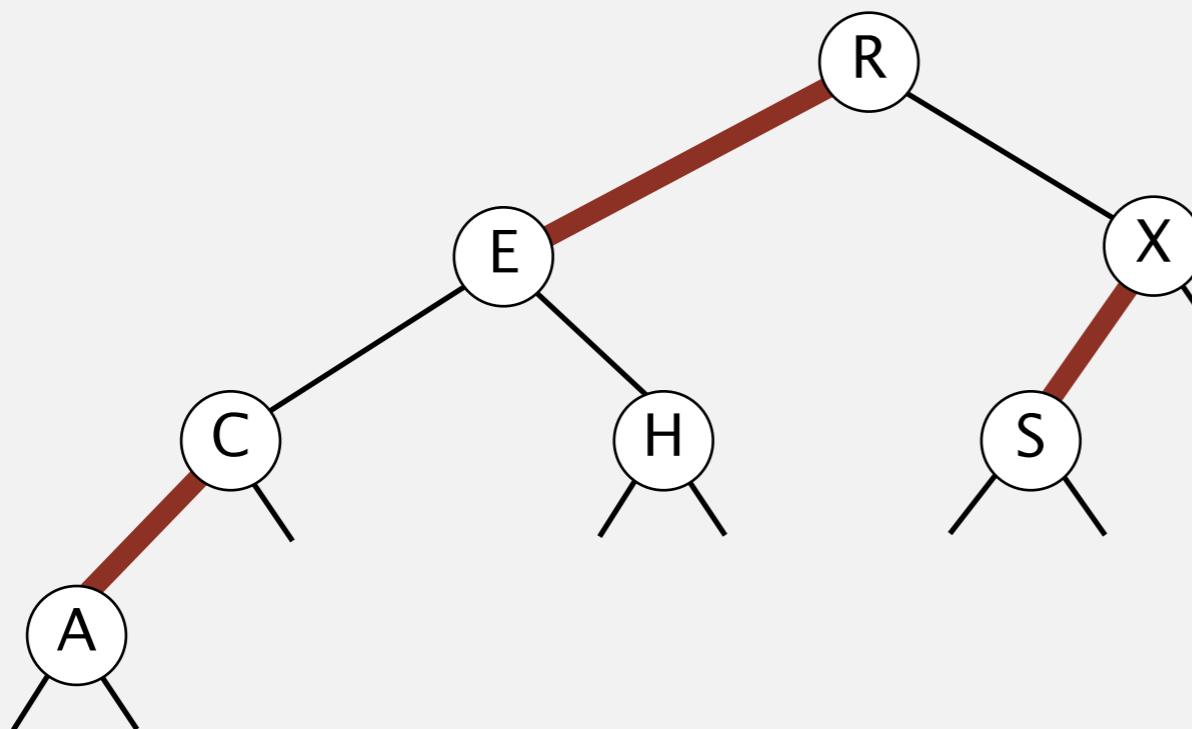
red-black BST



# Red-black BST construction demo

---

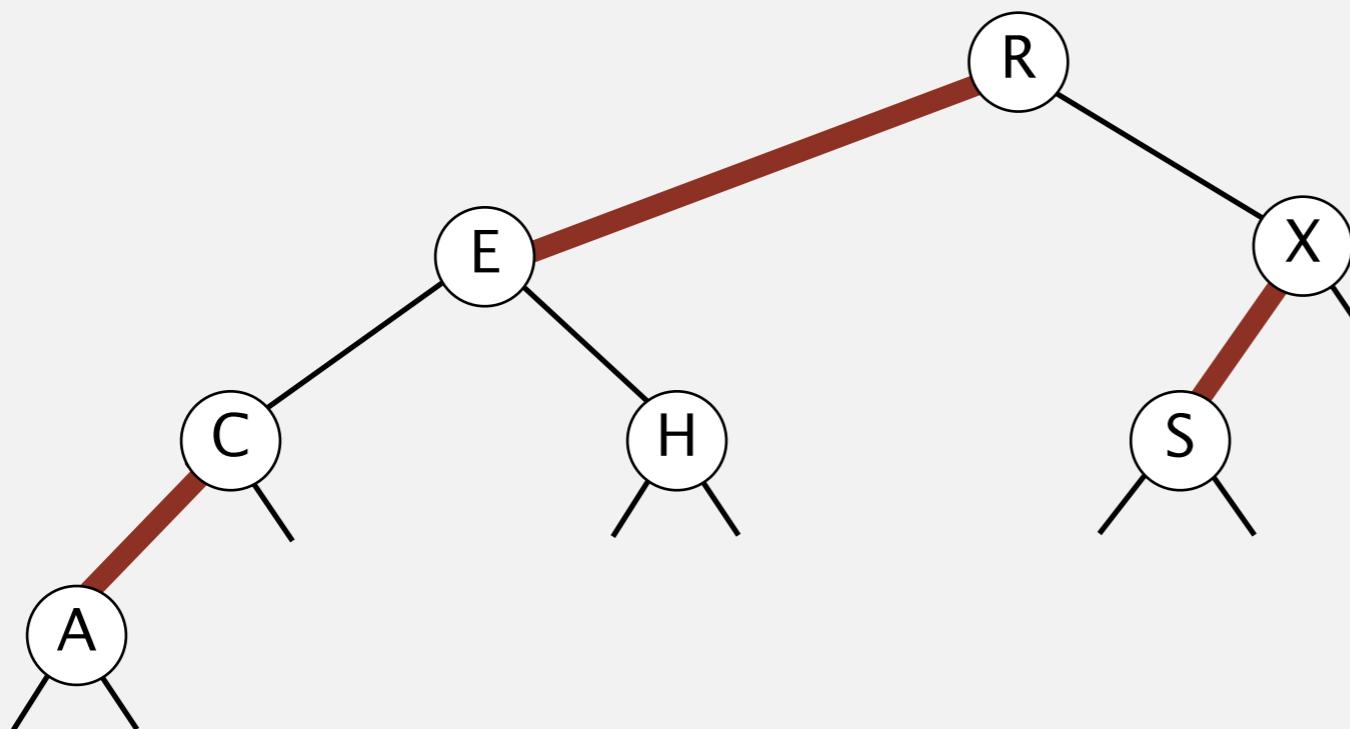
red-black BST



# Red-black BST construction demo

---

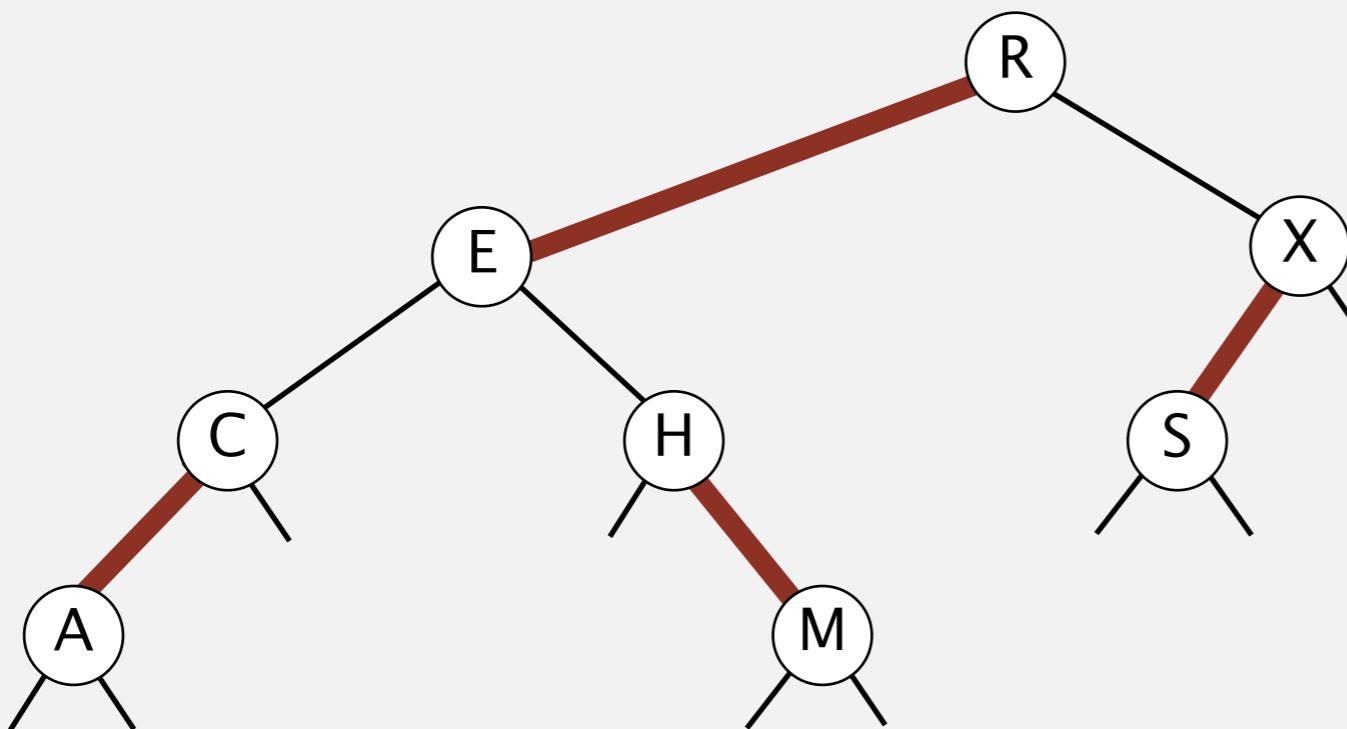
red-black BST



# Red-black BST construction demo

---

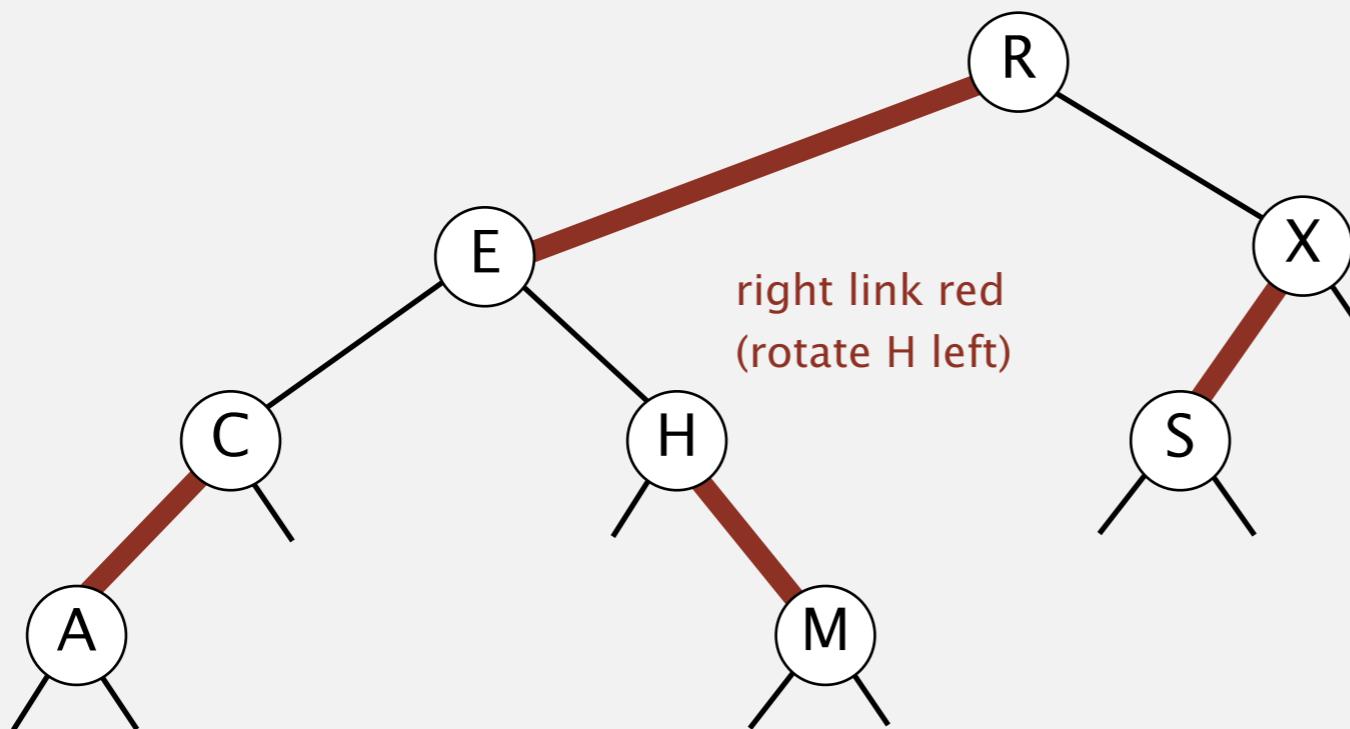
insert M



# Red-black BST construction demo

---

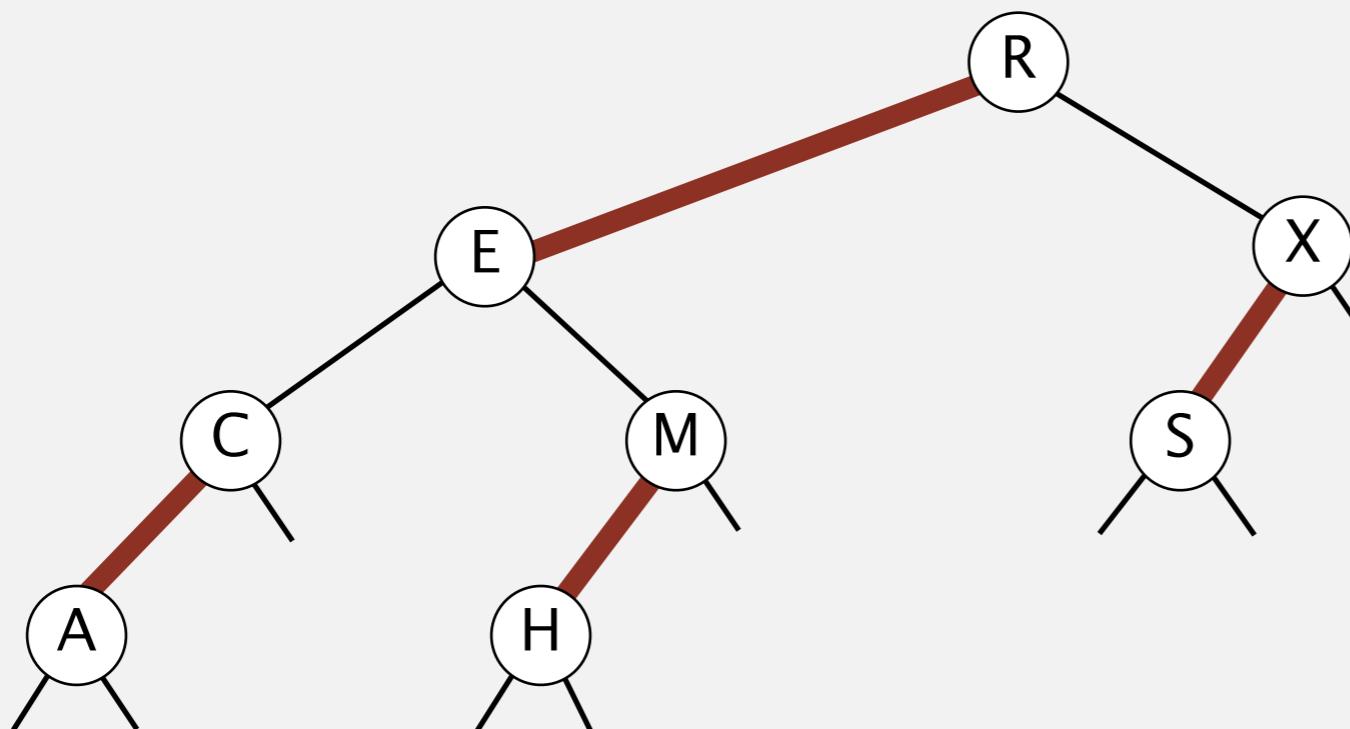
insert M



# Red-black BST construction demo

---

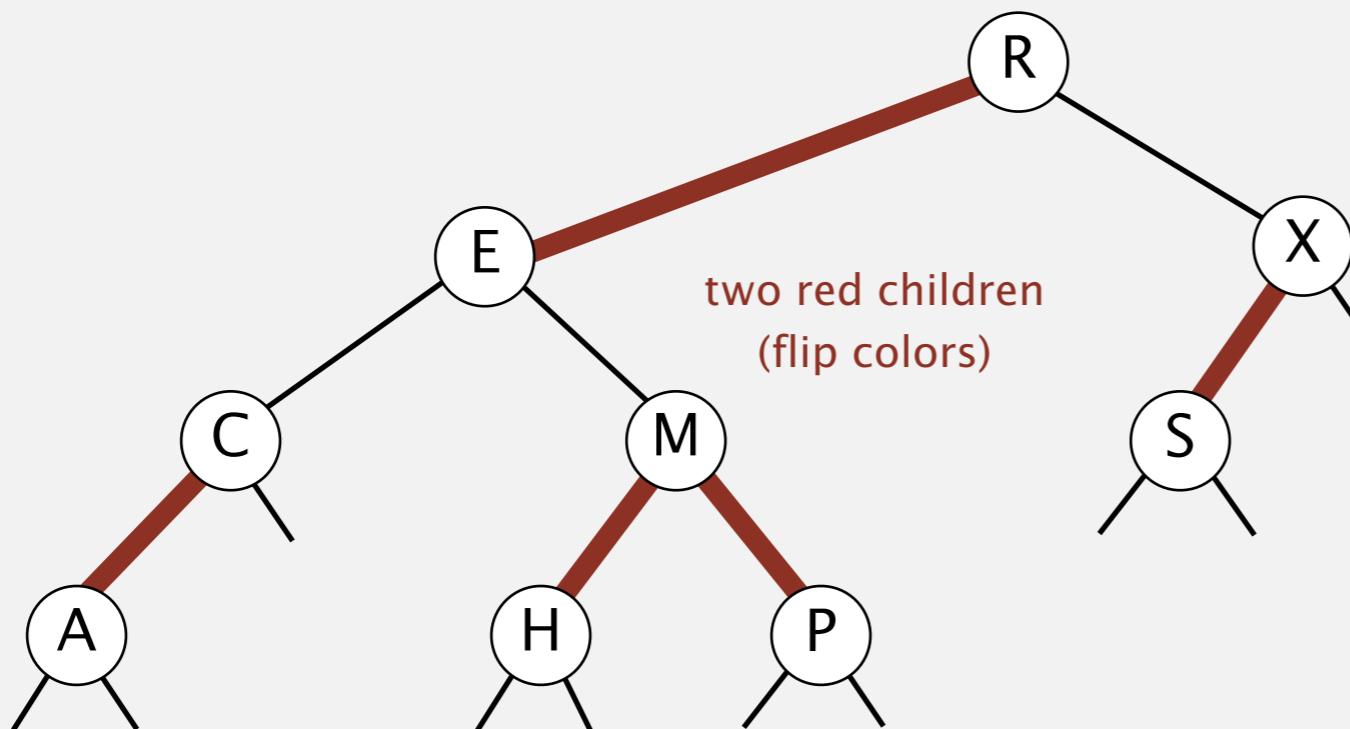
red-black BST



# Red-black BST construction demo

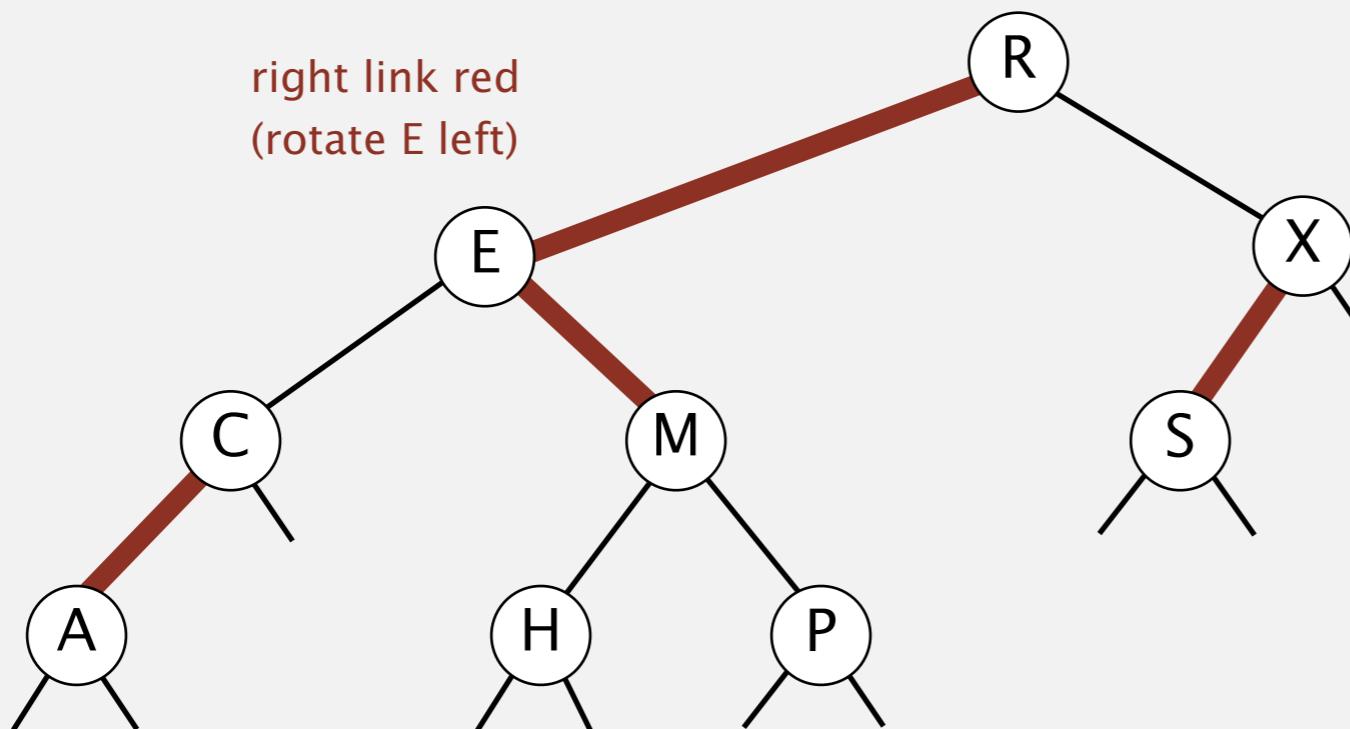
---

insert P



# Red-black BST construction demo

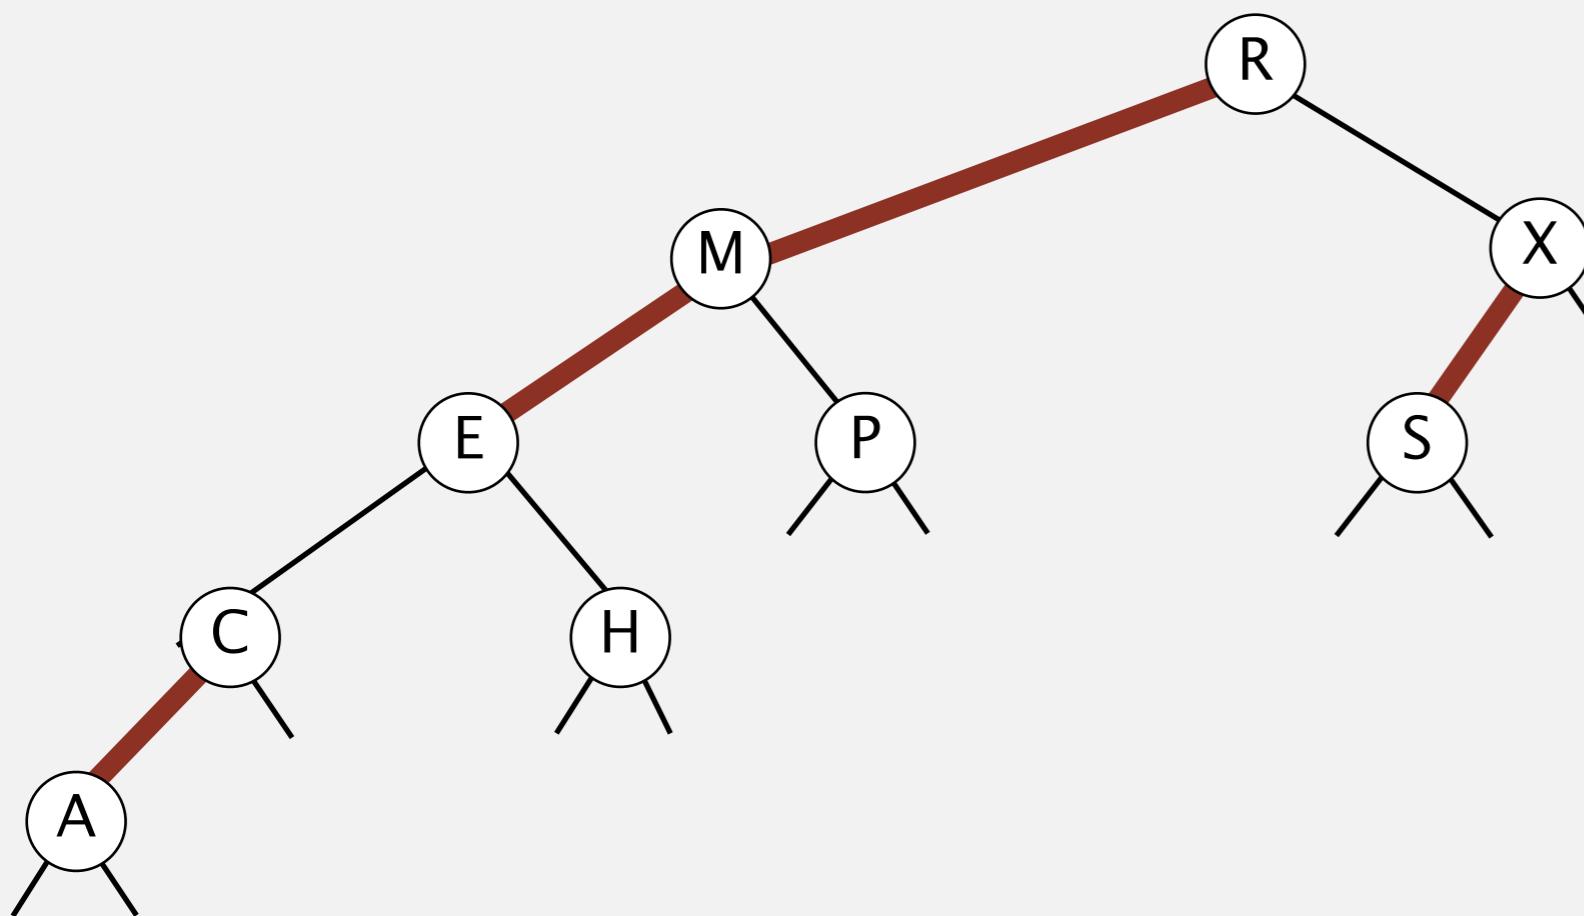
---



# Red-black BST construction demo

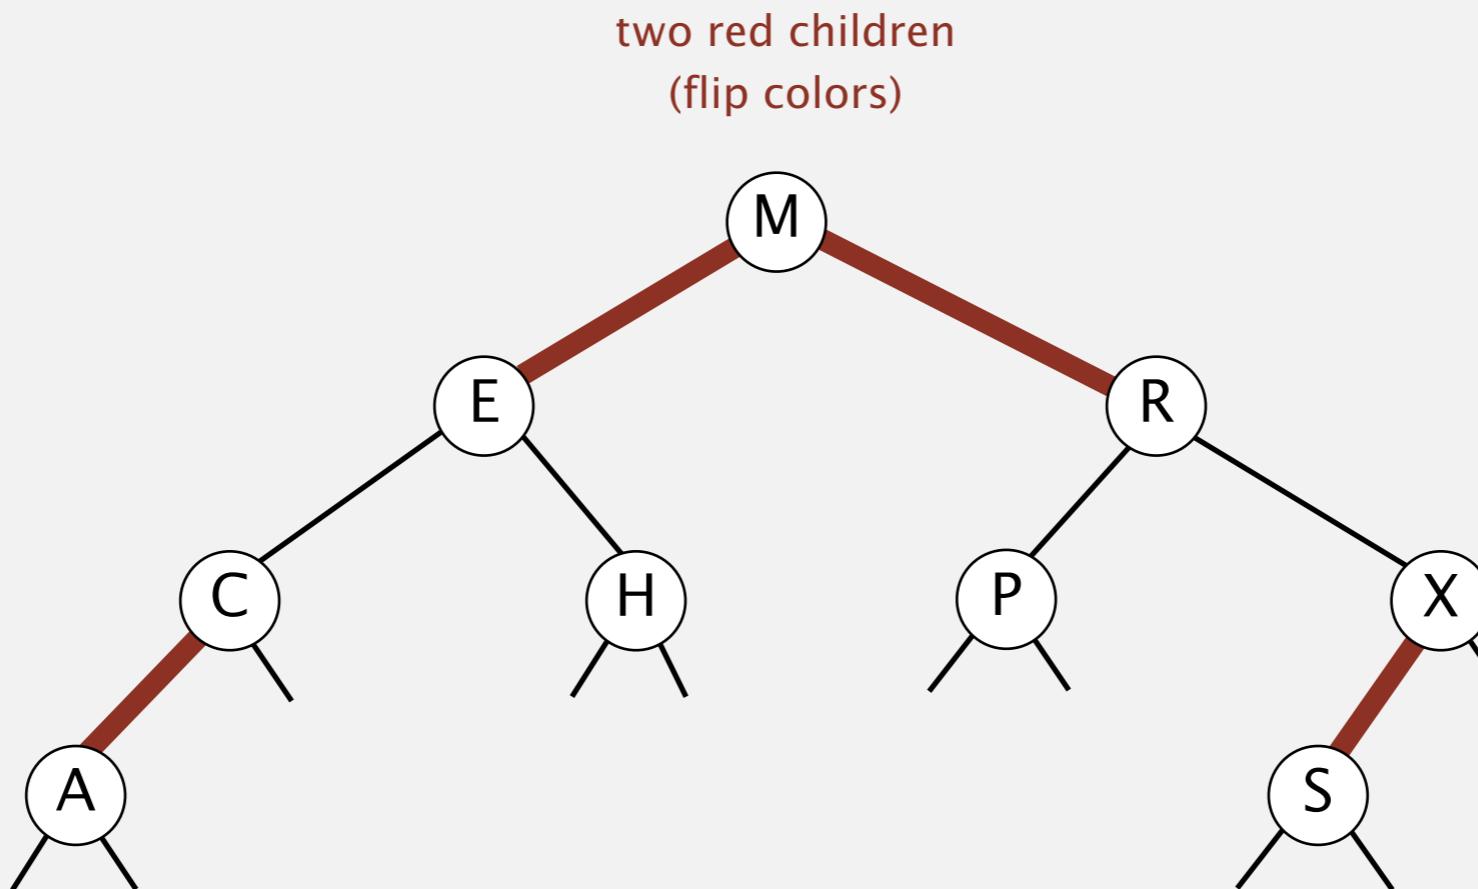
---

two left reds in a row  
(rotate R right)



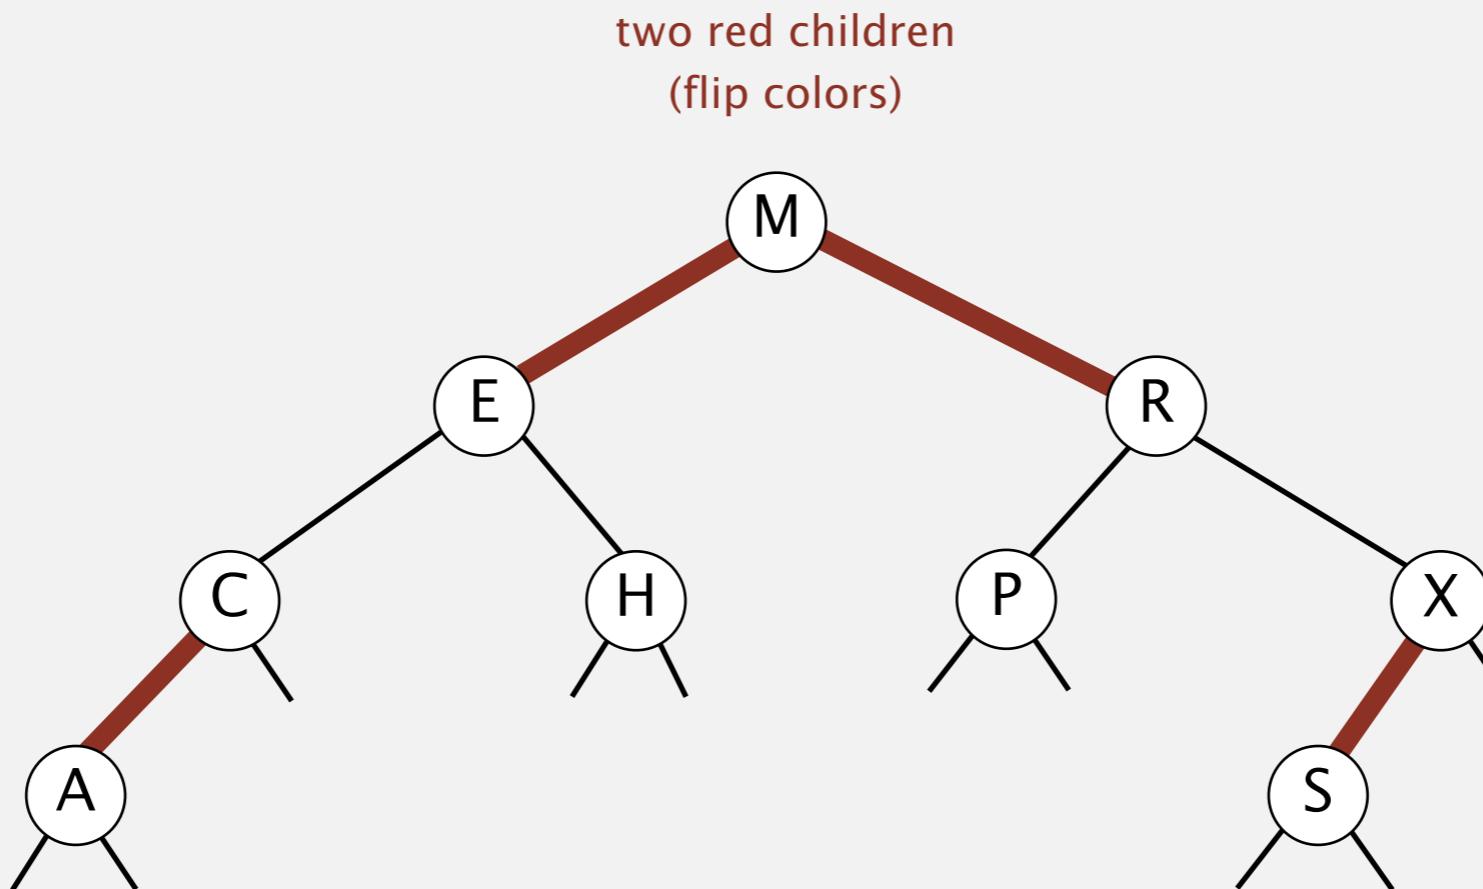
# Red-black BST construction demo

---



# Red-black BST construction demo

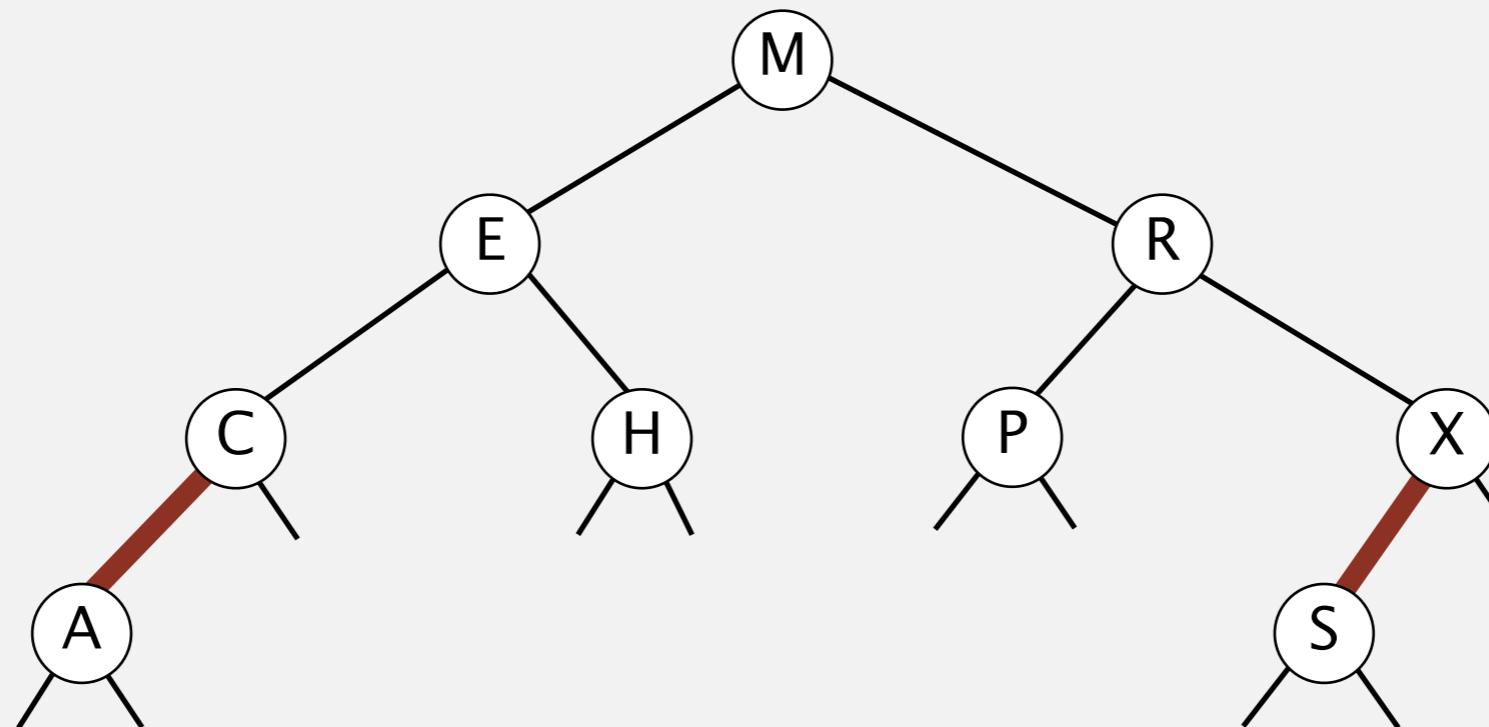
---



# Red-black BST construction demo

---

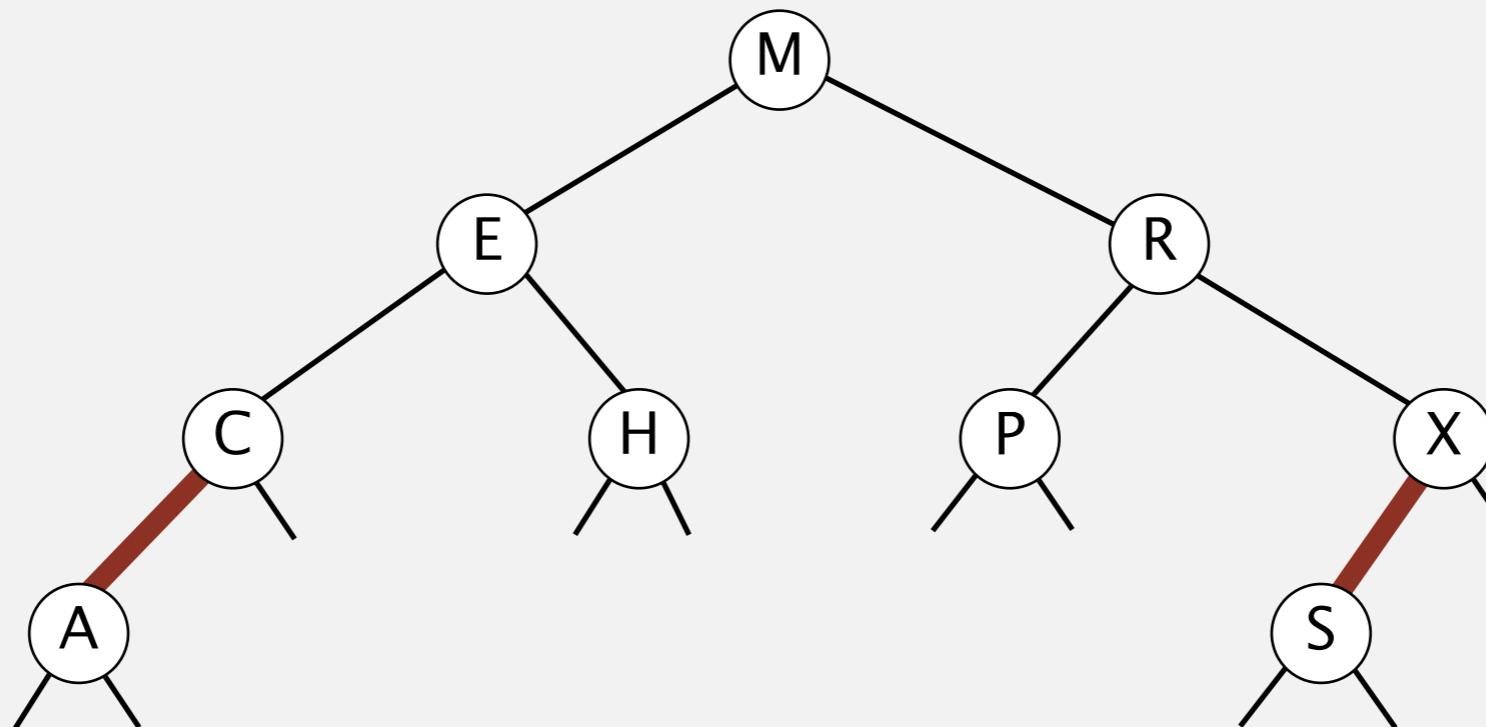
red-black BST



# Red-black BST construction demo

---

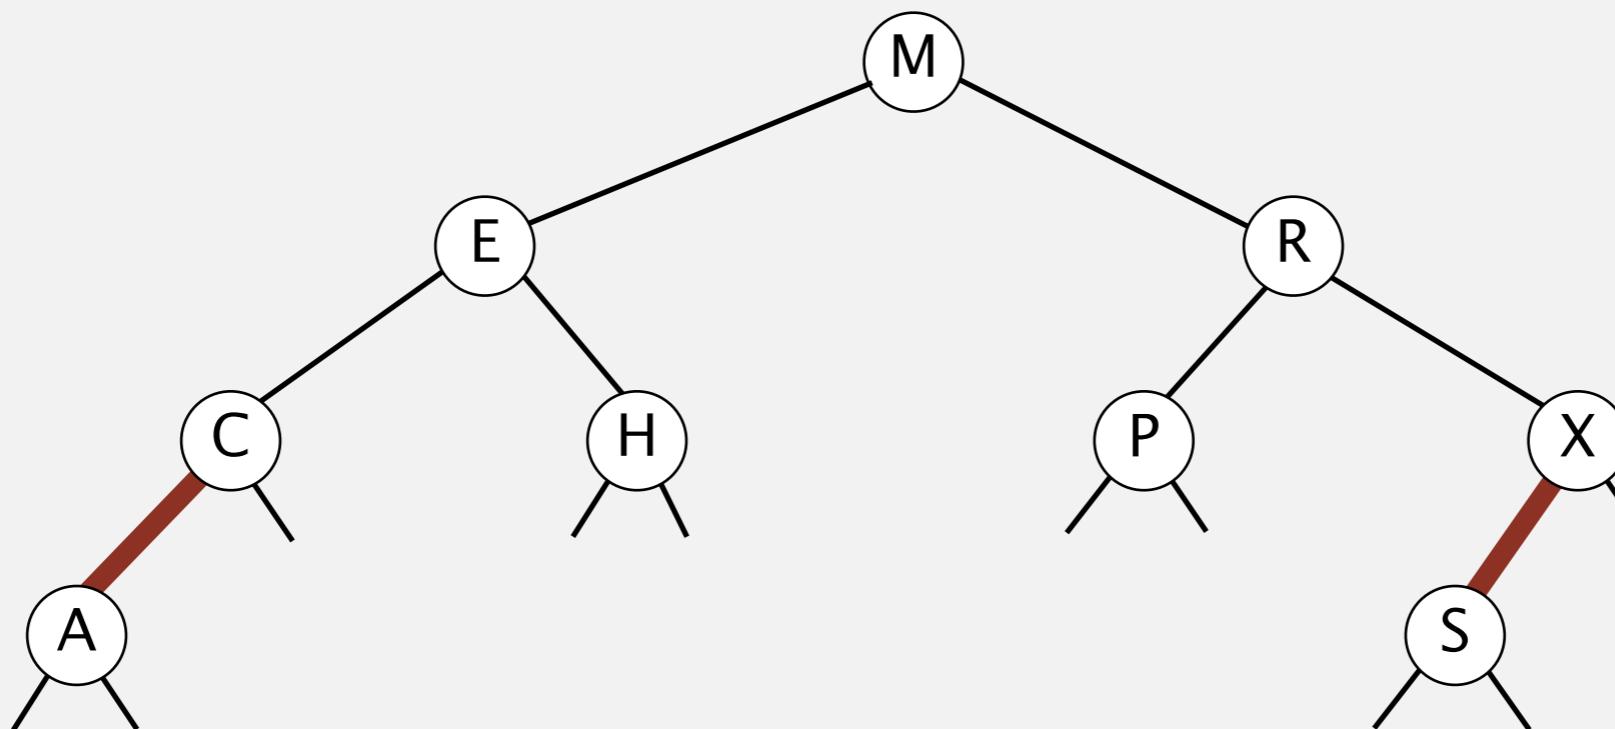
red-black BST



# Red-black BST construction demo

---

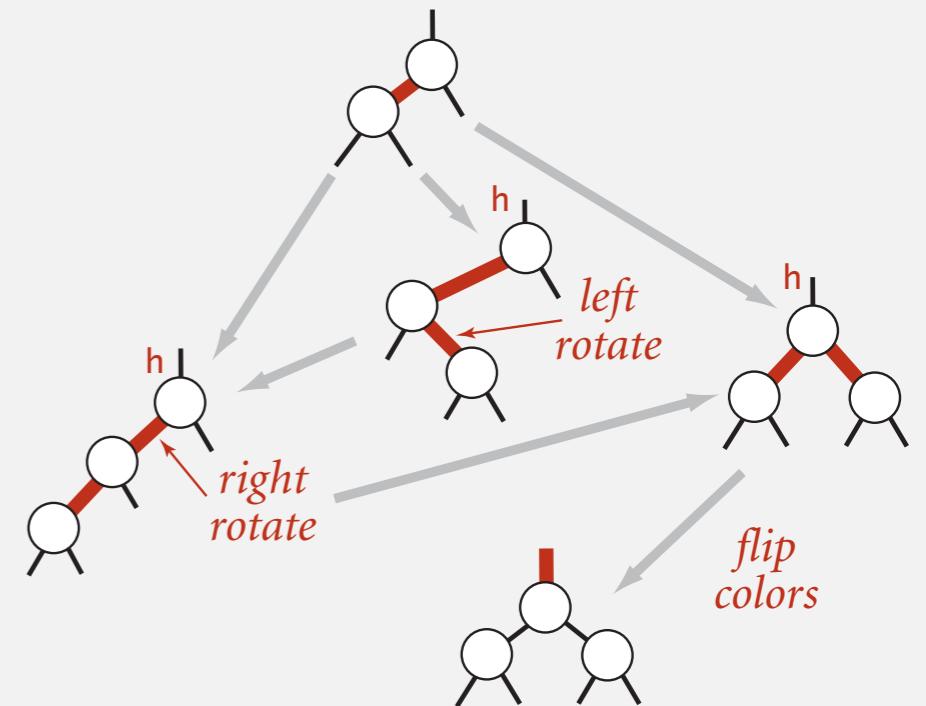
red-black BST



# Insertion into a LLRB tree: Java implementation

Same code for all cases.

- Right child red, left child black: **rotate left**.
- Left child, left-left grandchild red: **rotate right**.
- Both children red: **flip colors**.



```
private Node put(Node h, Key key, Value val)
{
    if (h == null) return new Node(key, val, RED);           ← insert at bottom
    int cmp = key.compareTo(h.key);                           (and color it red)
    if      (cmp < 0) h.left  = put(h.left,  key, val);
    else if (cmp > 0) h.right = put(h.right, key, val);
    else if (cmp == 0) h.val  = val;

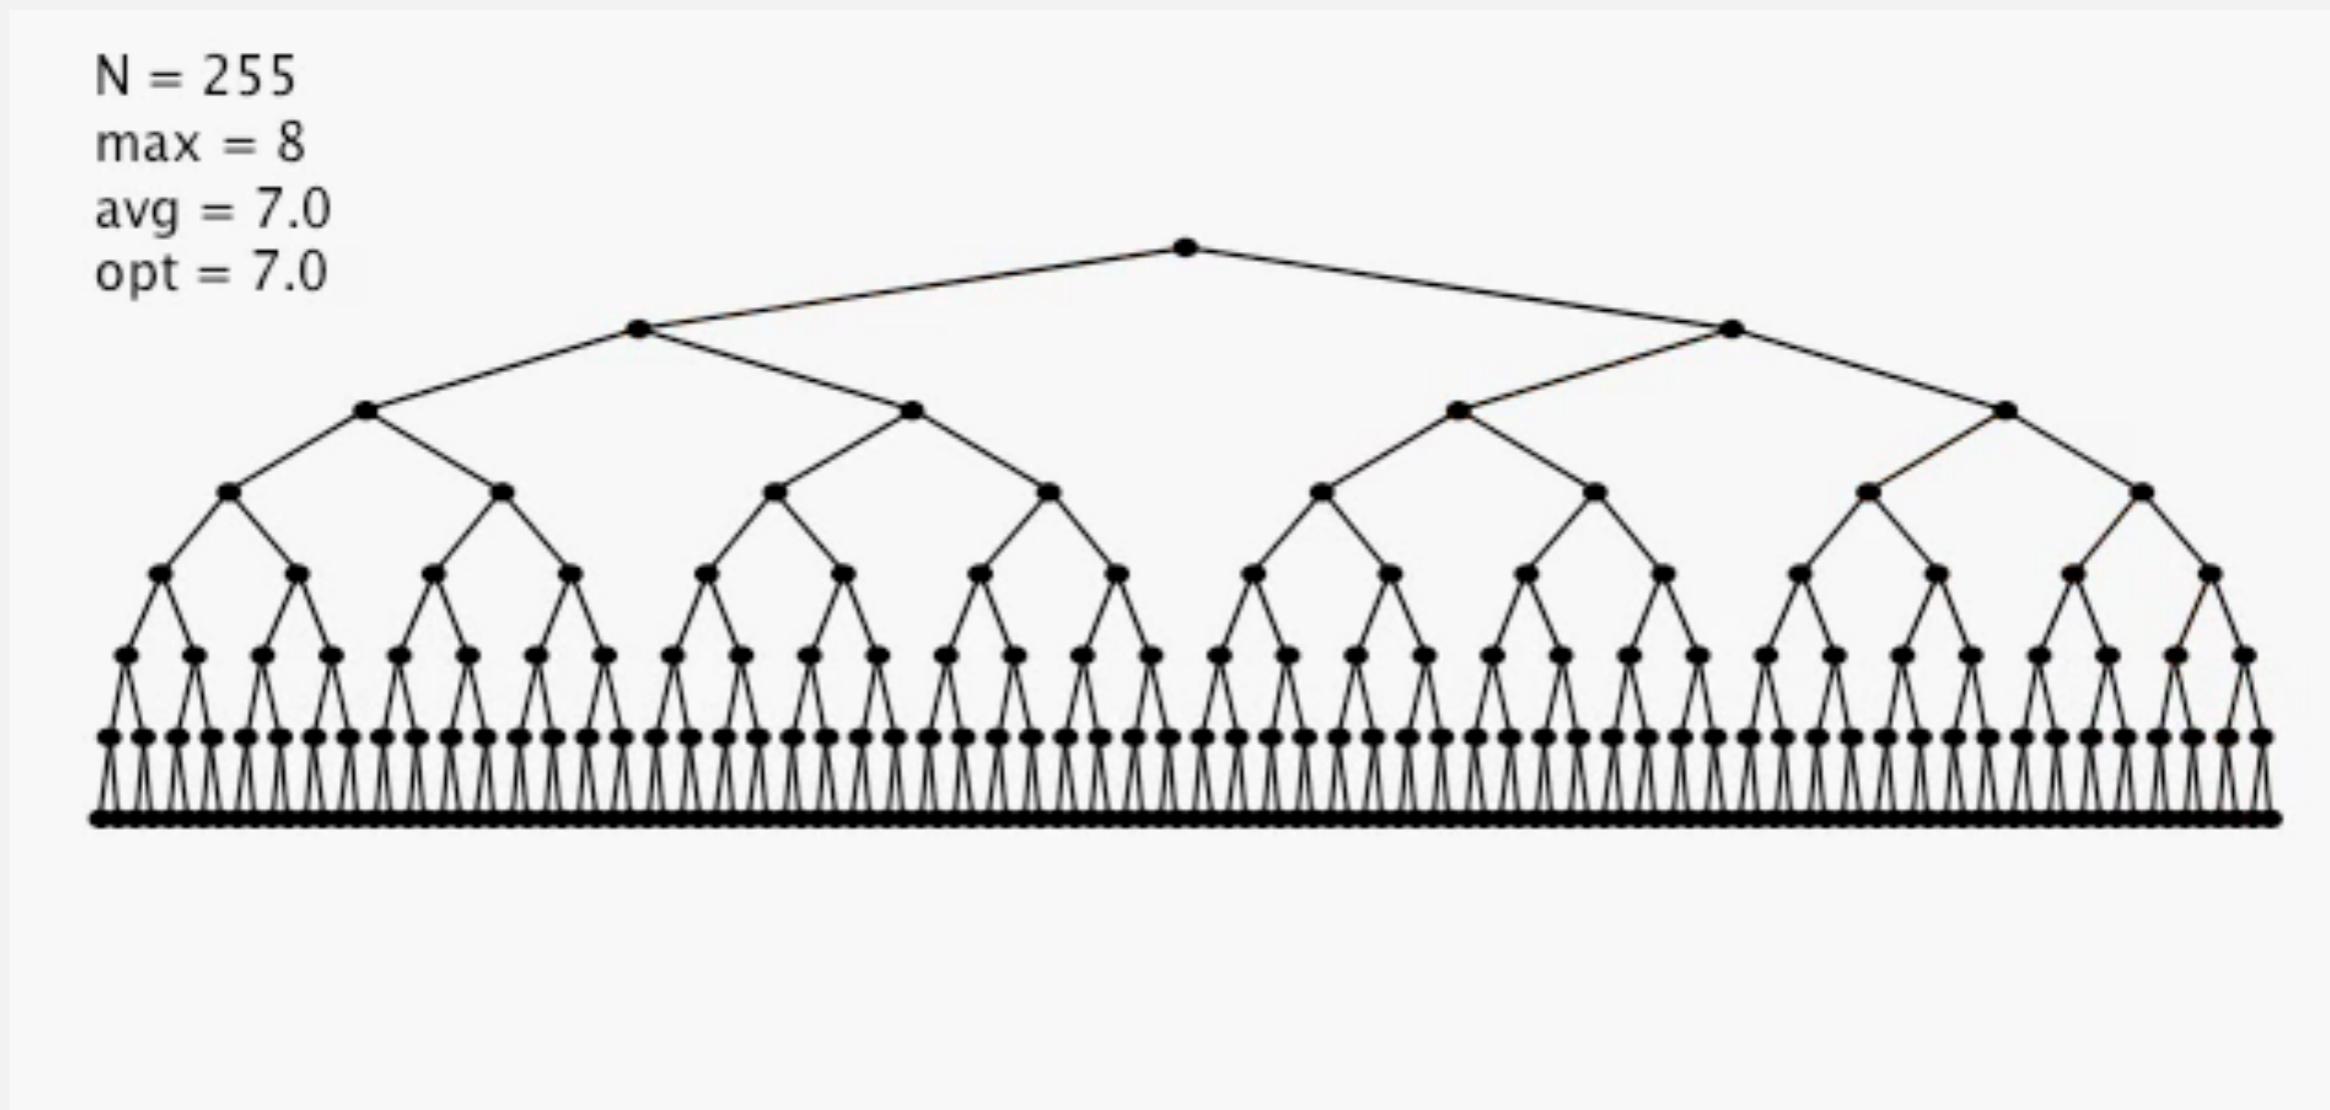
    if (isRed(h.right) && !isRed(h.left))     h = rotateLeft(h);   ← lean left
    if (isRed(h.left)  && isRed(h.left.left)) h = rotateRight(h);  ← balance 4-node
    if (isRed(h.left)  && isRed(h.right))     flipColors(h);       ← split 4-node

    return h;
}
```

only a few extra lines of code provides near-perfect balance

# Insertion into a LLRB tree: visualization

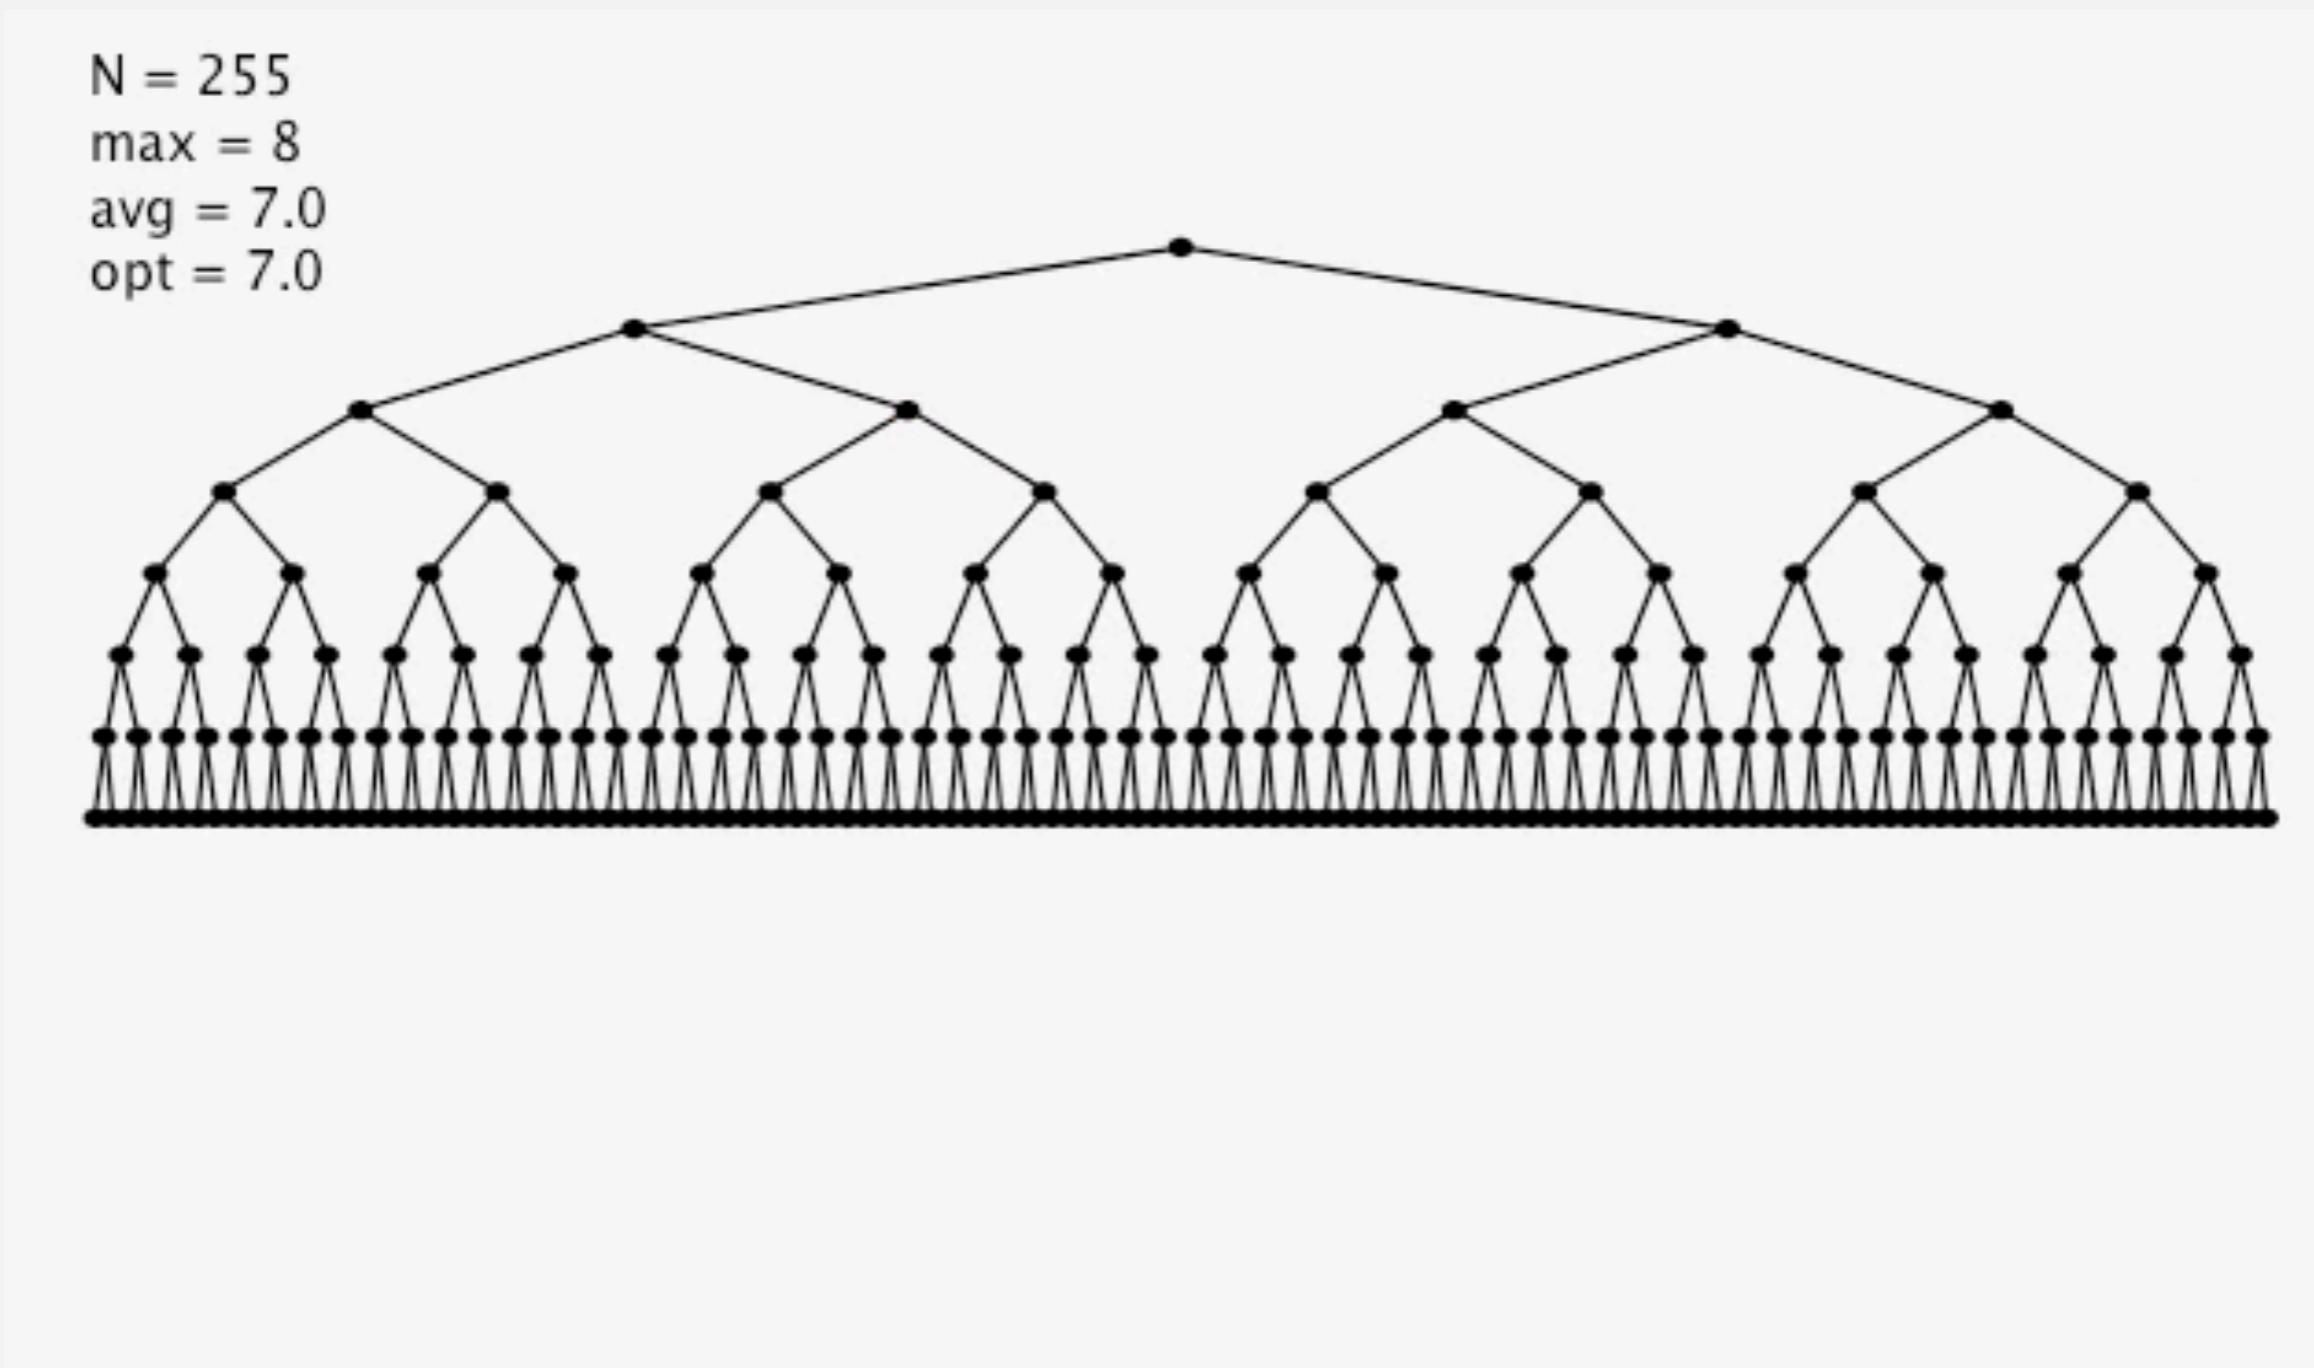
---



255 insertions in ascending order

# Insertion into a LLRB tree: visualization

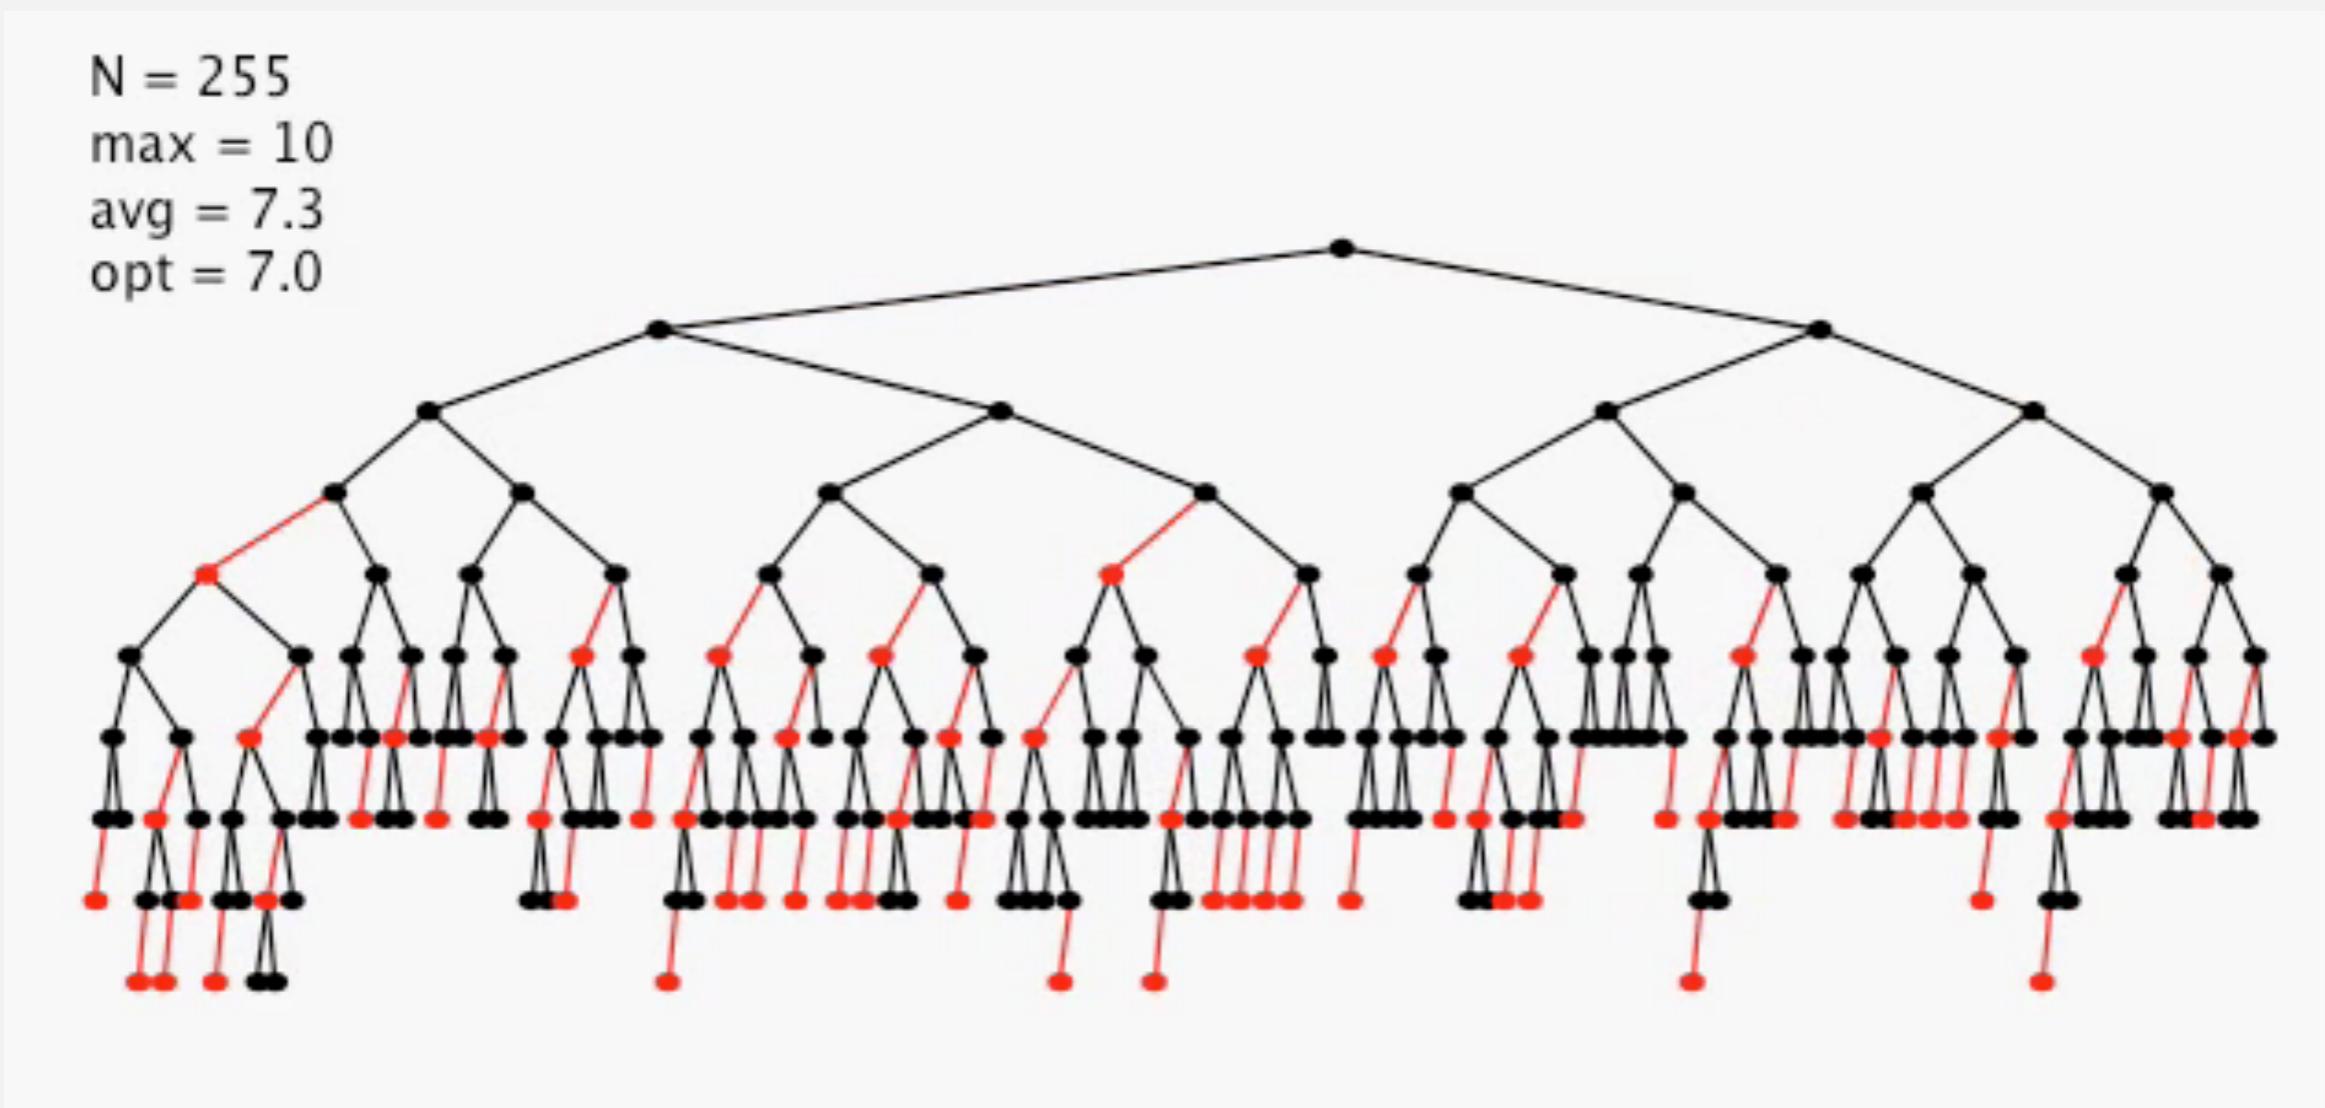
---



255 insertions in descending order

# Insertion into a LLRB tree: visualization

---



255 random insertions

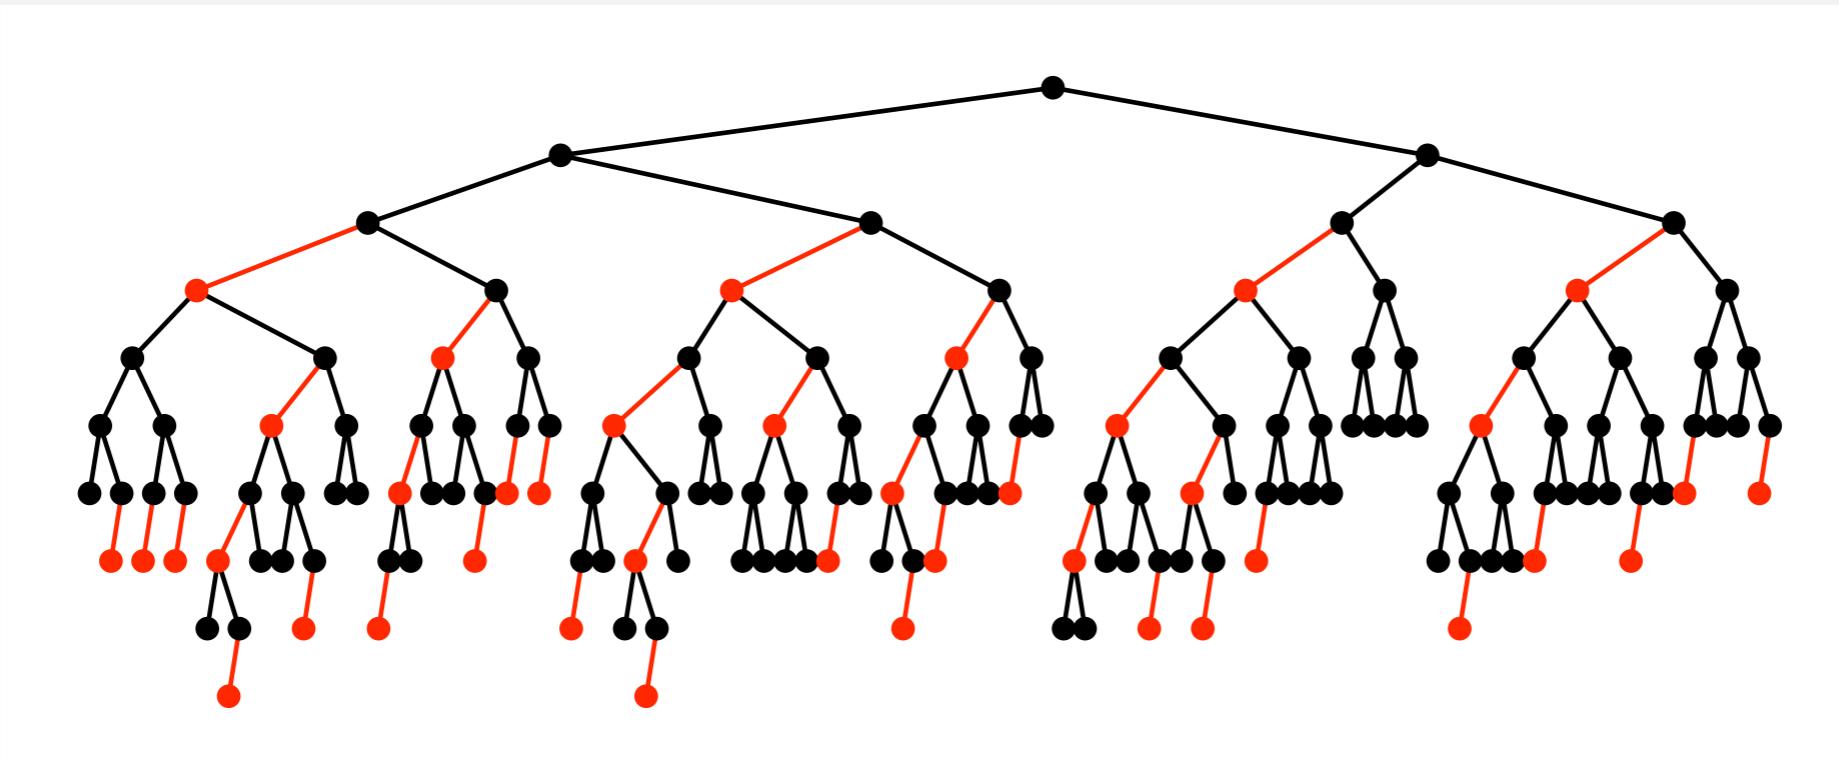
## Balance in LLRB trees

---

Proposition. Height of tree is  $\leq 2 \lg N$  in the worst case.

Pf.

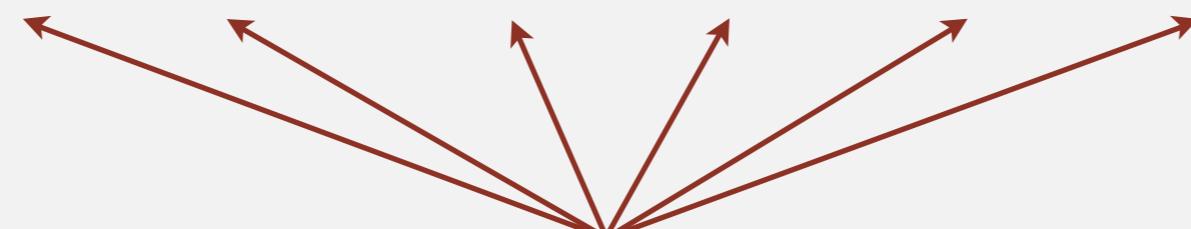
- Black height = height of corresponding 2–3 tree  $\leq \lg N$ .
- Never two red links in-a-row.



Property. Height of tree is  $\sim 1.0 \lg N$  in typical applications.

# ST implementations: summary

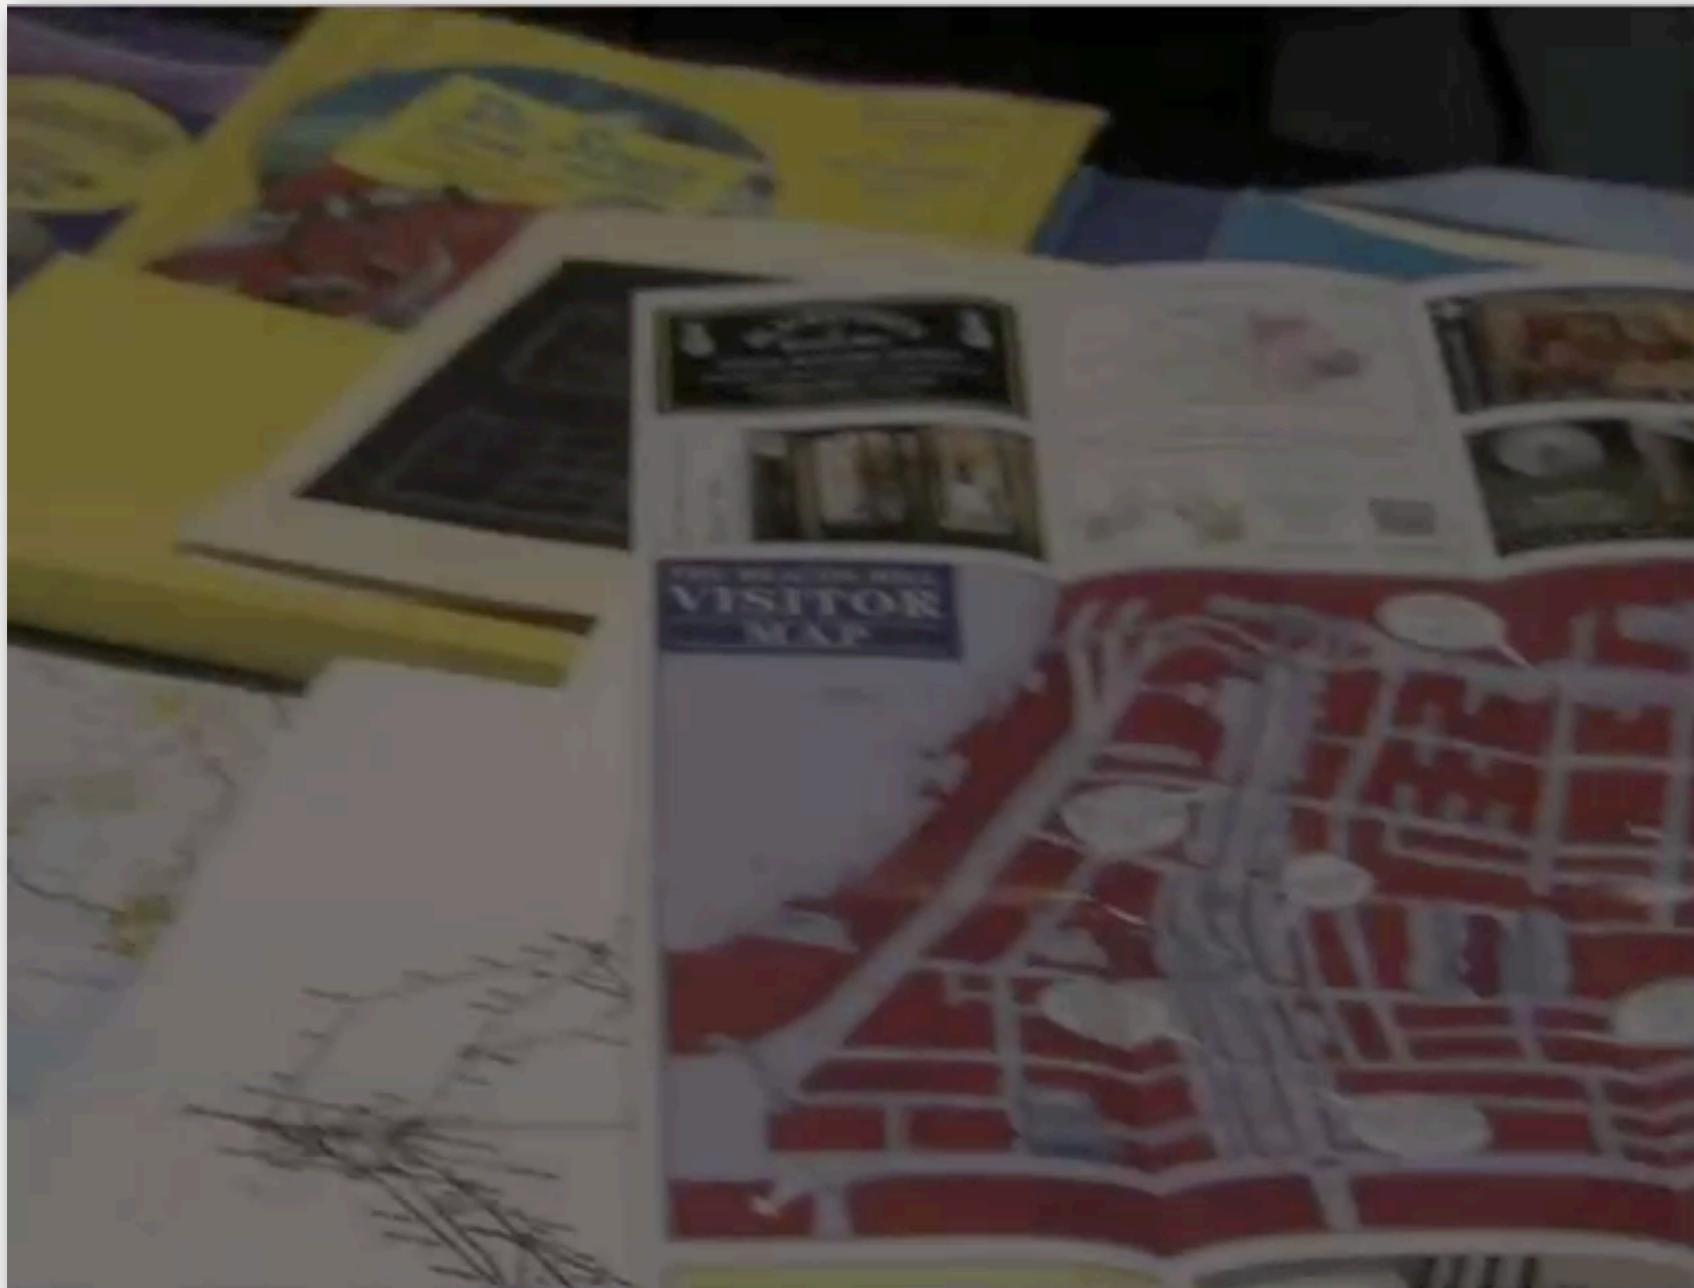
implementation	guarantee			average case			ordered ops?	key interface
	search	insert	delete	search hit	insert	delete		
<b>sequential search (unordered list)</b>	$N$	$N$	$N$	$N$	$N$	$N$		<code>equals()</code>
<b>binary search (ordered array)</b>	$\log N$	$N$	$N$	$\log N$	$N$	$N$	✓	<code>compareTo()</code>
<b>BST</b>	$N$	$N$	$N$	$\log N$	$\log N$	$\sqrt{N}$	✓	<code>compareTo()</code>
<b>2-3 tree</b>	$\log N$	$\log N$	$\log N$	$\log N$	$\log N$	$\log N$	✓	<code>compareTo()</code>
<b>red-black BST</b>	$\log N$	$\log N$	$\log N$	$\log N$	$\log N$	$\log N$	✓	<code>compareTo()</code>



hidden constant  $c$  is small  
(at most  $2 \lg N$  compares)

# Red-black BSTs in the wild

---



*Common sense. Sixth sense.  
Together they're the  
FBI's newest team.*

# Red-black BSTs in the wild

---

## ACT FOUR

FADE IN:

48 INT. FBI HQ - NIGHT

48

Antonio is at THE COMPUTER as Jess explains herself to Nicole and Pollock. The CONFERENCE TABLE is covered with OPEN REFERENCE BOOKS, TOURIST GUIDES, MAPS and REAMS OF PRINTOUTS.

JESS

It was the red door again.

POLLOCK

I thought the red door was the storage container.

JESS

But it wasn't red anymore. It was black.

ANTONIO

So red turning to black means... what?

POLLOCK

Budget deficits? Red ink, black ink?

NICOLE

Yes. I'm sure that's what it is. But maybe we should come up with a couple other options, just in case.

Antonio refers to his COMPUTER SCREEN, which is filled with mathematical equations.

ANTONIO

It could be an algorithm from a binary search tree. A red-black tree tracks every simple path from a node to a descendant leaf with the same number of black nodes.

JESS

Does that help you with girls?

# War story: why red-black?

## Xerox PARC innovations. [1970s]

- Alto.
- GUI.
- Ethernet.
- Smalltalk.
- InterPress.
- Laser printing.
- Bitmapped display.
- WYSIWYG text editor.
- ...



Xerox Alto

### A DICHROMATIC FRAMEWORK FOR BALANCED TREES

Leo J. Guibas  
*Xerox Palo Alto Research Center,  
Palo Alto, California, and  
Carnegie-Mellon University*

and  
Robert Sedgewick\*  
Program in Computer Science  
*Brown University*  
Providence, R. I.

#### ABSTRACT

In this paper we present a uniform framework for the implementation and study of balanced tree algorithms. We show how to imbed in this

the way down towards a leaf. As we will see, this has a number of significant advantages over the older methods. We shall examine a number of variations on a common theme and exhibit full implementations which are notable for their brevity. One implementation is examined carefully, and some properties about its

## War story: red-black BSTs

---

Telephone company contracted with database provider to build real-time database to store customer information.

### Database implementation.

- Red–Black BST.
- Exceeding height limit of 80 triggered error-recovery process.

show allow for up to  $2^{40}$  keys

### Extended telephone service outage.

- Main cause = height bound exceeded!
- Telephone company sues database provider.
- Legal testimony:

did not rebalance  
BST during delete



*“If implemented properly, the height of a red–black BST with  $N$  keys is at most  $2 \lg N$ . ” — expert witness*



## File system model

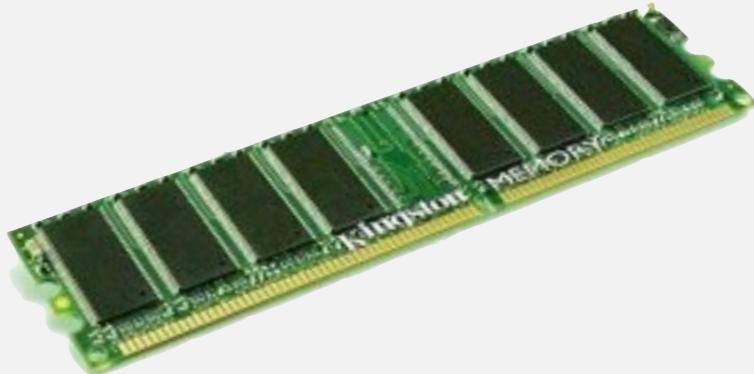
---

**Page.** Contiguous block of data (e.g., a 4,096-byte chunk).

**Probe.** First access to a page (e.g., from disk to memory).



slow



fast

**Property.** Time required for a probe is much larger than time to access data within a page.

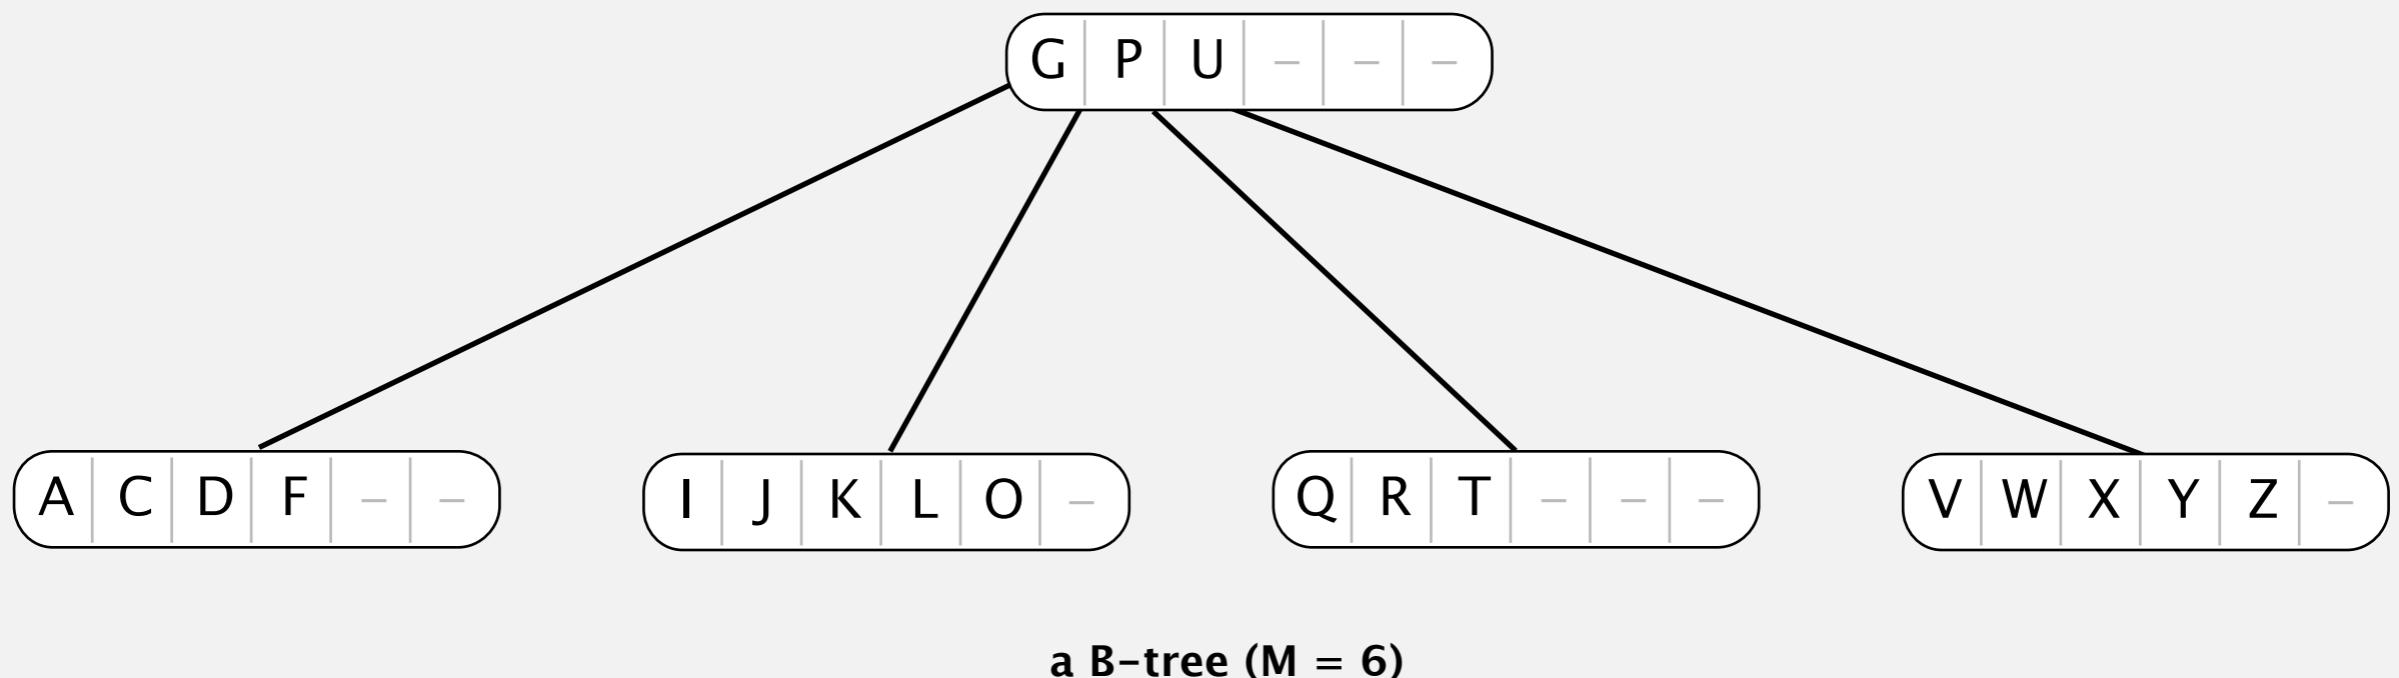
**Cost model.** Number of probes.

**Goal.** Access data using minimum number of probes.

**B-tree.** Generalize 2–3 trees by allowing up to  $M$  keys per node.

- At least  $\lfloor M/2 \rfloor$  keys in all nodes (except root).
- Every path from root to leaf has same number of links.

choose  $M$  as large as possible so that  $M$  keys fit in a page  
( $M = 1,024$  is typical)

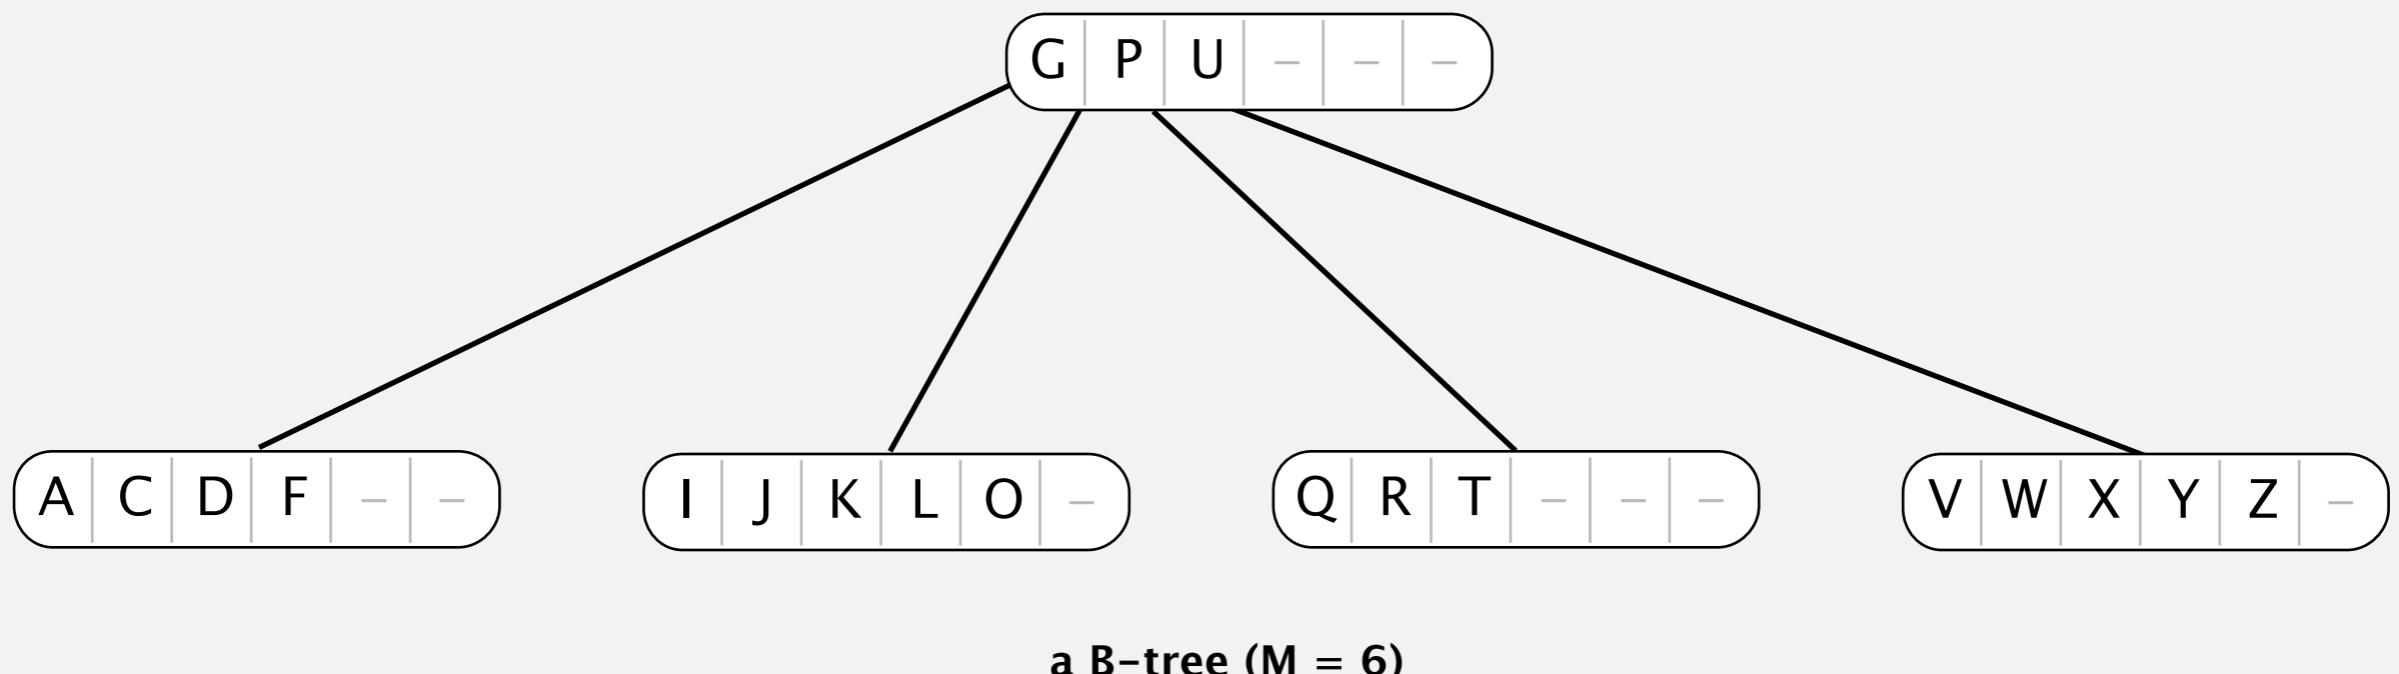


# Search in a B-tree

---

- Start at root.
- Check if node contains key.
- Otherwise, find interval for search key and take corresponding link.

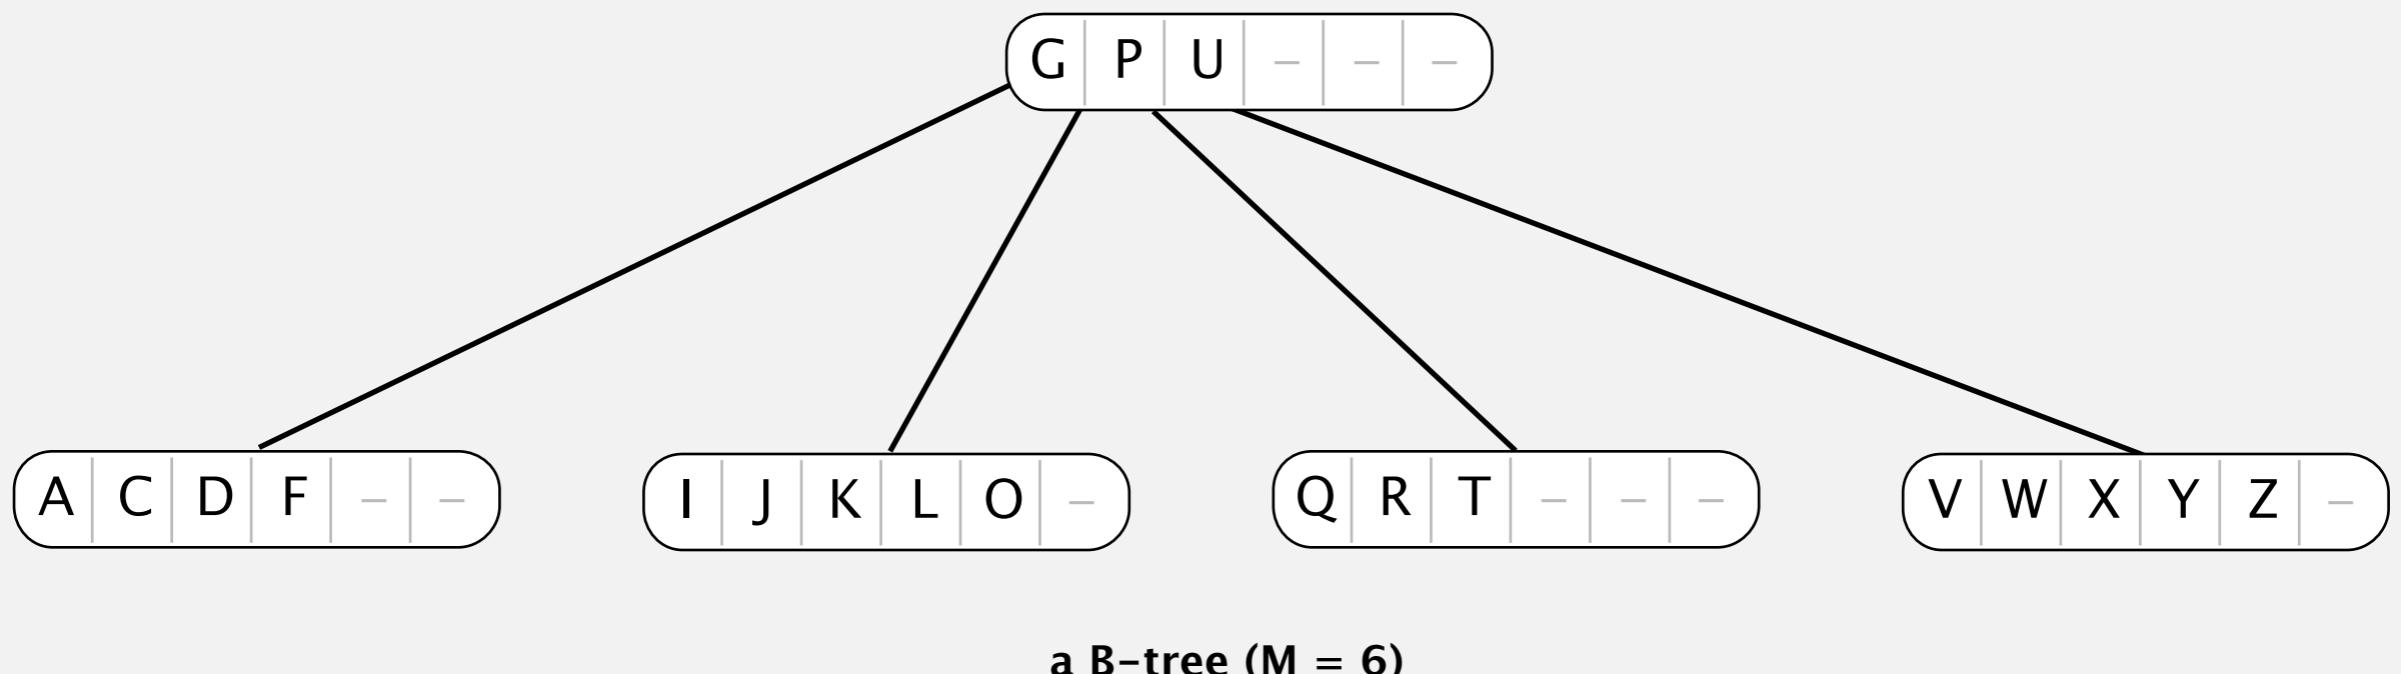
could use binary search  
(but all ops are considered free)



## Insertion in a B-tree

---

- Search for new key.
- Insert at bottom.
- Split nodes with  $M + 1$  keys on the way back up the B-tree  
(moving middle key to parent).



## Balance in B-tree

---

**Proposition.** A search or an insertion in a B-tree of order  $M$  with  $N$  keys requires between  $\sim \log_M N$  and  $\sim \log_{M/2} N$  probes.

**Pf.** All nodes (except possibly root) have between  $\lfloor M/2 \rfloor$  and  $M$  keys.

**In practice.** Number of probes is at most 4.  $\leftarrow$   $M = 1024$ ;  $N = 62$  billion  
 $\log_{M/2} N \leq 4$

# Balanced trees in the wild

---

Red–Black trees are widely used as system symbol tables.

- Java: `java.util.TreeMap`, `java.util.TreeSet`.
- C++ STL: `map`, `multimap`, `multiset`.
- Linux kernel: completely fair scheduler, `linux/rbtree.h`.
- Emacs: conservative stack scanning.

B-tree cousins. B+ tree, B\*tree, B# tree, ...

B-trees (and cousins) are widely used for file systems and databases.

- Windows: NTFS.
- Mac: HFS, HFS+.
- Linux: ReiserFS, XFS, Ext3FS, JFS, BTRFS.
- Databases: ORACLE, DB2, INGRES, SQL, PostgreSQL.

